MRI in Neuroscience and Neurology · WiSe 2026/27 · session 6

Brain region segmentations using FSL, FreeSurfer and Python

Lecture of 18 November 2026 · 53 slides

Grey matter (FreeSurfer, SynthSeg) · white matter lesions (BIANCA, U-Net, Dice) · microbleeds (U-Net, F1)

Online deck with voice, clips and the interactive pages:
mendeltem.github.io/mri-segmentation-lectures/lecture_1118/deck/

All example images come from public data sets; see the sources of each slide.

Contents

Opening

1. Three tasks and one lesson
2. Segmentation: one label per voxel

1 · Grey matter

3. What are brain volumes used for?
4. FreeSurfer: the cortex as a folded sheet
5. Cortical thickness, point by point
6. From folds to named regions
7. From scan to surface: five steps
8. One command, many hours
9. SynthSeg: trained on images that never existed
10. How one fake image is made
11. Any contrast, in minutes
12. FreeSurfer or SynthSeg?
13. How do we know it is right? The QC score
14. Try it: SynthSeg on a stroke scan interactive
15. FreeSurfer's warning number: Euler
16. The proof: test against a reference
17. Where it fails: stroke lesions
18. What to do about stroke lesions
19. Your turn: atlas quiz exercise

2 · White matter lesions

20. White matter hyperintensities
21. What is the lesion volume used for?
22. The reference: masks drawn by experts
23. BIANCA: ask the neighbours
24. From probability to mask: the cut-off
25. Try it: move the cut-off interactive
26. What matters when you train BIANCA
27. What matters when you use BIANCA
28. What BIANCA misses: small lesions
29. The U-Net in one picture
30. BIANCA or U-Net?
31. The Dice score
32. Even experts do not agree
33. Dice is hard on small lesions

34. Try it: shift the mask interactive
35. Your turn: segment the lesions exercise
36. Can you beat the algorithm?

3 · Microbleeds

37. Microbleeds: small dark dots
38. Why count microbleeds?
39. The network marks voxels. We need a count.
40. Clusters: 26 neighbours
41. Try it: neighbours and clusters interactive
42. F1: counting hits and misses
43. Try it: find the microbleeds interactive
44. Why the count matters: four
45. Microbleed or vessel? Scroll.
46. Your turn: find the microbleeds exercise
47. Can you beat the algorithm?

4 · New data, new scanner

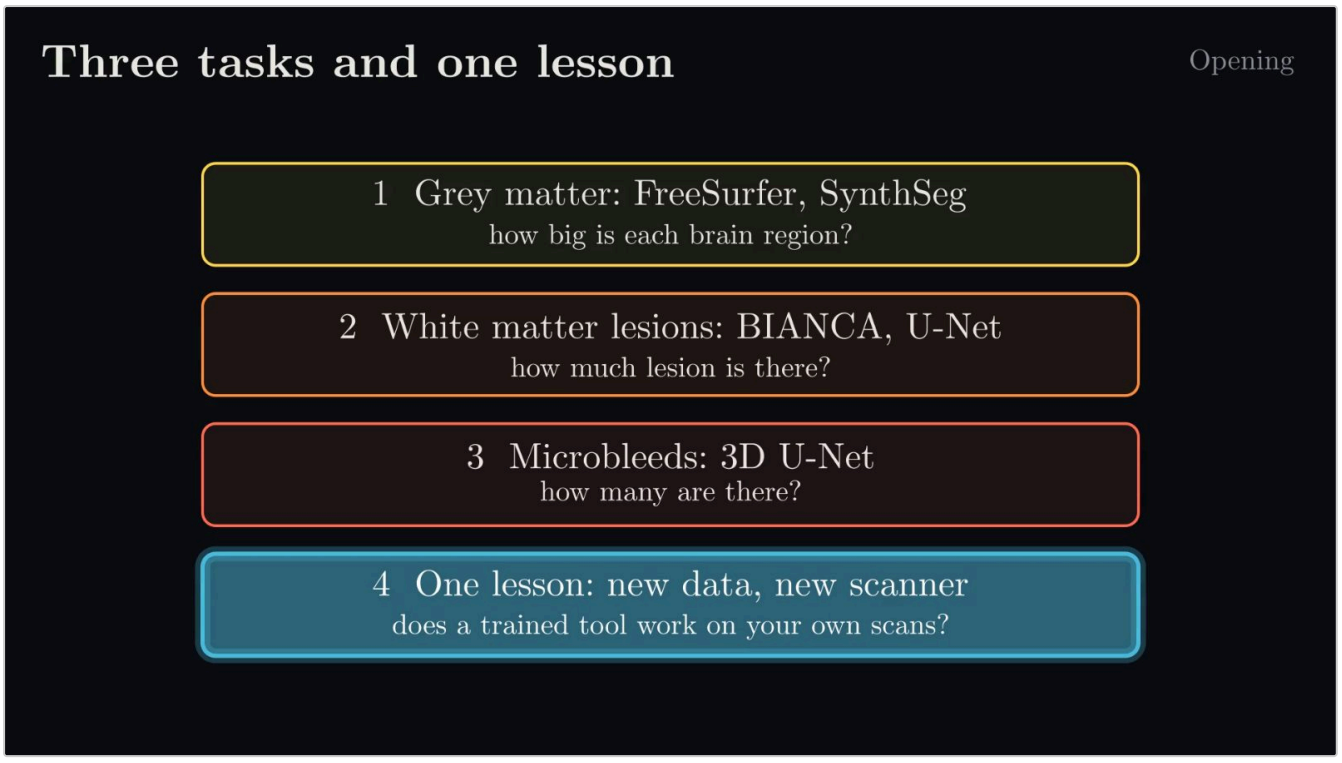
48. Does your data look like the training data?
49. Two ways to adapt
50. BIANCA: just add your scans to the pool
51. Fine-tuning: when you only have the model

Wrap-up

52. Conclusion
53. Films, slides and exercises stay online

All sources (48)

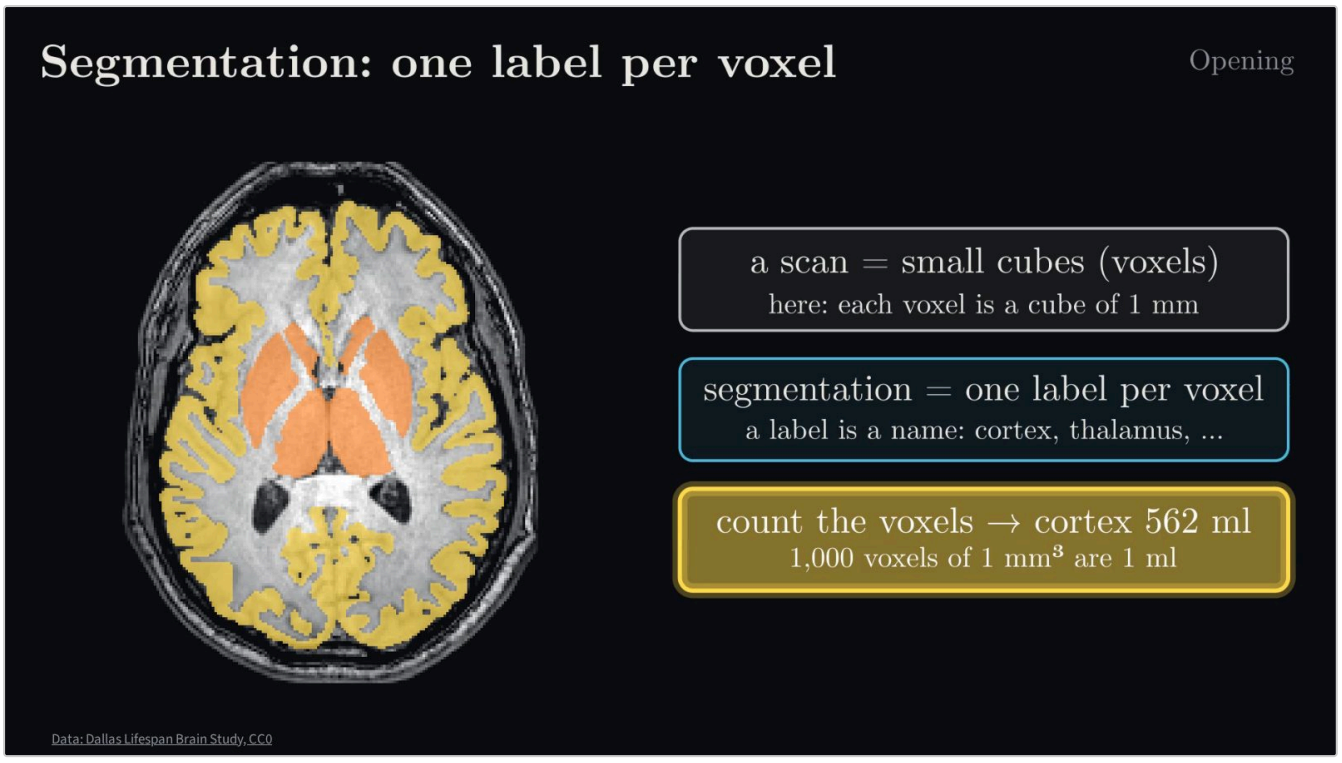
1. Three tasks and one lesson



KEY POINTS AND WHAT IS SAID

- **1 Grey matter: FreeSurfer, SynthSeg** – how big is each brain region?
  Three tasks today. Grey matter: how big is each brain region?
- **2 White matter lesions: BIANCA, U-Net** – how much lesion is there?
  White matter lesions: how much is there?
- **3 Microbleeds: 3D U-Net** – how many are there?
  And microbleeds: how many are there?
- **4 One lesson: new data, new scanner** – does a trained tool work on your own scans?
  And one lesson at the end: new data, new scanner.

2. Segmentation: one label per voxel



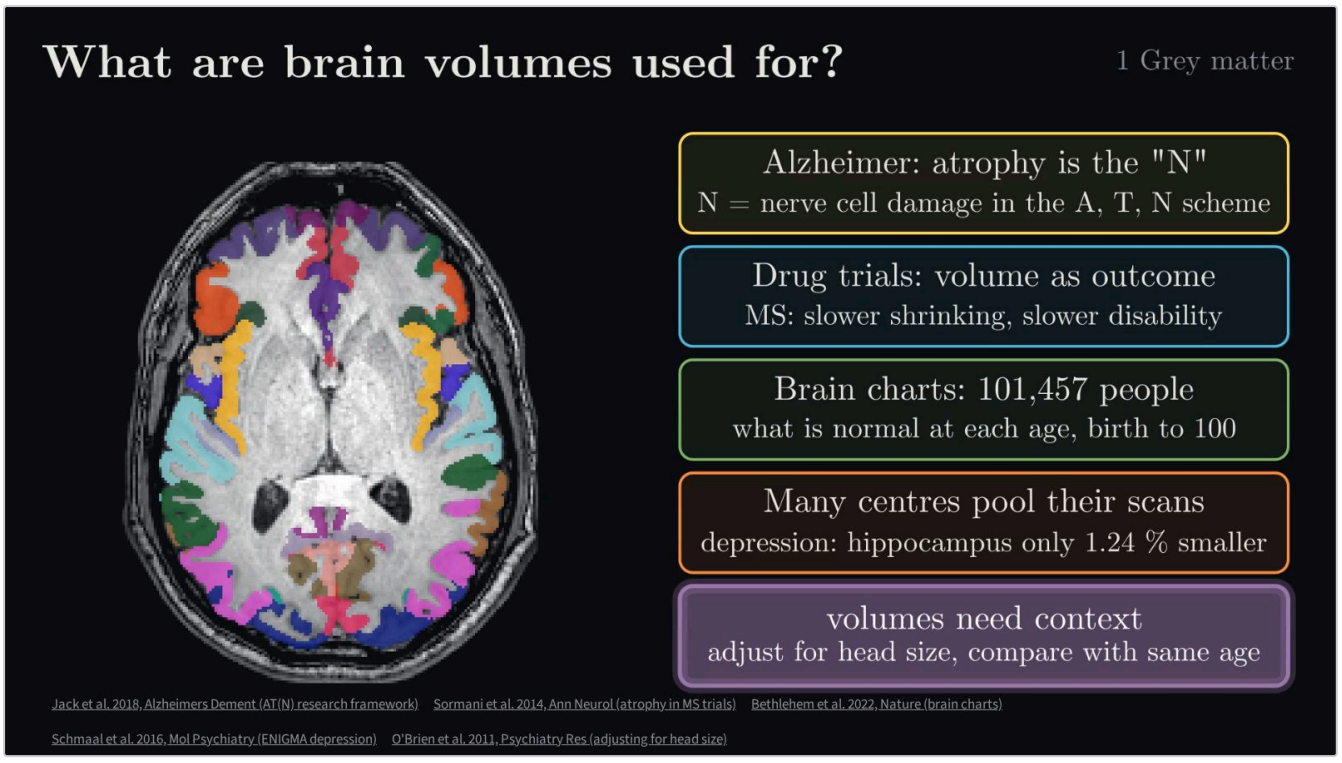
KEY POINTS AND WHAT IS SAID

- **a scan = small cubes (voxels)** – here: each voxel is a cube of 1 mm
  An MRI scan is made of small cubes, the voxels. Here, each one is a cube of one millimetre.
- **segmentation = one label per voxel** – a label is a name: cortex, thalamus, ...
  Segmentation gives every voxel a label: a name, like cortex or thalamus.
- **count the voxels → cortex 562 ml** – 1,000 voxels of 1 mm³ are 1 ml
  Count the voxels of one label, and you have its volume. A thousand voxels are one millilitre. This cortex: five hundred sixty two millilitres.

SOURCES

- Data: Dallas Lifespan Brain Study, CC0
  Dallas Lifespan Brain Study: public MRI data (CC0), used for our grey-matter pictures.

3. What are brain volumes used for?



KEY POINTS AND WHAT IS SAID

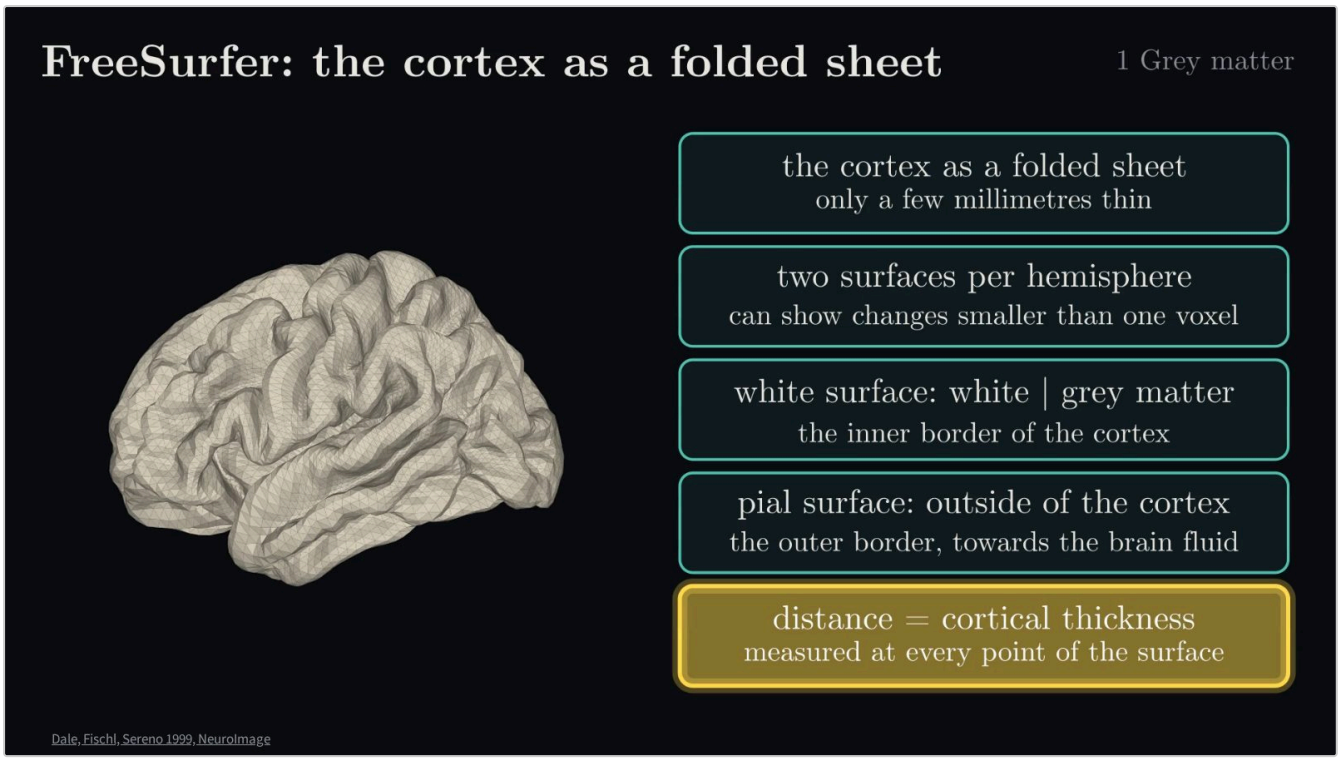
- **Alzheimer: atrophy is the "N"** – N = nerve cell damage in the A, T, N scheme
  Why measure volumes at all? Three real examples. In Alzheimer research, atrophy, the shrinking of the brain, is one of three markers of the disease.
- **Drug trials: volume as outcome** – MS: slower shrinking, slower disability
  Drug trials use it as an outcome. In multiple sclerosis, drugs that slowed the shrinking more also slowed disability more.
- **Brain charts: 101,457 people** – what is normal at each age, birth to 100
  Brain charts pool more than a hundred thousand people. Like growth charts for children, they show what is normal at each age.
- **Many centres pool their scans** – depression: hippocampus only 1.24 % smaller
  And many centres pool their scans. In depression, the hippocampus is about one percent smaller: too little to see in one patient, but clear in thousands.
- **volumes need context** – adjust for head size, compare with same age
  And a volume only means something corrected for head size, and compared with people of the same age.

SOURCES

- Jack et al. 2018, Alzheimers Dement (AT(N) research framework)
  Research definition of Alzheimer's disease by biomarkers: amyloid (A), tau (T), neurodegeneration (N). Atrophy on MRI is one (N) marker; not specific, and not meant for routine practice.
- Sormani et al. 2014, Ann Neurol (atrophy in MS trials)
  Meta-analysis of 13 multiple sclerosis trials (more than 13,500 patients): the more a drug slowed brain atrophy, the more it slowed disability (R2 = 0.48).

- Bethlehem et al. 2022, Nature (brain charts)
  Brain charts for the human lifespan: 123,984 scans of 101,457 people, from 115 days after conception to 100 years; gives centiles like a growth chart.
- Schmaal et al. 2016, Mol Psychiatry (ENIGMA depression)
  ENIGMA consortium: 1,728 patients with depression and 7,199 controls from 15 samples; hippocampus 1.24 % smaller (Cohen's d = -0.14).
- O'Brien et al. 2011, Psychiatry Res (adjusting for head size)
  Heads differ in size, so do brain structures: volumes must be adjusted for head (intracranial) size. Compares three ways (proportion, residual, covariance) and recommends a general model.

4. FreeSurfer: the cortex as a folded sheet



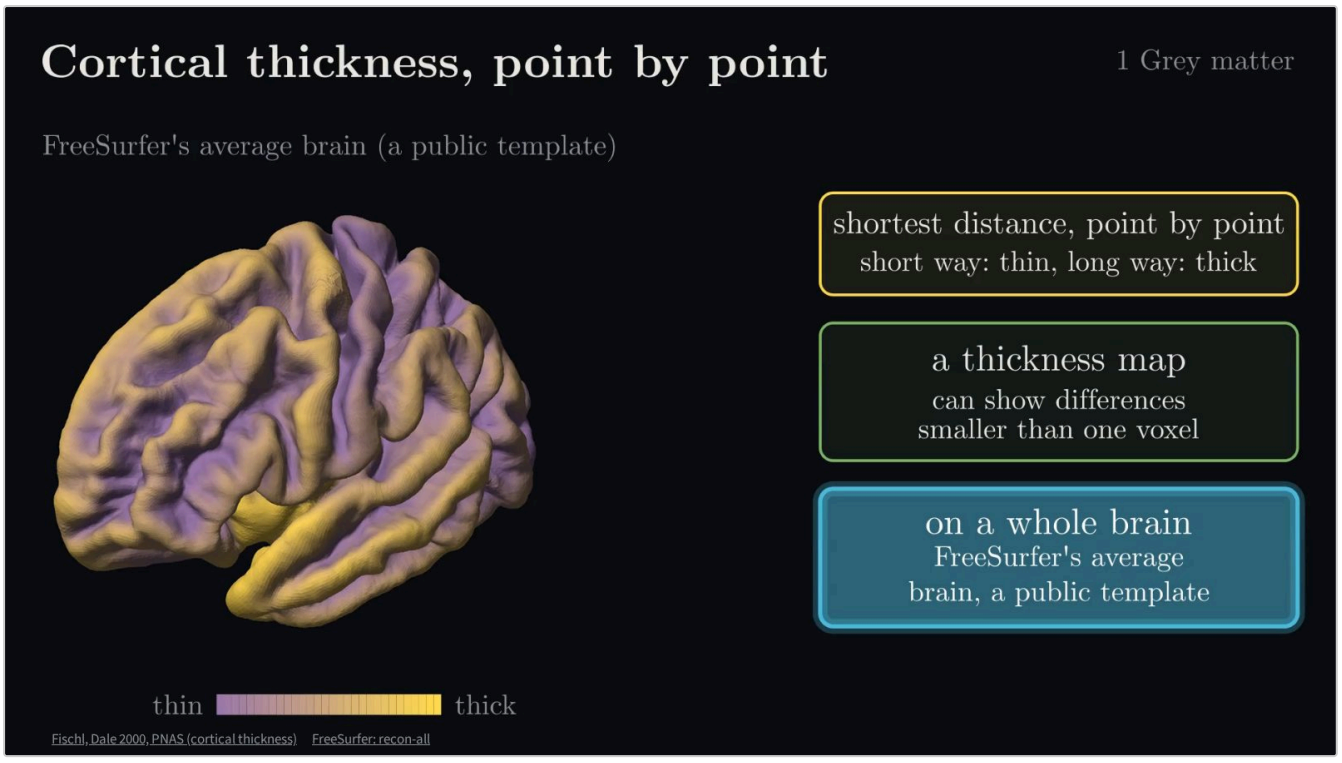
KEY POINTS AND WHAT IS SAID

- **the cortex as a folded sheet** – only a few millimetres thin
  FreeSurfer is the classic tool.
- **two surfaces per hemisphere** – can show changes smaller than one voxel
  It sees the cortex as a folded sheet, only a few millimetres thin. So it does not count voxels: it draws two surfaces.
- **white surface: white | grey matter** – the inner border of the cortex
  The white surface is the inner border, between white and grey matter.
- **pial surface: outside of the cortex** – the outer border, towards the brain fluid
  The pial surface is the outer border.
- **distance = cortical thickness** – measured at every point of the surface
  The distance between them is the cortical thickness.

SOURCES

- Dale, Fischl, Sereno 1999, NeuroImage
  FreeSurfer part I: segmentation and the white and pial surface of the cortex.

5. Cortical thickness, point by point



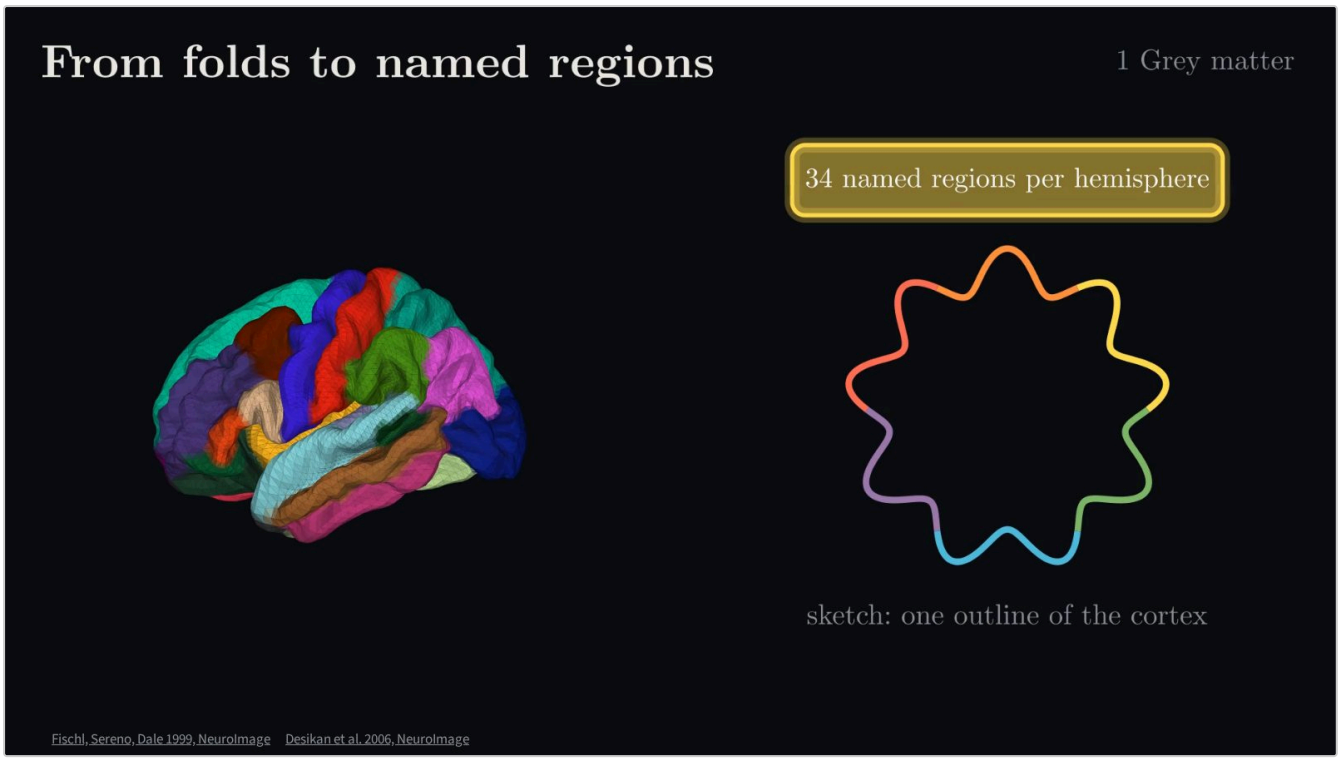
KEY POINTS AND WHAT IS SAID

- **shortest distance, point by point** – short way: thin, long way: thick
  How thick is the cortex at one point? From every point of one surface, FreeSurfer measures the shortest way to the other surface. A short way means thin cortex. A long way means thick cortex.
- **a thickness map** – can show differences smaller than one voxel
  Colour every point by its distance, and you get a thickness map. The surfaces run between the voxels. So the map can show differences smaller than one voxel. You see where the cortex gets thinner, not only that it shrinks.
- **on a whole brain** – FreeSurfer's average brain, a public template
  Now the same on a whole brain: FreeSurfer's average brain. The white surface grows outwards, to the pial surface. Then every point gets its colour. Purple is thin cortex. Yellow is thick cortex.

SOURCES

- Fischl, Dale 2000, PNAS (cortical thickness)
  How FreeSurfer measures cortical thickness from the two surfaces; differences between people below 0.5 mm.
- FreeSurfer: recon-all
  FreeSurfer documentation of recon-all, the command that runs the whole pipeline.

6. From folds to named regions



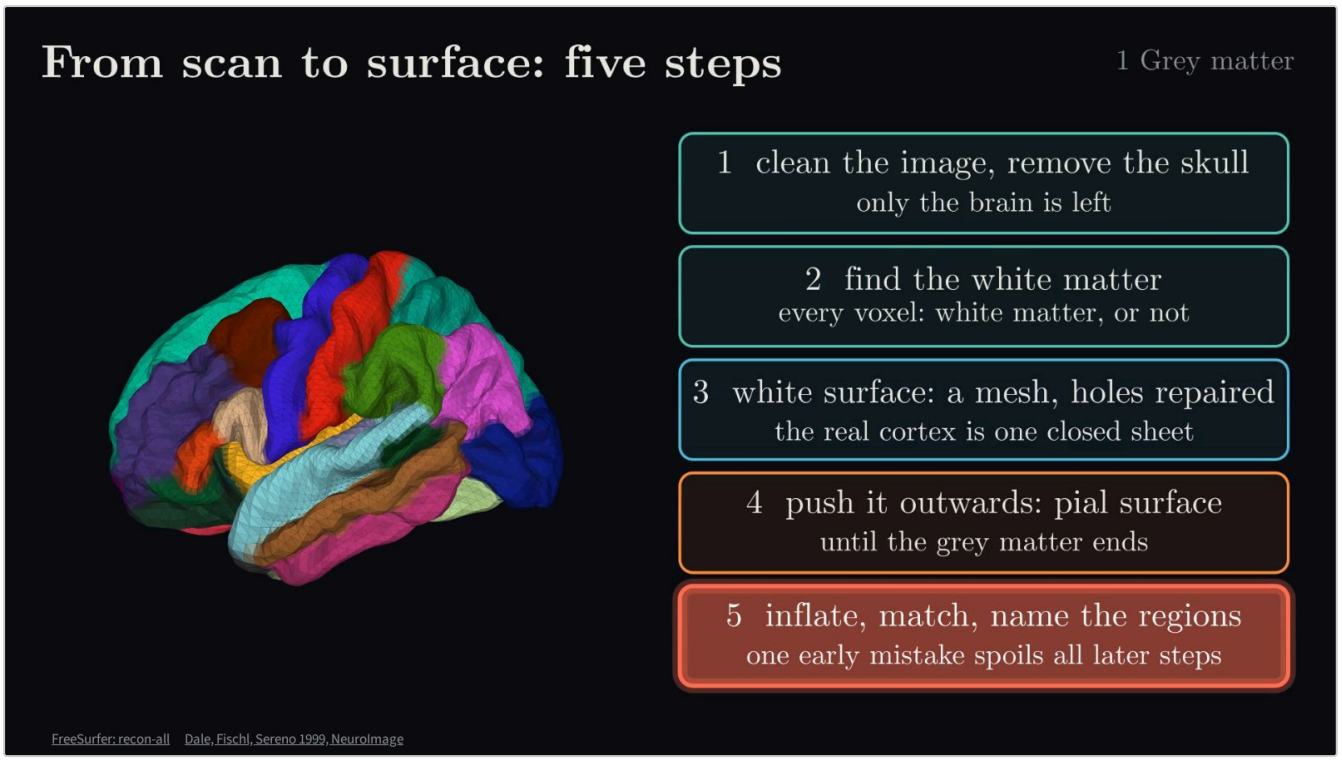
KEY POINTS AND WHAT IS SAID

- **inflate the surface**
  To name the folds, FreeSurfer blows the surface up
- **… then a sphere**
  into a sphere.
- **match the folds to an average brain**
  There it matches the folds to an average brain.
- **every fold gets a name: 34 regions per hemisphere**
  So every fold gets a name: thirty four regions per hemisphere.

SOURCES

- Fischl, Sereno, Dale 1999, NeuroImage
  FreeSurfer part II: the cortex inflated and mapped to a sphere; folding patterns aligned between brains.
- Desikan et al. 2006, NeuroImage
  Automatic labelling of 34 cortical regions per hemisphere, learned from 40 hand-labelled brains.

7. From scan to surface: five steps



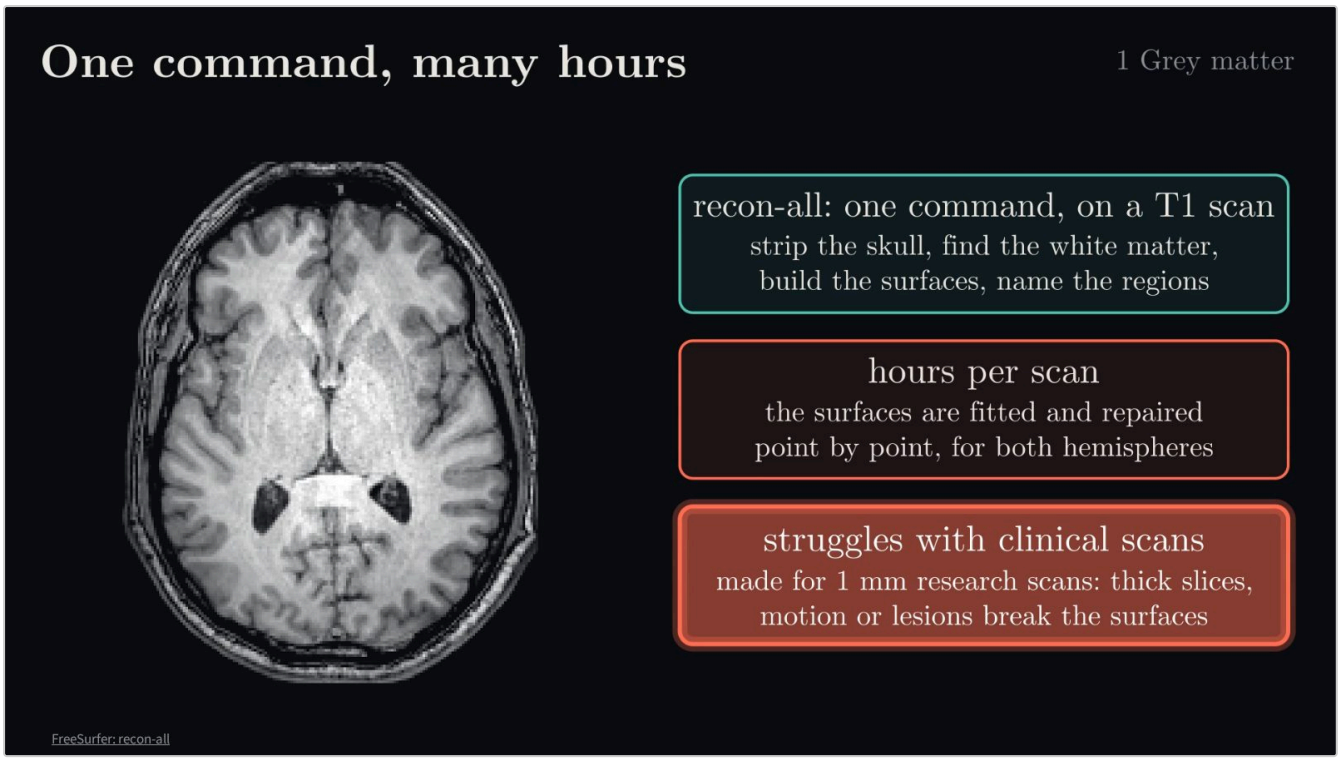
KEY POINTS AND WHAT IS SAID

- **1 clean the image, remove the skull** – only the brain is left
  How does FreeSurfer get from a scan to these surfaces? In a fixed chain of steps. First it cleans the image and removes the skull. Only the brain is left.
- **2 find the white matter** – every voxel: white matter, or not
  Then every voxel is sorted: white matter, or not. The white matter becomes one solid body in each hemisphere.
- **3 white surface: a mesh, holes repaired** – the real cortex is one closed sheet
  Around this body it lays a mesh of small triangles: the white surface. Holes are repaired, because the real cortex is one closed sheet.
- **4 push it outwards: pial surface** – until the grey matter ends
  This surface is pushed outwards, until the grey matter ends. That is the pial surface.
- **5 inflate, match, name the regions** – one early mistake spoils all later steps
  Last, the surface is inflated, matched to the average brain, and the regions get their names. Every step needs the one before. So one early mistake, such as a piece of skull left behind, spoils everything after it.

SOURCES

- FreeSurfer: recon-all
  FreeSurfer documentation of recon-all, the command that runs the whole pipeline.
- Dale, Fischl, Sereno 1999, NeuroImage
  FreeSurfer part I: segmentation and the white and pial surface of the cortex.

8. One command, many hours



KEY POINTS AND WHAT IS SAID

- **recon-all: one command, on a T1 scan** – strip the skull, find the white matter, build the surfaces, name the regions

  One command does it all: recon all, on a T one scan. Behind it are many steps, one after the other.

- **hours per scan** – the surfaces are fitted and repaired point by point, for both hemispheres

  It needs time: hours per scan, because the surfaces are fitted and repaired point by point.

- **struggles with clinical scans** – made for 1 mm research scans: thick slices, motion or lesions break the surfaces

  And it struggles with clinical scans. It was made for research scans, with thin slices and good contrast. Thick slices, motion or lesions break the surfaces.

SOURCES

- FreeSurfer: recon-all
  FreeSurfer documentation of recon-all, the command that runs the whole pipeline.

9. SynthSeg: trained on images that never existed



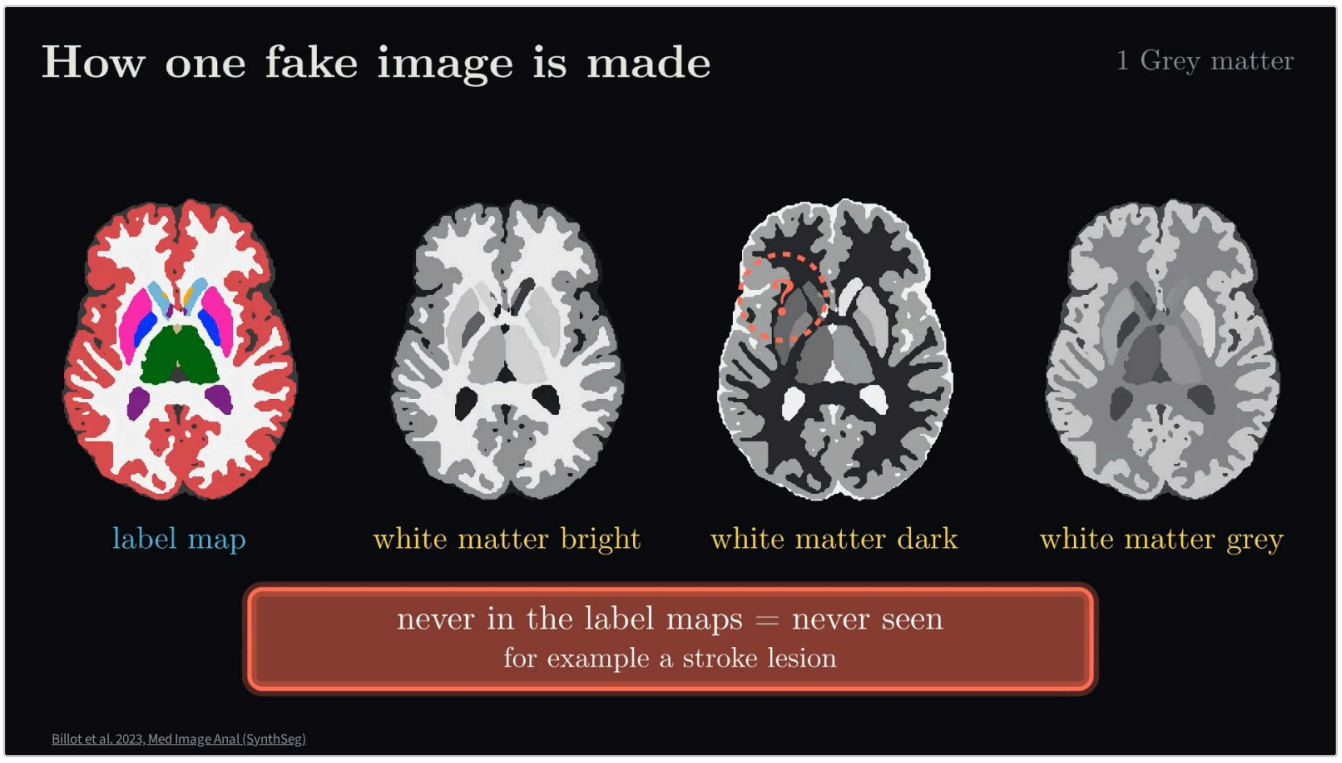
KEY POINTS AND WHAT IS SAID

- **a neural network that never saw a real scan**
  SynthSeg is a neural network that never saw a real scan.
- **trained on fake images: random contrast**
  It trained on fake images with random contrast and resolution.
- **a real scan = one more variation**
  So a real scan is just one more variation.

SOURCES

- Billot et al. 2023, Med Image Anal (SynthSeg)
  SynthSeg: a network trained only on synthetic images; segments scans of any contrast and resolution.

10. How one fake image is made



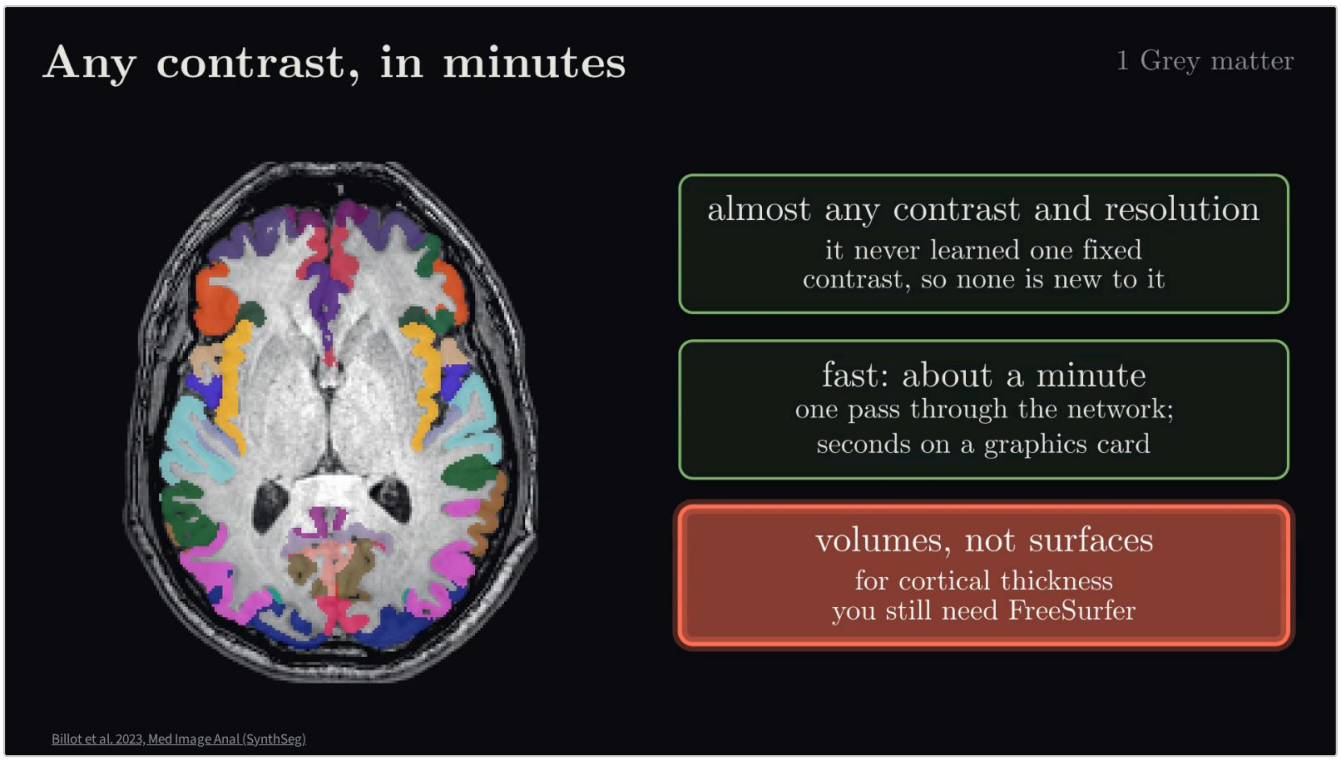
KEY POINTS AND WHAT IS SAID

- **start: only a label map** – no scan needed (one slice of a public scan)
  How is one fake image made? It starts with a label map of a brain: no scan, only the labels.
- **a random grey value per label** – white matter bright, dark or grey
  Then every label gets a random grey value. White matter may be bright, dark, or grey. So the contrast tells the network nothing.
- **the contrast tells nothing** – only the shape does
  Only the shape does.
- **never in the label maps = never seen** – for example a stroke lesion
  What was never in the label maps, the network has never seen, for example a stroke lesion.

SOURCES

- Billot et al. 2023, Med Image Anal (SynthSeg)
  SynthSeg: a network trained only on synthetic images; segments scans of any contrast and resolution.

11. Any contrast, in minutes



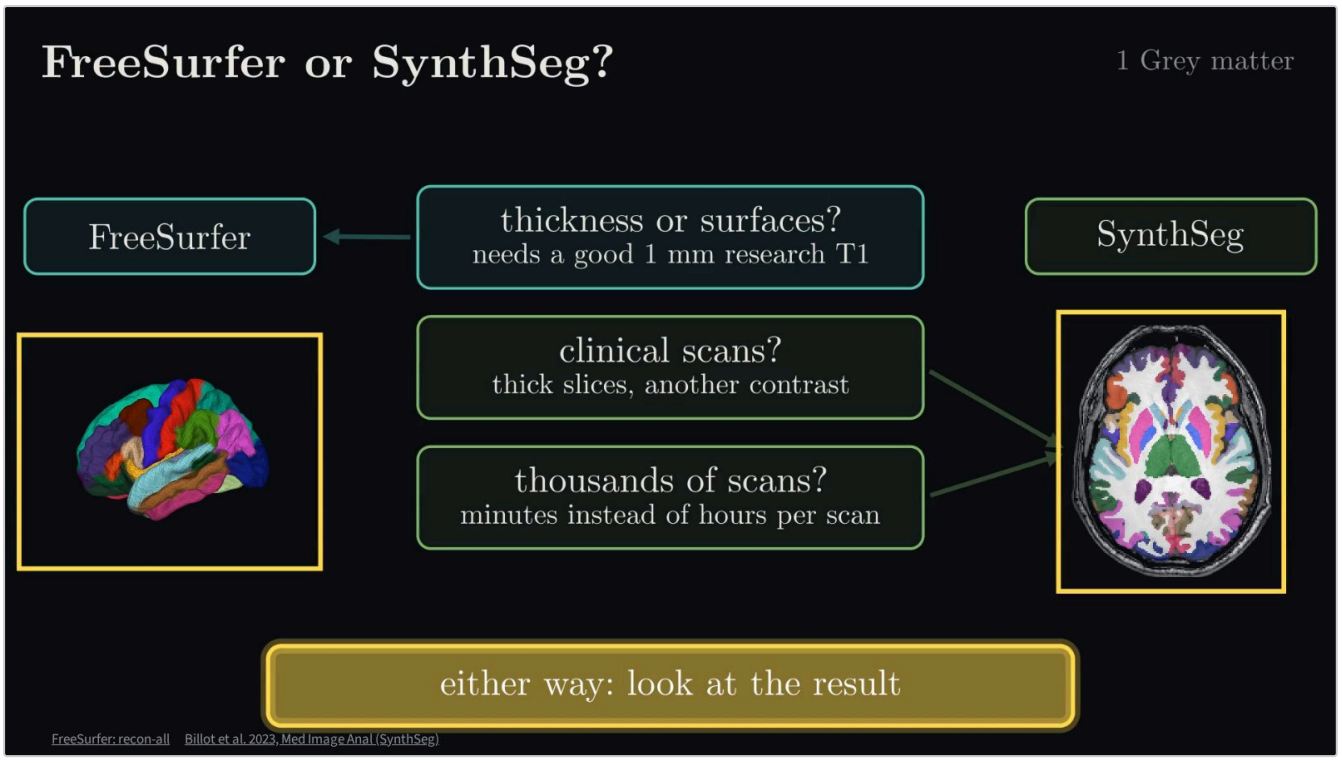
KEY POINTS AND WHAT IS SAID

- **almost any contrast and resolution** – it never learned one fixed contrast, so none is new to it
  Because of this training, it works on almost any contrast and resolution, without retraining.
- **fast: about a minute** – one pass through the network; seconds on a graphics card
  And it is fast: the scan goes through the network once. About a minute on a normal computer.
- **volumes, not surfaces** – for cortical thickness you still need FreeSurfer
  But it gives volumes, not surfaces. For cortical thickness, you still need FreeSurfer.

SOURCES

- Billot et al. 2023, Med Image Anal (SynthSeg)
  SynthSeg: a network trained only on synthetic images; segments scans of any contrast and resolution.

12. FreeSurfer or SynthSeg?



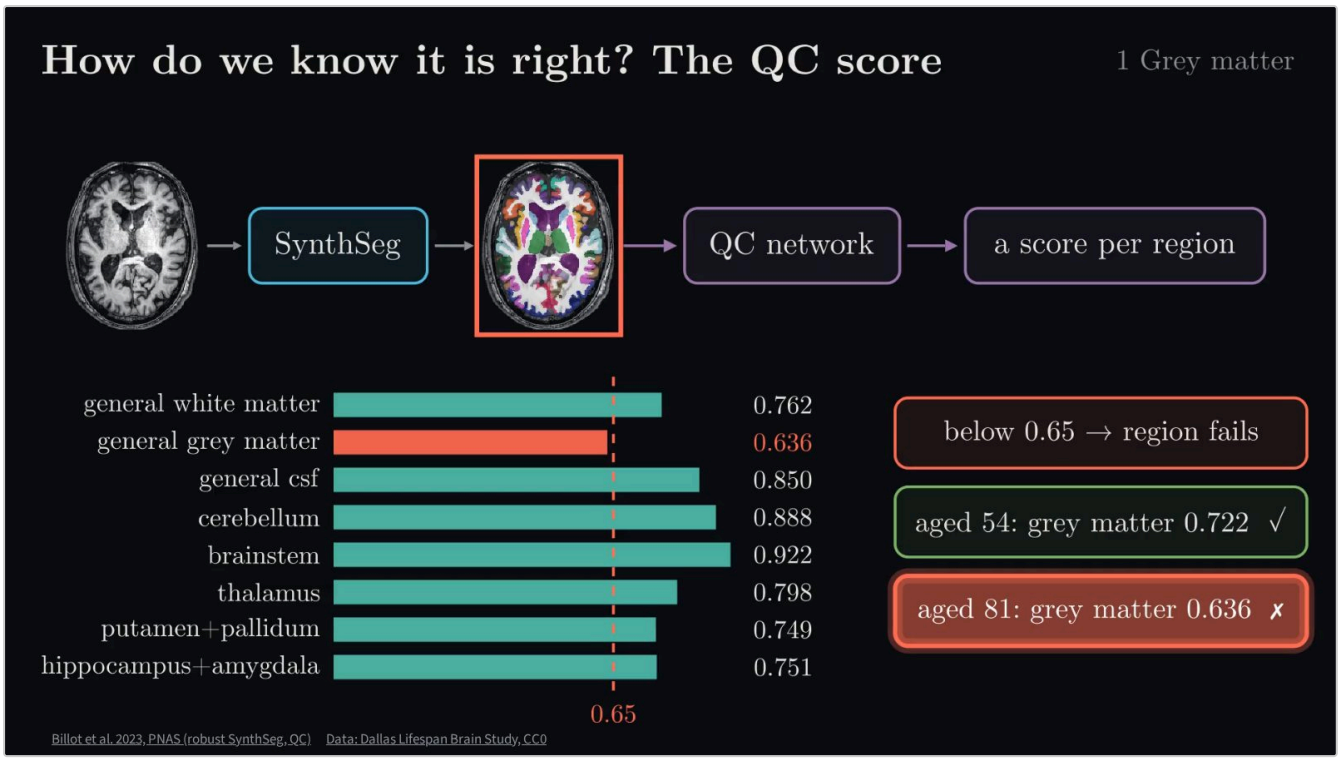
KEY POINTS AND WHAT IS SAID

- **thickness or surfaces? FreeSurfer** – needs a good 1 mm research T1
  FreeSurfer or SynthSeg? Ask three questions. One: do you need cortical thickness, or surfaces? Then FreeSurfer. And then you need a good research scan.
- **clinical scans? SynthSeg** – thick slices, another contrast
  Two: are your scans clinical, with thick slices or another contrast? Then SynthSeg.
- **thousands of scans? SynthSeg** – minutes instead of hours per scan
  Three: do you have thousands of scans, and little time? Also SynthSeg.
- **either way: look at the result**
  And whichever you choose: look at the result. Both tools always give an answer. Not always a right one.

SOURCES

- FreeSurfer: recon-all
  FreeSurfer documentation of recon-all, the command that runs the whole pipeline.
- Billot et al. 2023, Med Image Anal (SynthSeg)
  SynthSeg: a network trained only on synthetic images; segments scans of any contrast and resolution.

13. How do we know it is right? The QC score



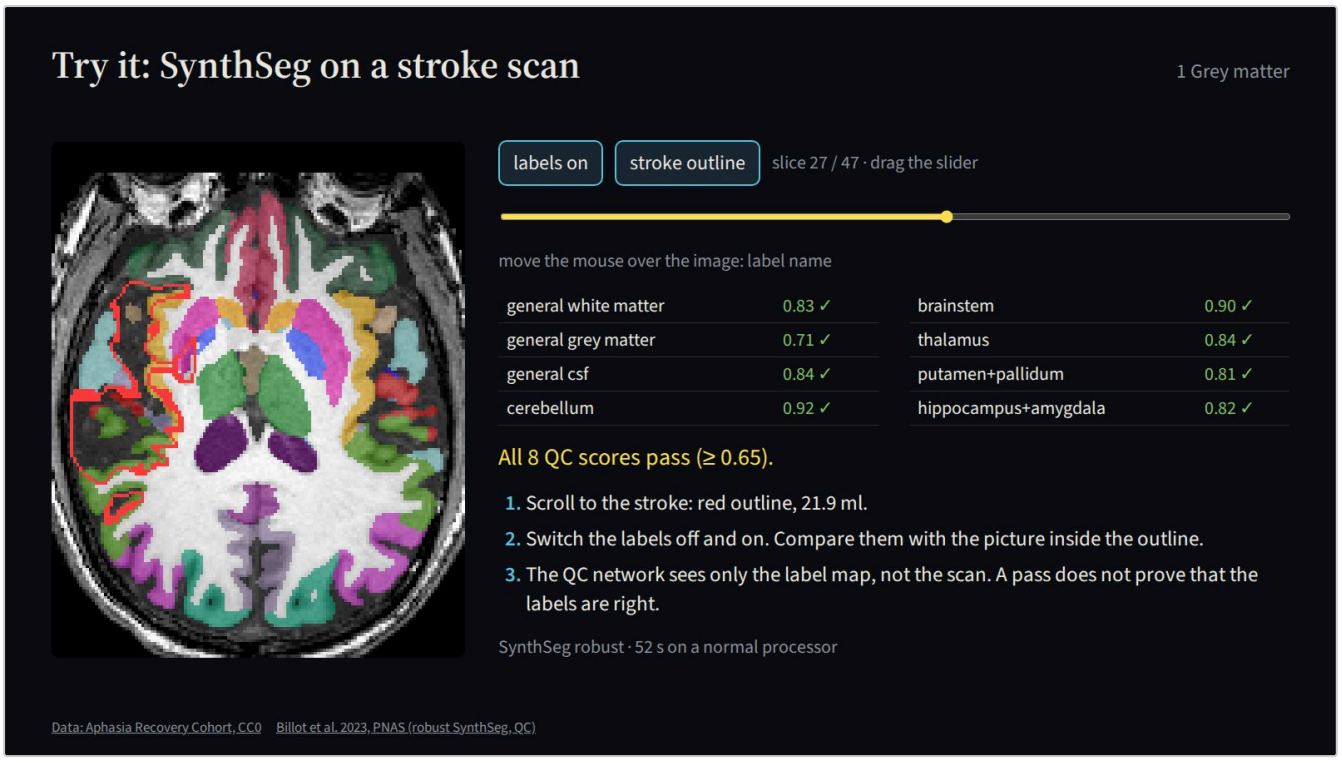
KEY POINTS AND WHAT IS SAID

- **a second small network gives the score**
  How do we know a segmentation is right? Nobody can look at ten thousand scans. So SynthSeg scores its own result. A second small network looks
- **it sees only the label map, not the scan**
  only at the label map, not at the scan.
- **it has learned: good and broken label maps**
  It has learned how good label maps look, and how broken ones look.
- **it guesses the quality: a score from 0 to 1**
  So for each region, it guesses the quality: a score between zero and one.
- **below 0.65 → region fails**
  Below zero point six five, a region fails.
- **aged 54: grey matter 0.722 ✓**
  This scan passes.
- **aged 81: grey matter 0.636 ✗**
  This one fails. Look at it.

SOURCES

- Billot et al. 2023, PNAS (robust SynthSeg, QC)
  SynthSeg+ (robust version) with automatic QC scores per region; below 0.65 a region counts as failed.
- Data: Dallas Lifespan Brain Study, CC0
  Dallas Lifespan Brain Study: public MRI data (CC0), used for our grey-matter pictures.

14. Try it: SynthSeg on a stroke scan



> Interactive page: try it in the online deck – mendeltem.github.io/mri-segmentation-lectures/lecture_1118/deck/#try_qc

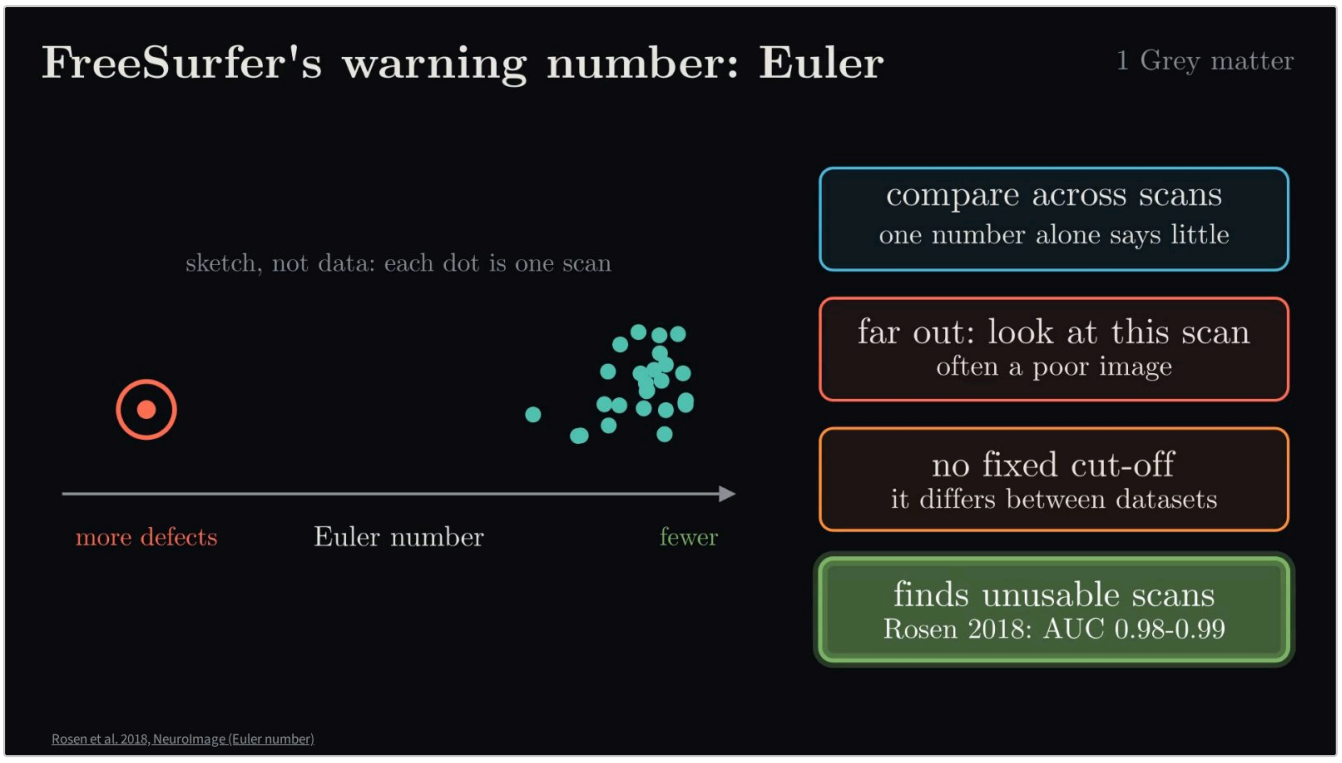
WHAT IS SAID

Scroll through a stroke scan with SynthSeg's labels. All QC values pass, yet the labels around the lesion are not right. A passing QC score is necessary, not sufficient. SynthSeg took 52 seconds on a normal processor.

SOURCES

- Data: Aphasia Recovery Cohort, CC0
  Aphasia Recovery Cohort: public stroke MRI with expert lesion masks (CC0).
- Billot et al. 2023, PNAS (robust SynthSeg, QC)
  SynthSeg+ (robust version) with automatic QC scores per region; below 0.65 a region counts as failed.

15. FreeSurfer's warning number: Euler



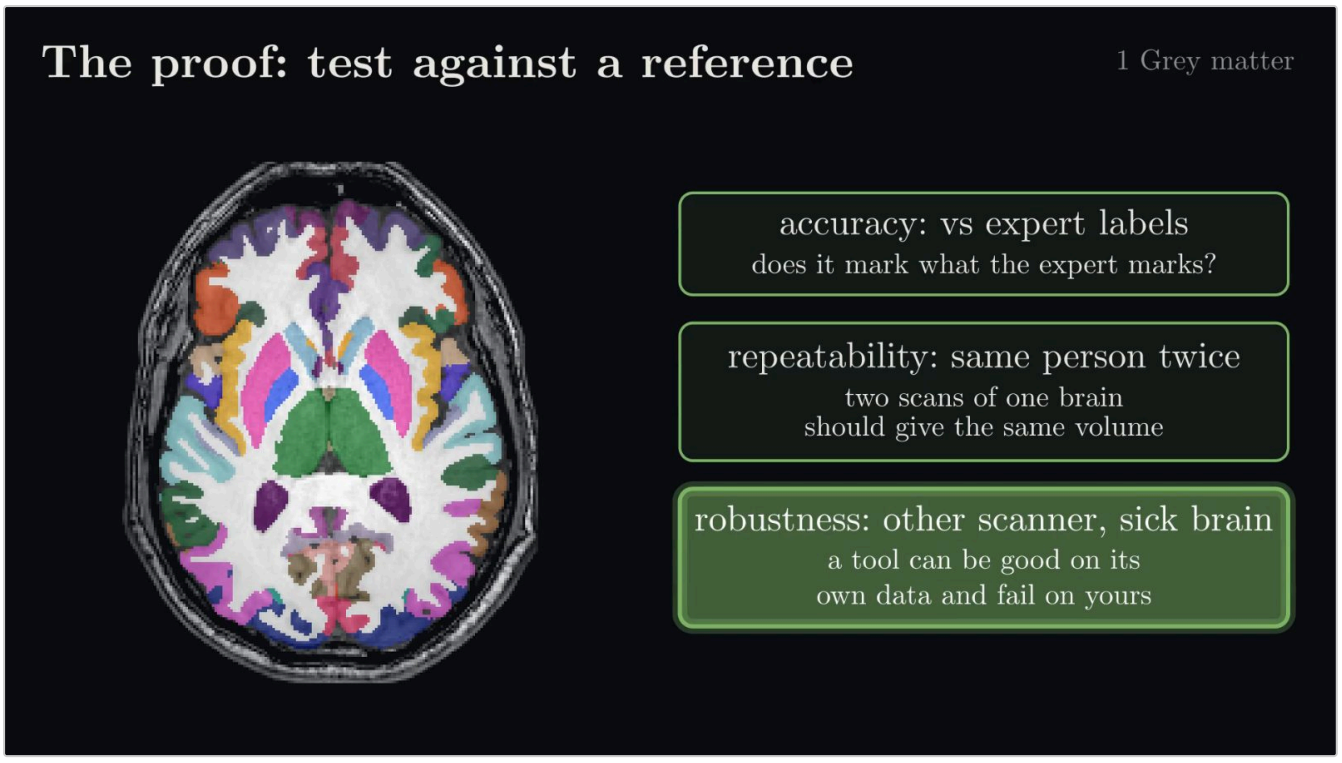
KEY POINTS AND WHAT IS SAID

- **corners − edges + faces** – count them on any surface made of flat pieces
  FreeSurfer has its own warning number: the Oiler number. Take a closed surface, like this cube. Count the corners, subtract the edges, add the faces. Eight minus twelve plus six is two.
- **closed, no holes: always 2** – a cube, a ball, a perfect cortex surface
  For a closed surface without holes, the answer is always two.
- **each handle: minus 2** – a handle is like the handle of a cup
  Every handle lowers the number by two. FreeSurfer's first cortex surface is never perfect.
- **first surface: 26 defects** – public stroke scan, left side
  It has faults, called defects.
- **Euler number −38** – after the repair: 2
  In this public stroke scan, the left side had twenty-six defects, and an Oiler number of minus thirty-eight. FreeSurfer repairs them, and the number is two again.
- **compare across scans** – one number alone says little
  One number alone says little. Compare it across your scans.
- **far out: look at this scan** – often a poor image
  A scan with far more defects than the others is often a poor image: look at it.
- **no fixed cut-off** – it differs between datasets
  There is no fixed cut-off.
- **finds unusable scans** – Rosen 2018: AUC 0.98-0.99
  In a large study, this number found the unusable scans very well.

SOURCES

- Rosen et al. 2018, NeuroImage (Euler number)
  The Euler number of FreeSurfer's surface flags unusable scans very well (AUC 0.98-0.99).

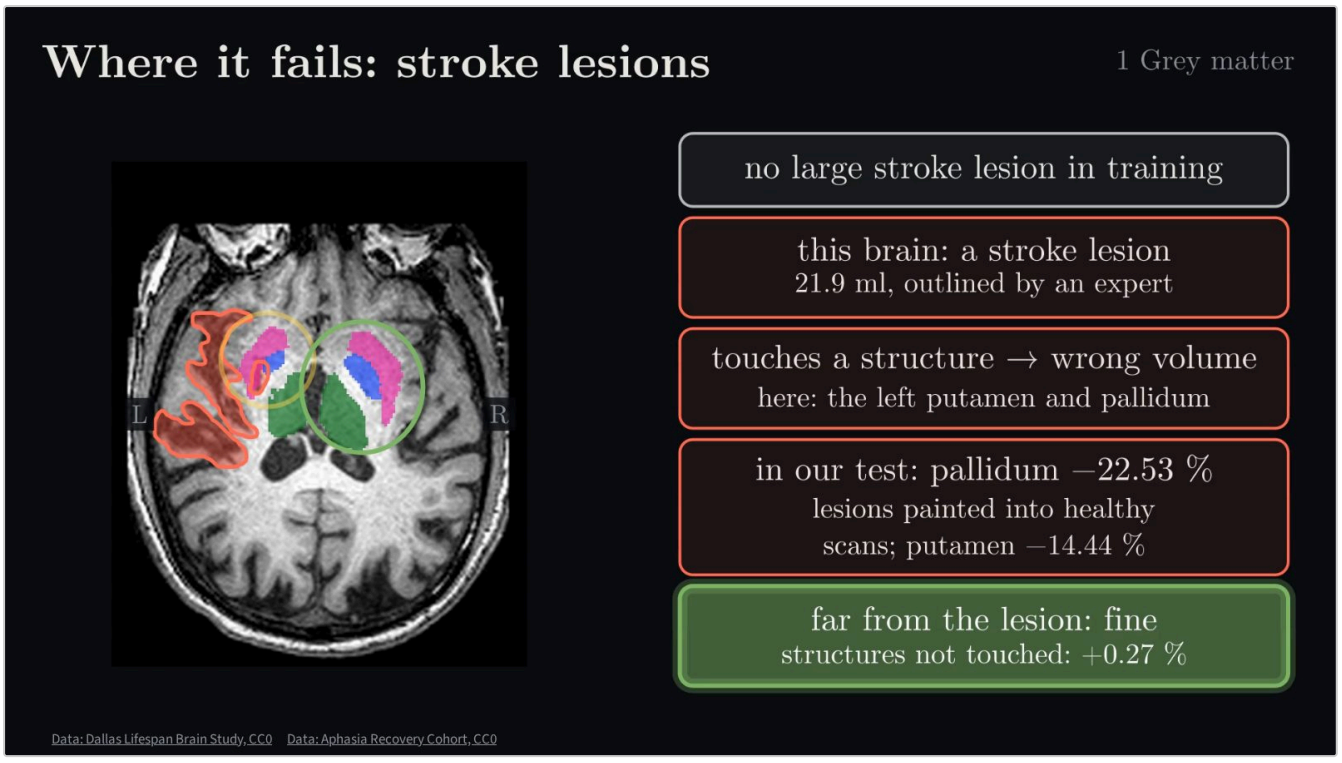
16. The proof: test against a reference



KEY POINTS AND WHAT IS SAID

- **accuracy: vs expert labels** – does it mark what the expert marks?
  The strongest proof: compare with labels drawn by an expert. That is accuracy.
- **repeatability: same person twice** – two scans of one brain should give the same volume
  Same person scanned twice: same result? That is repeatability. And does it still work on another scanner, or on a sick brain?
- **robustness: other scanner, sick brain** – a tool can be good on its own data and fail on yours
  Robustness.

17. Where it fails: stroke lesions



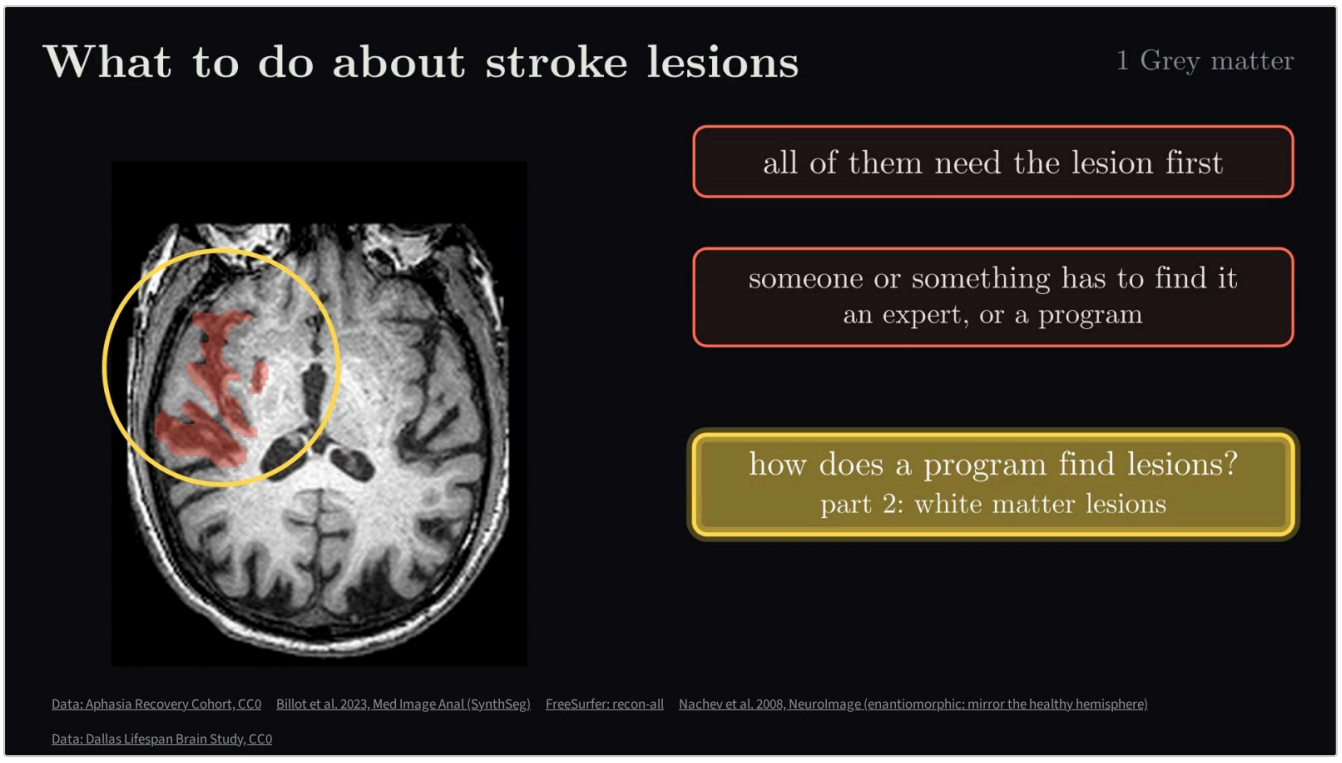
KEY POINTS AND WHAT IS SAID

- **no large stroke lesion in training**
  None of these tools saw a large stroke lesion in training. For them, it is something new.
- **this brain: a stroke lesion**
  This brain has a stroke lesion.
- **lesion touches a structure → wrong volume**
  Where the lesion touches a structure, the program gets its volume wrong.
- **pallidum −22.53 %**
  In our test, the pallidum came out about a fifth too small.
- **far from the lesion: fine**
  Structures far from the lesion are fine.

SOURCES

- Data: Dallas Lifespan Brain Study, CC0
  Dallas Lifespan Brain Study: public MRI data (CC0), used for our grey-matter pictures.
- Data: Aphasia Recovery Cohort, CC0
  Aphasia Recovery Cohort: public stroke MRI with expert lesion masks (CC0).

18. What to do about stroke lesions



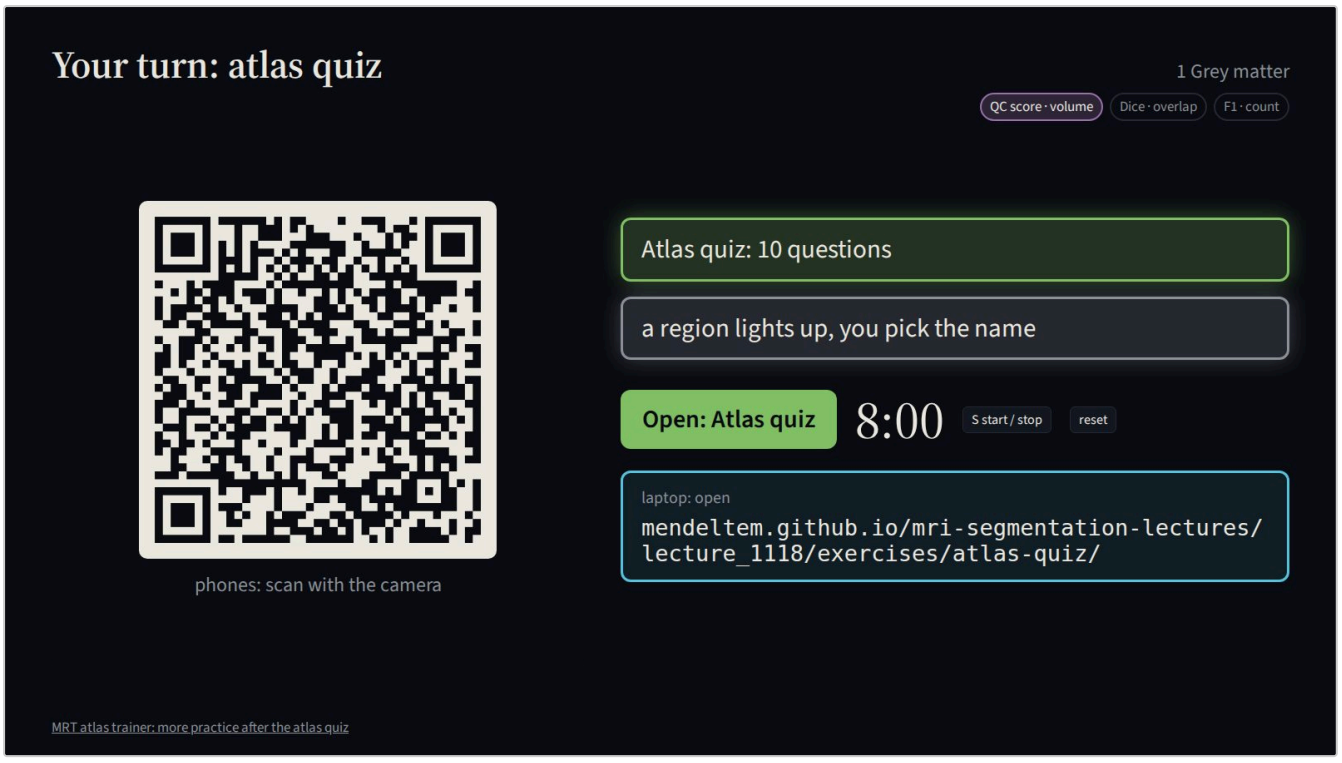
KEY POINTS AND WHAT IS SAID

- **1 fill it from the mirror image** – white-matter values do not fit: a stroke often includes cortex
  There are several ways to deal with these lesions, for SynthSeg and for FreeSurfer. One: fill it. For a stroke, copy the grey values from the mirror image of the healthy hemisphere. White matter values would not fit, because a stroke often includes cortex.
- **2 null it** – lesion voxels set to zero
  Two: null it. The lesion voxels are set to zero.
- **3 leave lesion + margin out** – the most pragmatic way
  Three: leave the lesion and a margin around it out of the count.
- **labels changed: 39 % at 1–2 mm** – 7 % at 8–10 mm, 1 % far away
  In our test, right next to the lesion, two in five labels changed. Ten millimetres away, fewer than one in ten.
- **5 mm margin: 85 % within 5 %** – mirror filling 82 %, nothing done 67 %
  With a five millimetre margin, eighty five percent of the structures on the stroke side stayed within five percent, as good as mirror filling. This is the most pragmatic way.
- **4 very large: healthy hemisphere only**
  Four: if the lesion is very large, use only the healthy hemisphere.
- **all of them need the lesion first** – someone or something has to find it
  All of these ways need the lesion first: someone, or something, has to find it.
- **how does a program find lesions?** – part 2
  How does a program find lesions? That is part two.

SOURCES

- Data: Aphasia Recovery Cohort, CC0
  Aphasia Recovery Cohort: public stroke MRI with expert lesion masks (CC0).
- Billot et al. 2023, Med Image Anal (SynthSeg)
  SynthSeg: a network trained only on synthetic images; segments scans of any contrast and resolution.
- FreeSurfer: recon-all
  FreeSurfer documentation of recon-all, the command that runs the whole pipeline.
- Nachev et al. 2008, NeuroImage (enantiomorphic: mirror the healthy hemisphere)
  Lesions distort the normalisation of brain scans. The enantiomorphic method replaces the lesion with the mirror image of the healthy homologous region; tested on healthy brains with artificial lesions, it beat cost-function masking.
- Data: Dallas Lifespan Brain Study, CC0
  Dallas Lifespan Brain Study: public MRI data (CC0), used for our grey-matter pictures.

19. Your turn: atlas quiz



> **Exercise: Atlas quiz** (8 min) – 10 questions: a region lights up, you pick its name.
> mendeltem.github.io/mri-segmentation-lectures/lecture_1118/exercises/atlas-quiz/

KEY POINTS

- **Atlas quiz: 10 questions**
- **a region lights up, you pick the name**

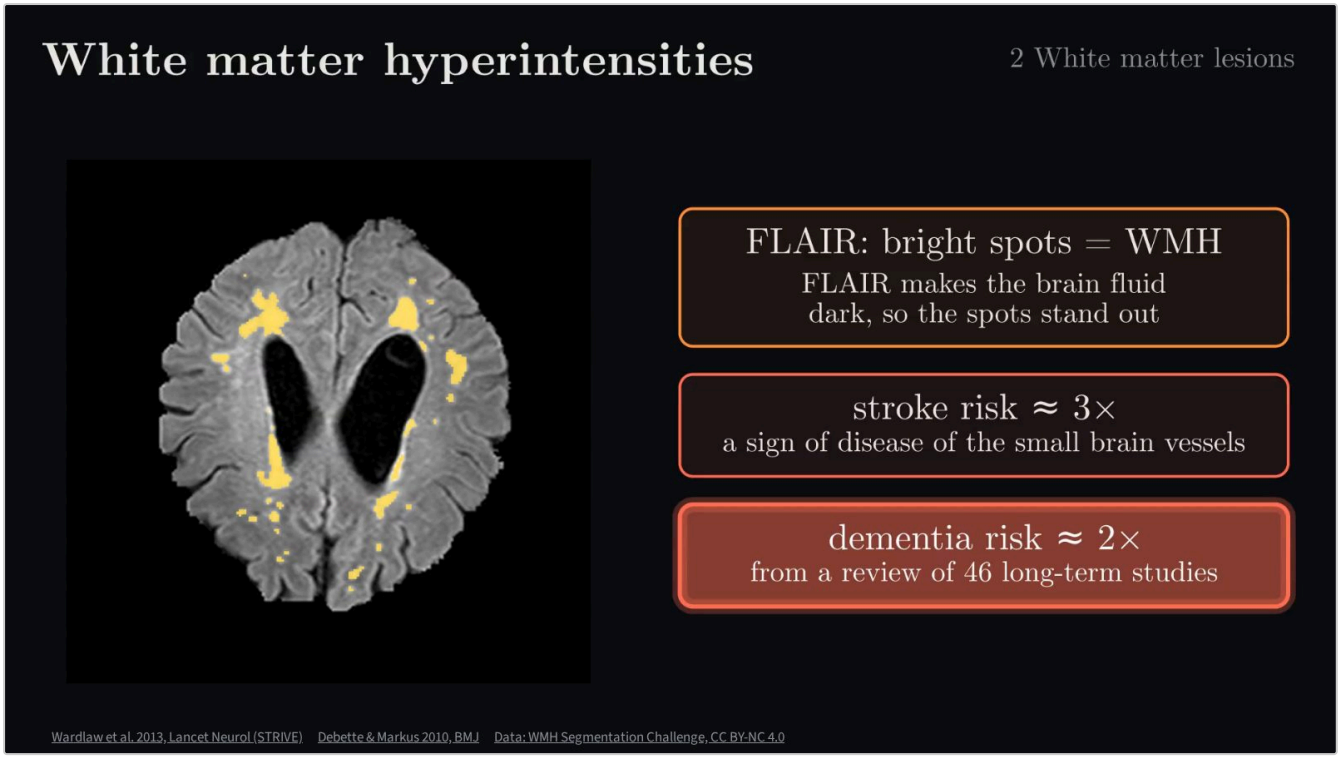
WHAT IS SAID

Now it is your turn. Open the atlas quiz on your phone or laptop. Ten questions: a region lights up, and you pick its name. You have eight minutes.

SOURCES

- MRT atlas trainer: more practice after the atlas quiz
  The full atlas trainer for home: explore, quiz and study modes with many more brain regions (interface in German).

20. White matter hyperintensities



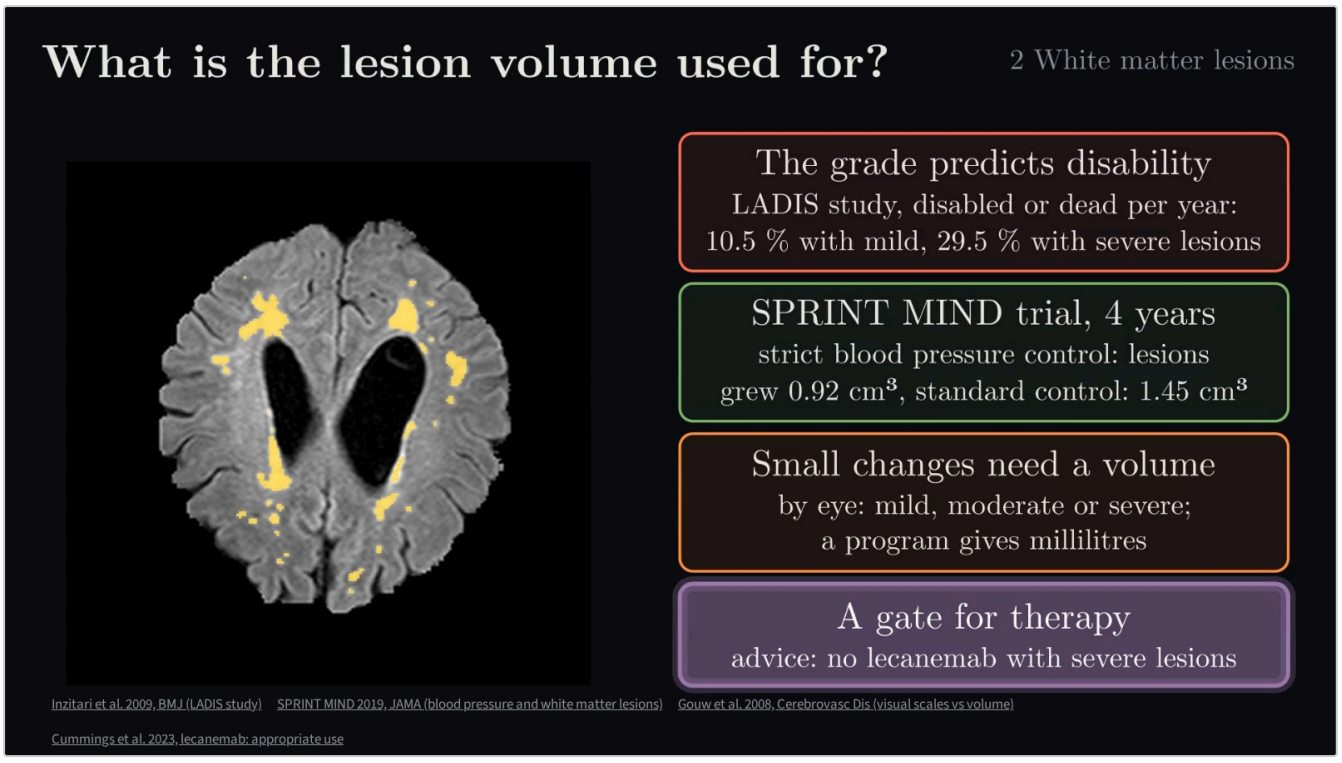
KEY POINTS AND WHAT IS SAID

- **FLAIR: bright spots = WMH** – FLAIR makes the brain fluid dark, so the spots stand out
  Part two: how do we find lesions? FLAIR is a scan that makes the brain fluid dark. So damaged white matter stands out as bright spots: white matter hyperintensities.
- **stroke risk ≈ 3×** – a sign of disease of the small brain vessels
  They are a sign of small vessel disease. People with many of them have a higher risk of stroke
- **dementia risk ≈ 2×** – from a review of 46 long-term studies
  and dementia.

SOURCES

- Wardlaw et al. 2013, Lancet Neurol (STRIVE)
  STRIVE consensus: names, definitions and imaging standards for small vessel disease (WMH, microbleeds).
- Debette & Markus 2010, BMJ
  Meta-analysis: white matter hyperintensities raise the risk of stroke (about 3x) and dementia (about 2x).
- Data: WMH Segmentation Challenge, CC BY-NC 4.0
  The public WMH Challenge data (DataverseNL), licence CC BY-NC 4.0; used in the lesion exercise.

21. What is the lesion volume used for?



KEY POINTS AND WHAT IS SAID

- **The grade predicts disability** – LADIS study, disabled or dead per year: 10.5 % with mild, 29.5 % with severe lesions

  What do we do with this? In one study of older people, doctors graded the lesions by eye: mild, moderate or severe. With severe lesions, almost three times as many people lost their independence or died each year.

- **SPRINT MIND trial, 4 years** – strict blood pressure control: lesions grew 0.92 cm³, standard control: 1.45 cm³

  Trials use the volume as an outcome: strict blood pressure control slowed the growth of the lesions.

- **Small changes need a volume** – by eye: mild, moderate or severe; a program gives millilitres

  But the difference was small: about half a millilitre in four years. No grade by eye can show that. You need a measured volume.

- **A gate for therapy** – advice: no lecanemab with severe lesions

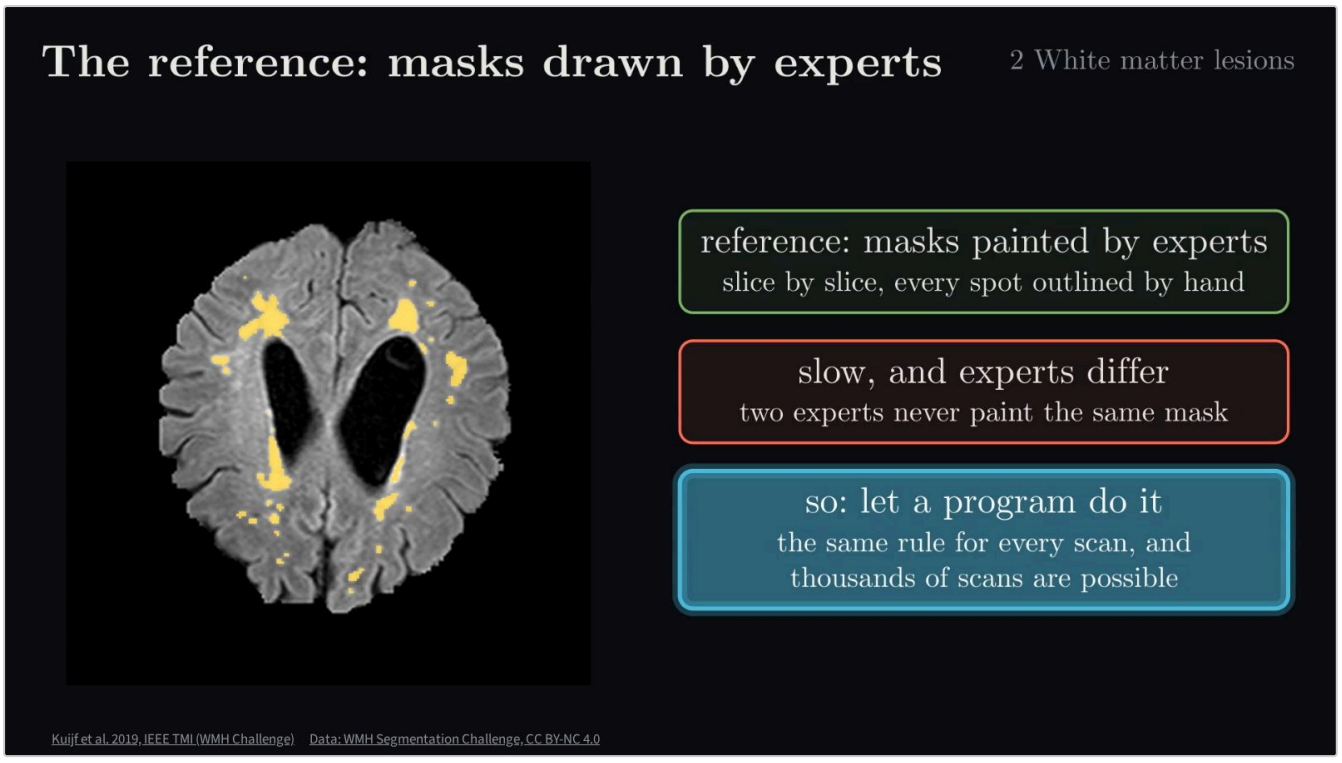
  And for the new Alzheimer antibodies, severe lesions speak against the treatment.

SOURCES

- Inzitari et al. 2009, BMJ (LADIS study)
  LADIS: older people living independently, 3 years follow-up; disability or death per year 10.5 % (mild), 15.1 % (moderate), 29.5 % (severe white matter changes).
- SPRINT MIND 2019, JAMA (blood pressure and white matter lesions)
  SPRINT MIND: intensive vs standard blood pressure control; lesion volume 4.57 to 5.49 cm3 vs 4.40 to 5.85 cm3 after about 4 years (difference -0.54 cm3, small).

- Gouw et al. 2008, Cerebrovasc Dis (visual scales vs volume)
  Measuring the change of lesion volume was reliable (ICC 0.88); judging progression with visual scales was less reliable (20 scan pairs).
- Cummings et al. 2023, lecanemab: appropriate use
  Appropriate use recommendations for lecanemab: no treatment with more than 4 microhaemorrhages.

22. The reference: masks drawn by experts



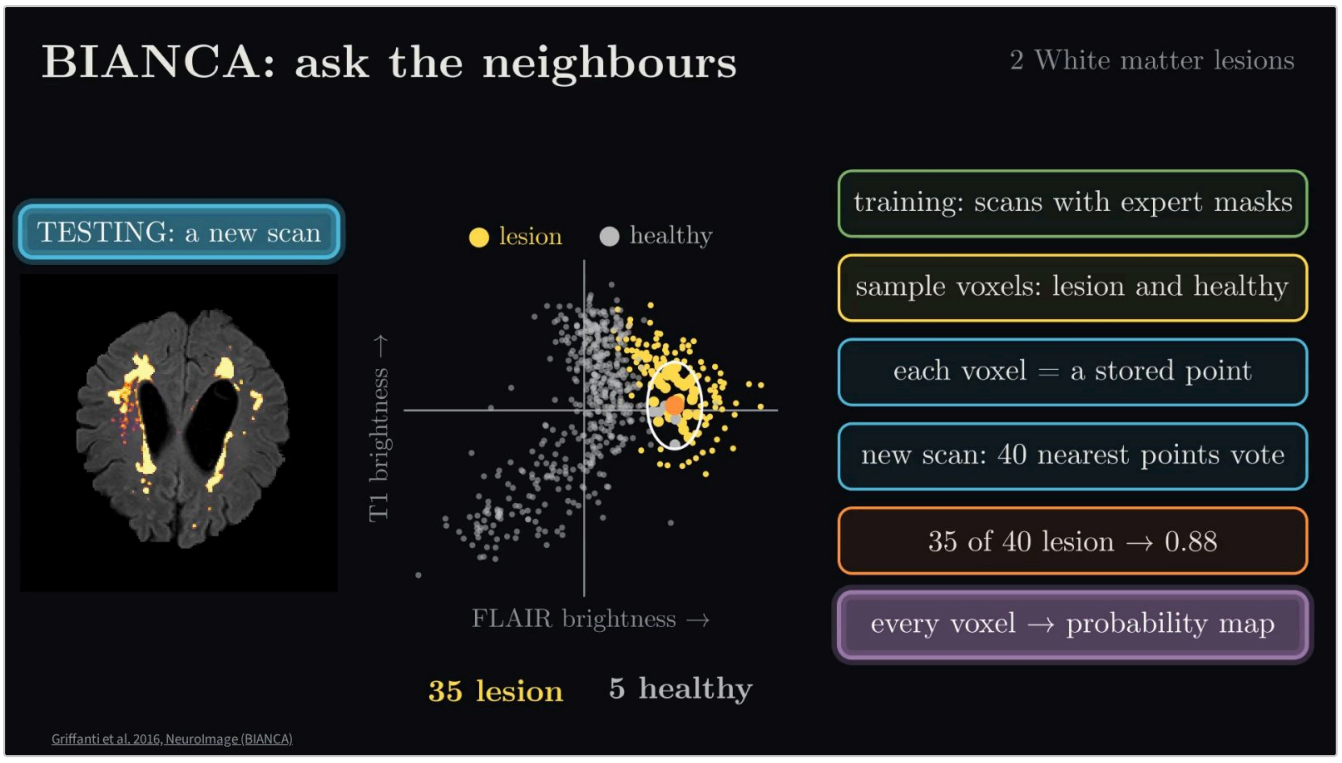
KEY POINTS AND WHAT IS SAID

- **reference: masks painted by experts** – slice by slice, every spot outlined by hand
  The reference is a mask painted by an expert, by hand, slice by slice.
- **slow, and experts differ** – two experts never paint the same mask
  That is slow, and two experts never paint exactly the same.
- **so: let a program do it** – the same rule for every scan, and thousands of scans are possible
  So we want a program to do it.

SOURCES

- Kuijf et al. 2019, IEEE TMI (WMH Challenge)
  WMH Segmentation Challenge: 60 training and 110 test scans, 5 scanners; best method Dice 0.80, second experts 0.77-0.79.
- Data: WMH Segmentation Challenge, CC BY-NC 4.0
  The public WMH Challenge data (DataverseNL), licence CC BY-NC 4.0; used in the lesion exercise.

23. BIANCA: ask the neighbours



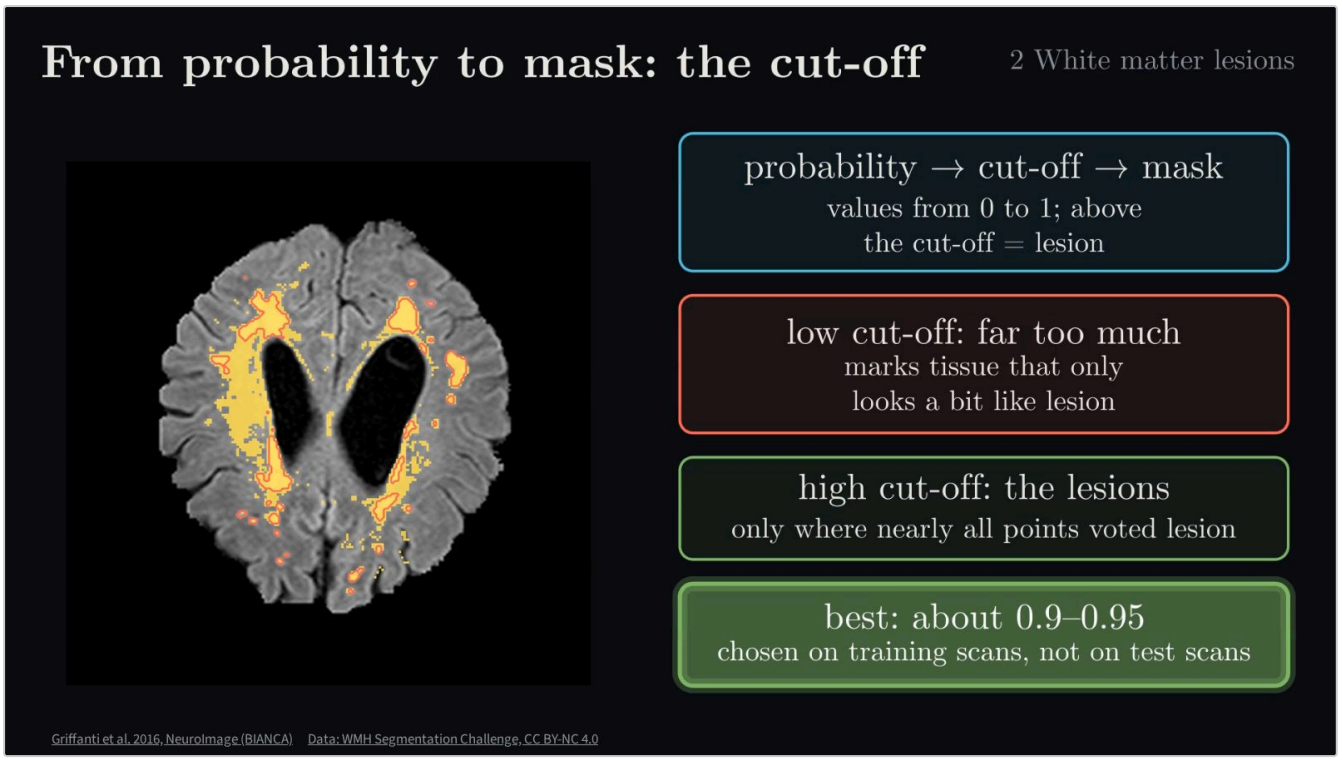
KEY POINTS AND WHAT IS SAID

- **training: scans with expert masks**
  BIANCA learns from examples. In training scans, an expert has marked the lesions.
- **sample voxels: lesion and healthy**
  BIANCA takes sample voxels, lesion and healthy,
- **each voxel = a point**
  and stores each one as a point: its brightness and its position.
- **new scan: 40 nearest points vote**
  Now a new scan, without a mask. For one voxel, the forty nearest stored points vote.
- **35 of 40 lesion → 0.88**
  Thirty five say lesion: probability zero point eight eight.
- **every voxel → probability map**
  BIANCA does this for every voxel.

SOURCES

- Griffanti et al. 2016, NeuroImage (BIANCA)
  BIANCA: k-nearest-neighbour WMH segmentation in FSL; Dice 0.76 in one cohort, 0.52 in another.

24. From probability to mask: the cut-off



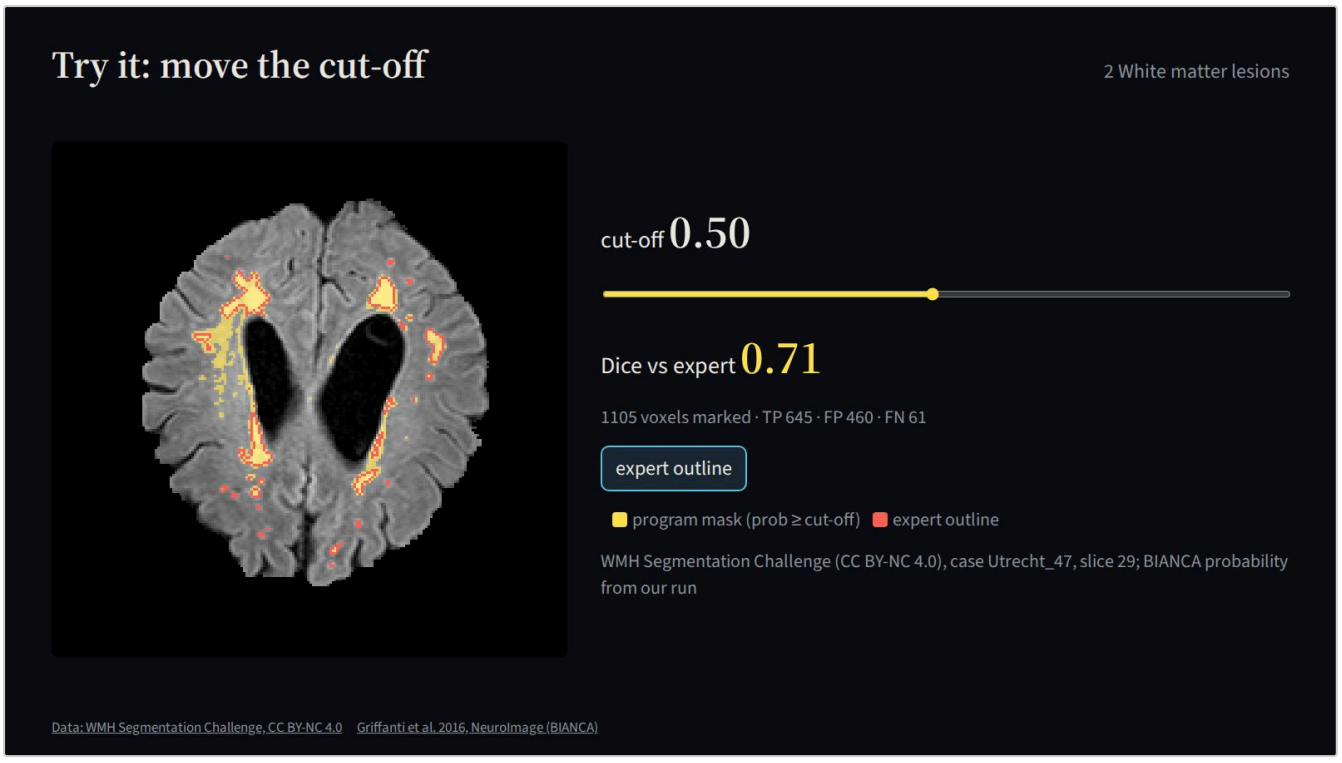
KEY POINTS AND WHAT IS SAID

- **probability → cut-off → mask** – values from 0 to 1; above the cut-off = lesion
  This gives a probability map: every voxel has a value between zero and one. A cut off turns it into a mask: above the cut off counts as lesion.
- **low cut-off: far too much** – marks tissue that only looks a bit like lesion
  A low cut off also marks tissue that only looks a little like a lesion.
- **high cut-off: the lesions** – only where nearly all points voted lesion
  A high one keeps the clear lesions.
- **best: about 0.9–0.95** – chosen on training scans, not on test scans
  About zero point nine to zero point nine five was best, chosen on the training scans.

SOURCES

- Griffanti et al. 2016, NeuroImage (BIANCA)
  BIANCA: k-nearest-neighbour WMH segmentation in FSL; Dice 0.76 in one cohort, 0.52 in another.
- Data: WMH Segmentation Challenge, CC BY-NC 4.0
  The public WMH Challenge data (DataverseNL), licence CC BY-NC 4.0; used in the lesion exercise.

25. Try it: move the cut-off



> Interactive page: try it in the online deck – mendeltem.github.io/mri-segmentation-lectures/lecture_1118/deck/#try_cutoff

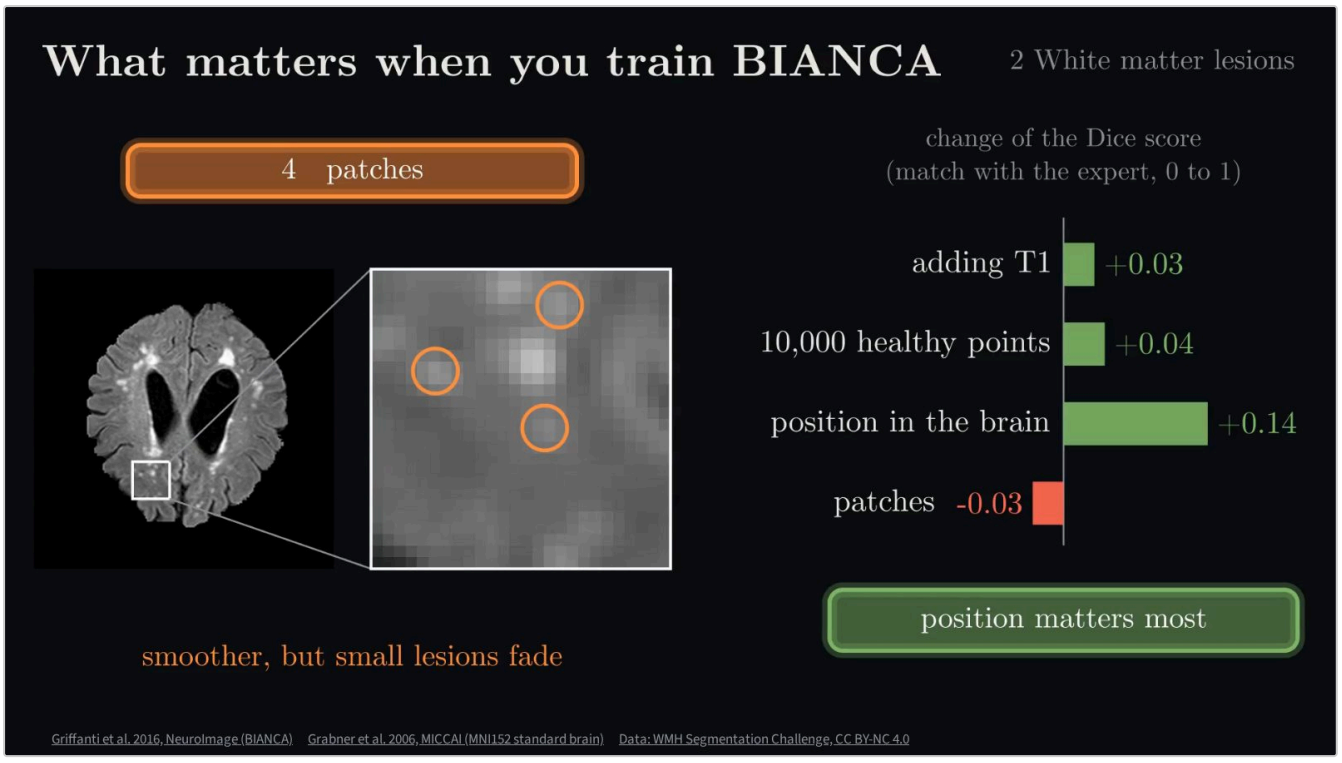
WHAT IS SAID

Move the slider. Low cut-off: the mask floods the white matter and Dice drops. High cut-off: only the lesions stay. Find the best Dice on this slice; the cut-off chosen on the training scans was about 0.9 to 0.95.

SOURCES

- Data: WMH Segmentation Challenge, CC BY-NC 4.0
  The public WMH Challenge data (DataverseNL), licence CC BY-NC 4.0; used in the lesion exercise.
- Griffanti et al. 2016, NeuroImage (BIANCA)
  BIANCA: k-nearest-neighbour WMH segmentation in FSL; Dice 0.76 in one cohort, 0.52 in another.

26. What matters when you train BIANCA



KEY POINTS AND WHAT IS SAID

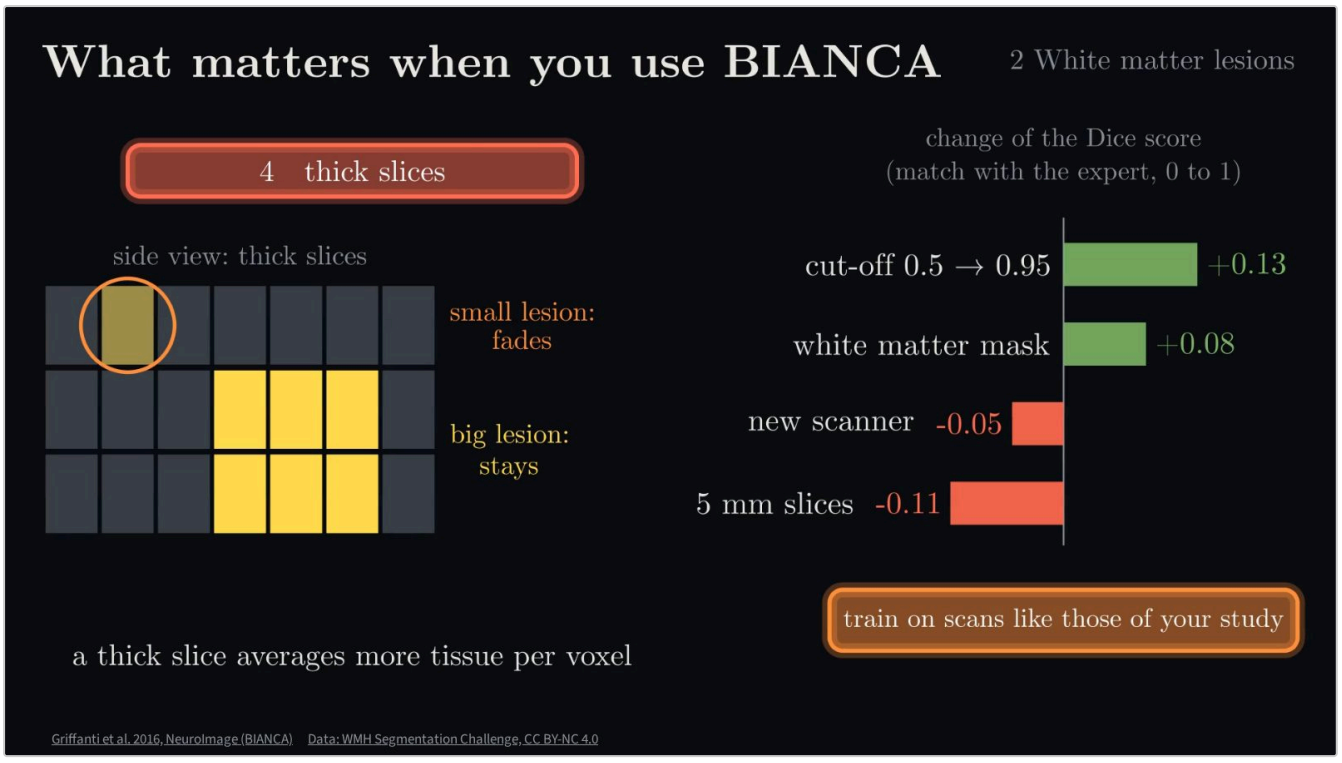
- **adding T1: +0.03**
  BIANCA has settings. We tried them, and each time measured how well the result matches the expert. First: which images? With FLAIR only, each voxel is one number. Adding the T one gives it a second number. That helped a little.
- **10,000 healthy points: +0.04**
  Second: how many healthy examples? Healthy tissue looks different from place to place. With ten thousand examples per scan instead of two thousand, the result got a little better.
- **position in the brain: +0.14**
  Third: position. Lesions have typical places. This map adds up the lesions of all training scans. Most lie around the ventricles. When BIANCA also knows where a voxel lies, the score rises the most.
- **patches: −0.03**
  Last: patches. BIANCA can also look at the three by three neighbourhood of a voxel, and use its average. That smooths noise, but small lesions fade. In our runs it made the result slightly worse.

SOURCES

- Griffanti et al. 2016, NeuroImage (BIANCA)
  BIANCA: k-nearest-neighbour WMH segmentation in FSL; Dice 0.76 in one cohort, 0.52 in another.
- Grabner et al. 2006, MICCAI (MNI152 standard brain)
  The MNI152 standard brain (average of 152 healthy brains): scans are aligned to it, so every voxel gets a position that means the same place in every person.
- Data: WMH Segmentation Challenge, CC BY-NC 4.0

The public WMH Challenge data (DataverseNL), licence CC BY-NC 4.0; used in the lesion exercise.

27. What matters when you use BIANCA



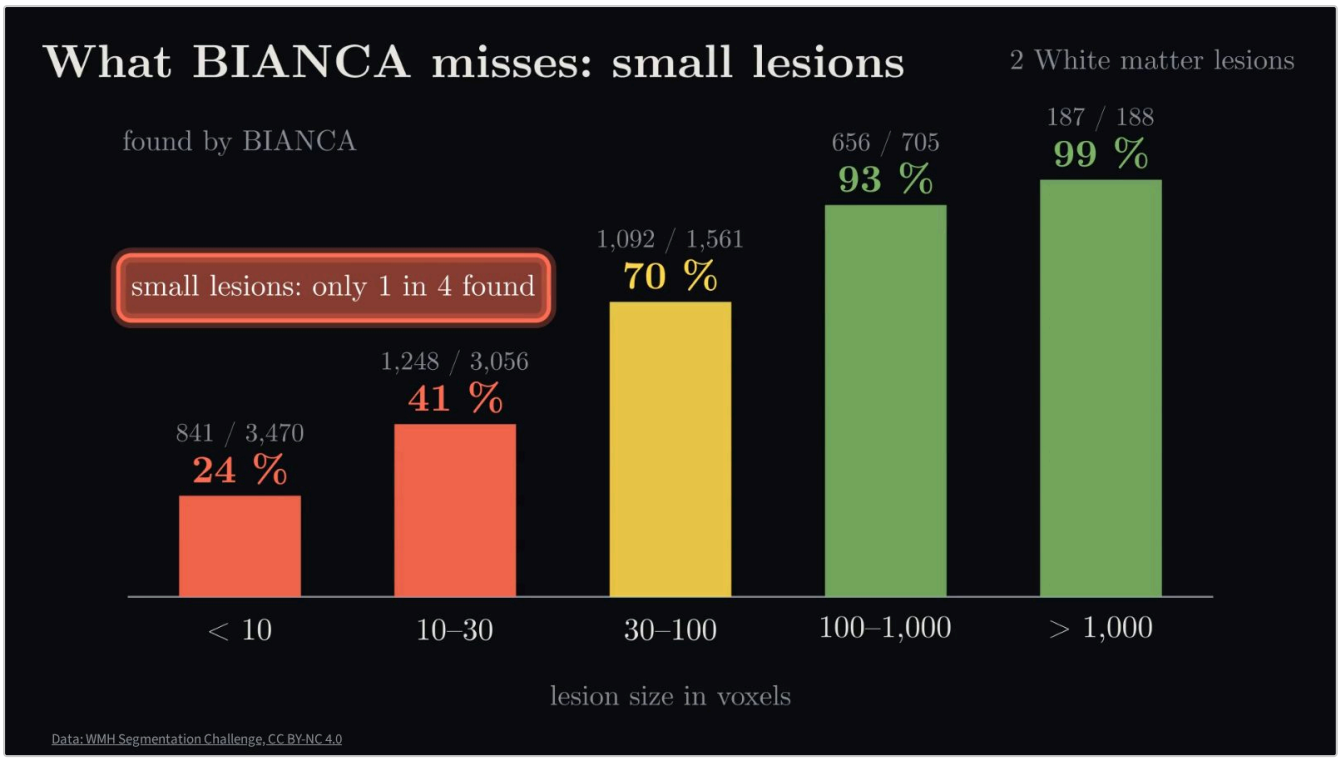
KEY POINTS AND WHAT IS SAID

- **cut-off 0.5 → 0.95: +0.13**
  After training, BIANCA gives a probability for every voxel. The cut off decides what counts as lesion. At zero point five, too much is marked. A high cut off, about zero point nine to zero point nine five, was much better.
- **white matter mask: +0.08**
  Some marks are still wrong: they lie outside the white matter, for example in the cortex. A white matter mask removes them. The true lesions stay.
- **new scanner: −0.05**
  Now your data. BIANCA learned on scans from certain scanners. A new scanner gives another brightness and another contrast. The stored examples fit less well, and the score drops.
- **5 mm slices: −0.11**
  And slice thickness. A thick slice averages more tissue in every voxel. Big lesions stay. Small lesions fade away. That cost the most. So train BIANCA on scans like the ones in your study.

SOURCES

- Griffanti et al. 2016, NeuroImage (BIANCA)
  BIANCA: k-nearest-neighbour WMH segmentation in FSL; Dice 0.76 in one cohort, 0.52 in another.
- Data: WMH Segmentation Challenge, CC BY-NC 4.0
  The public WMH Challenge data (DataverseNL), licence CC BY-NC 4.0; used in the lesion exercise.

28. What BIANCA misses: small lesions



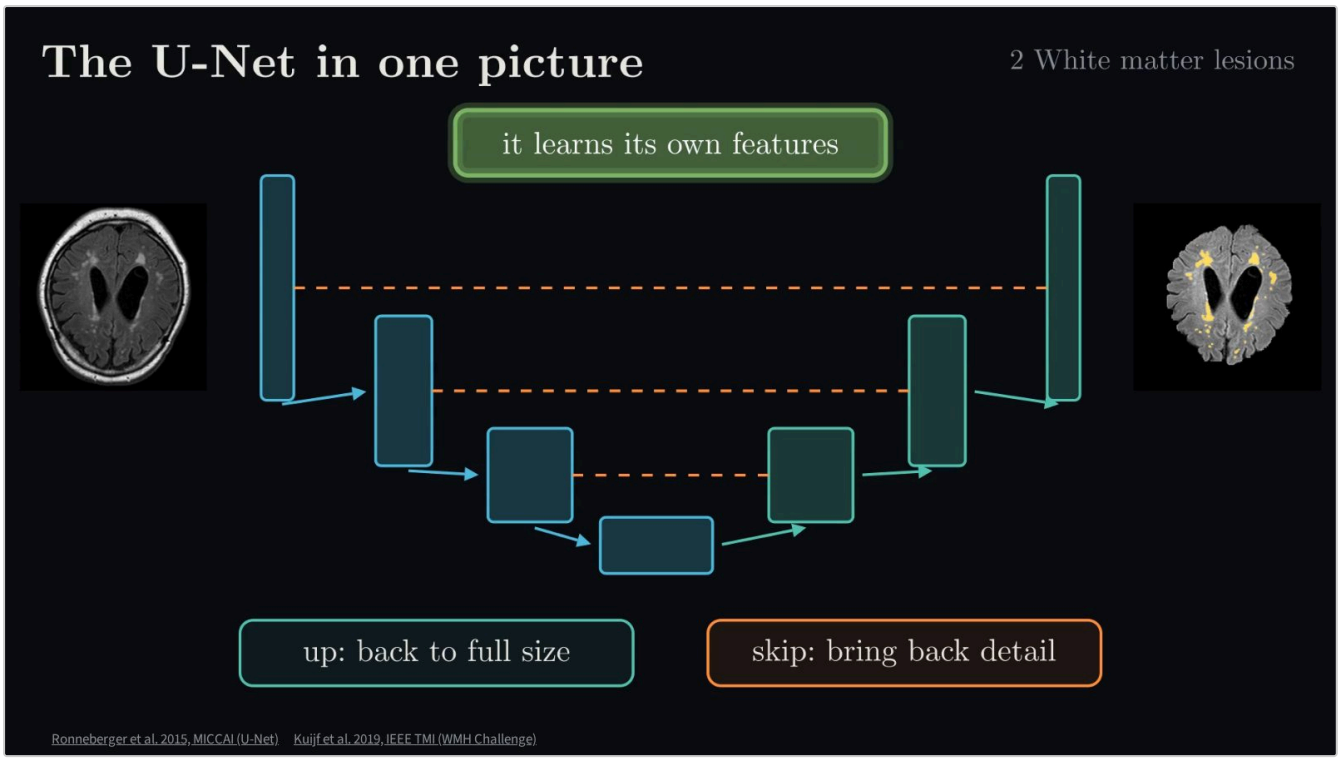
KEY POINTS AND WHAT IS SAID

- **over 1,000 voxels: 99 % found**
  Which lesions does BIANCA find? The big ones, almost always.
- **under 10 voxels: 24 % found**
  But of the smallest lesions, under ten voxels, only one in four. That is why small lesions pull Dice down.

SOURCES

- Data: WMH Segmentation Challenge, CC BY-NC 4.0
  The public WMH Challenge data (DataverseNL), licence CC BY-NC 4.0; used in the lesion exercise.

29. The U-Net in one picture



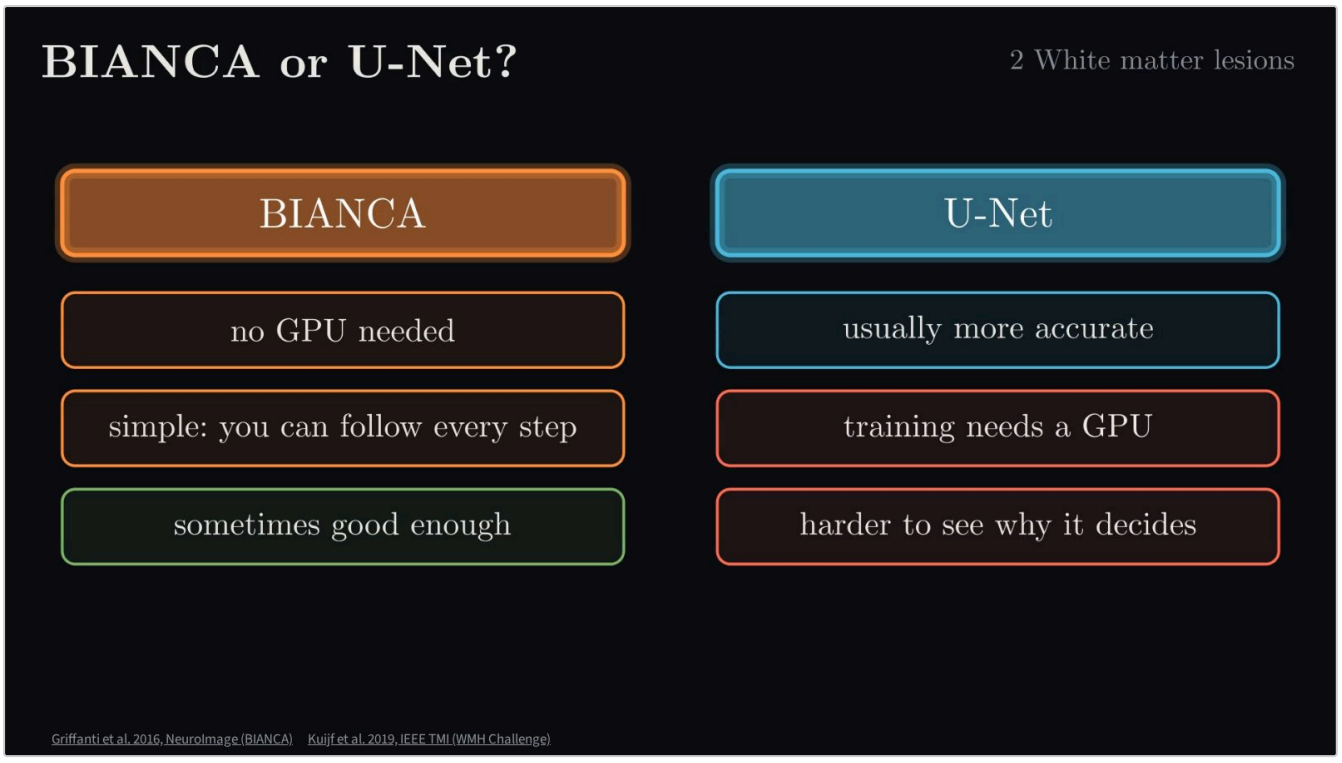
KEY POINTS AND WHAT IS SAID

- **down: shrink, collect context**
  A U-Net goes down, shrinking the image to see context.
- **up: back to full size**
  Then it goes up again,
- **skip connections: fine detail**
  and skip connections bring back the detail.
- **learns its own features**
  It learns its own features. Today, the best lesion tools work like this.

SOURCES

- Ronneberger et al. 2015, MICCAI (U-Net)
  U-Net: the encoder-decoder network with skip connections for biomedical image segmentation.
- Kuijf et al. 2019, IEEE TMI (WMH Challenge)
  WMH Segmentation Challenge: 60 training and 110 test scans, 5 scanners; best method Dice 0.80, second experts 0.77-0.79.

30. BIANCA or U-Net?



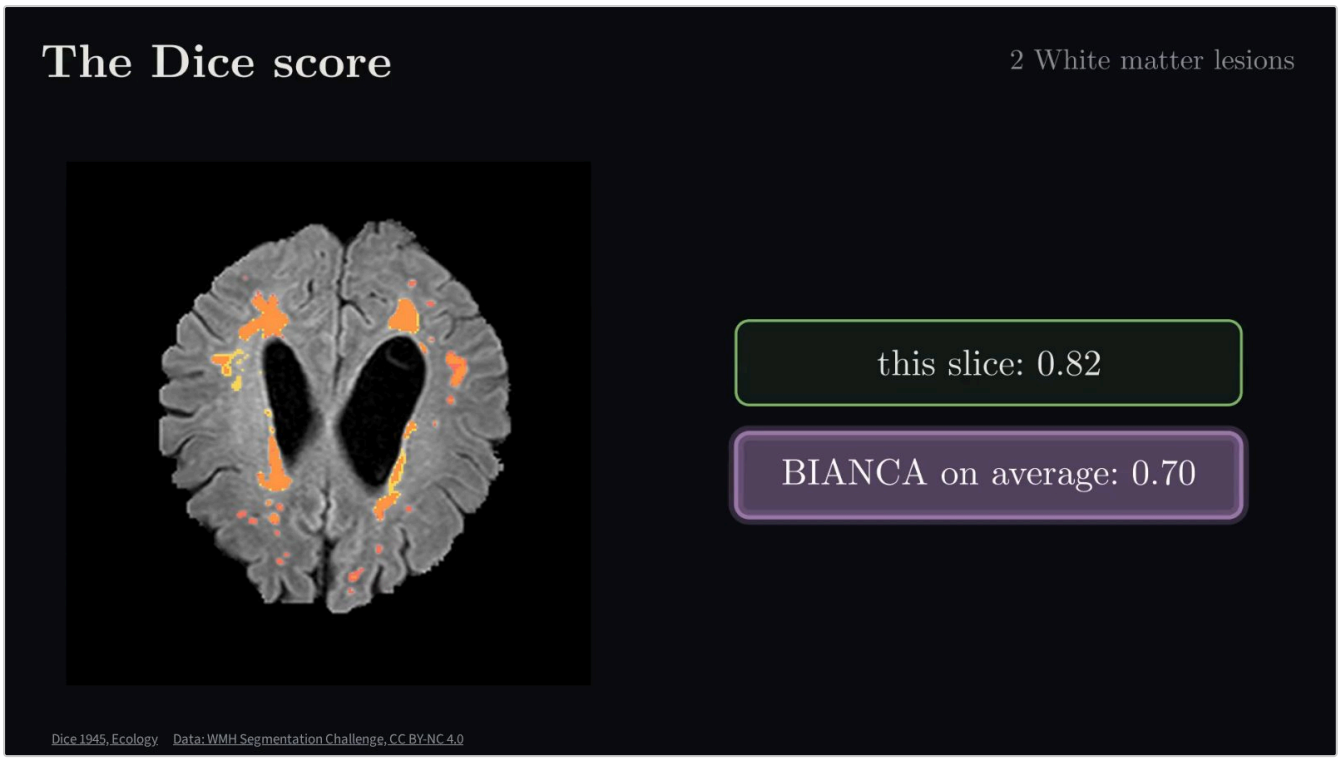
KEY POINTS AND WHAT IS SAID

- **BIANCA: no GPU needed**
  So, BIANCA or U-Net? BIANCA needs no graphics card.
- **simple: you can follow every step**
  It is simple: you can follow every step.
- **sometimes good enough**
  And sometimes, it is good enough.
- **U-Net: usually more accurate**
  A U-Net is usually more accurate.
- **training needs a GPU**
  But training it needs a graphics card,
- **harder to see why it decides**
  and it is harder to see why it decides.

SOURCES

- Griffanti et al. 2016, NeuroImage (BIANCA)
  BIANCA: k-nearest-neighbour WMH segmentation in FSL; Dice 0.76 in one cohort, 0.52 in another.
- Kuijf et al. 2019, IEEE TMI (WMH Challenge)
  WMH Segmentation Challenge: 60 training and 110 test scans, 5 scanners; best method Dice 0.80, second experts 0.77-0.79.

31. The Dice score



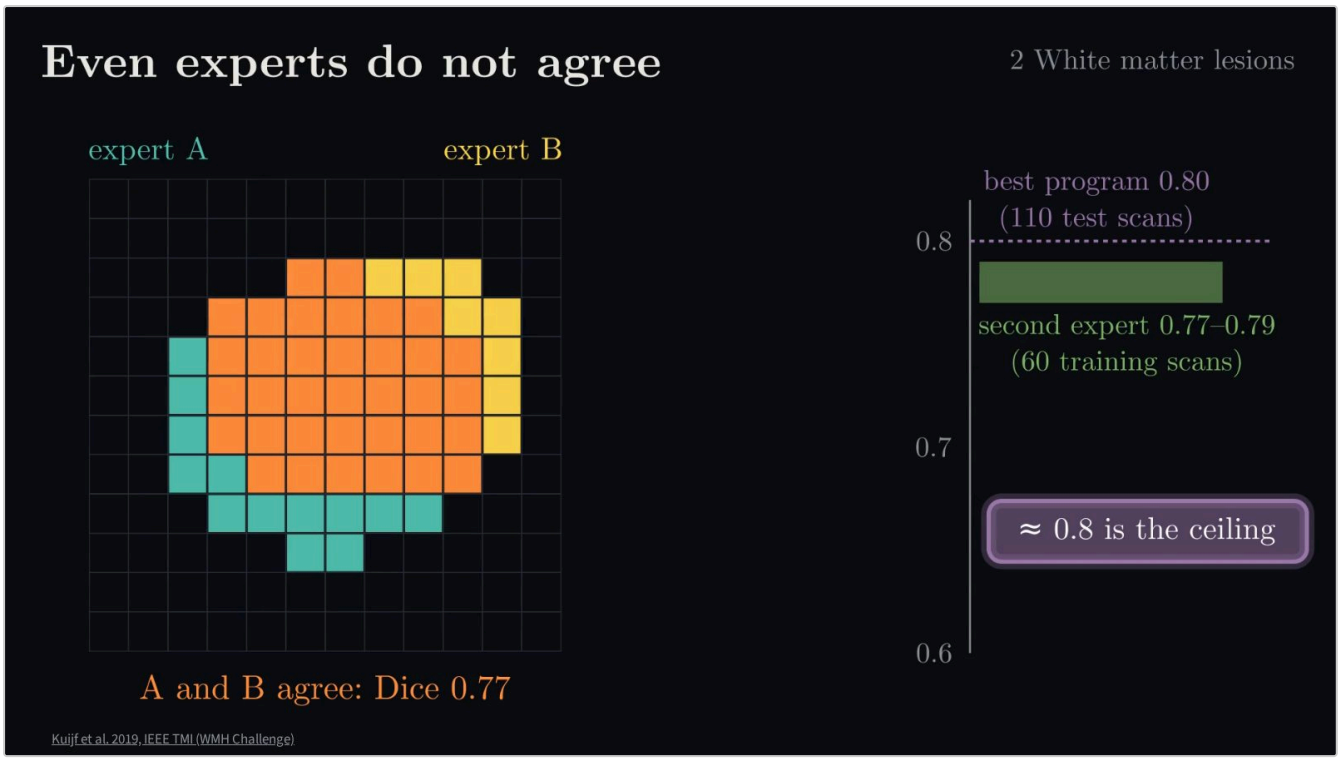
KEY POINTS AND WHAT IS SAID

- **TP: in both**
  Put the mask on the reference: hits in both,
- **FP: only the program**
  extra marks,
- **FN: missed**
  and misses.
- **Dice = 2·TP / (2·TP + FP + FN)**
  Dice is two times the hits, divided by both mask sizes together. One is perfect.
- **this slice: 0.82**
  This slice: zero point eight two.
- **BIANCA on average: 0.70**
  BIANCA on average: zero point seven.

SOURCES

- Dice 1945, Ecology
  Dice's 1945 overlap coefficient, first used to measure how often two species occur together.
- Data: WMH Segmentation Challenge, CC BY-NC 4.0
  The public WMH Challenge data (DataverseNL), licence CC BY-NC 4.0; used in the lesion exercise.

32. Even experts do not agree



KEY POINTS AND WHAT IS SAID

- **second expert vs reference: 0.77–0.79** – on the 60 training scans
  Even experts do not agree. A second expert reached only zero point seven seven to zero point seven nine against the reference, on the training scans.
- **best program (U-Net-like): 0.80** – on the 110 test scans: about the same level
  The best program, a U-Net, reached zero point eight on the test scans: about the same level. So that is the ceiling.

SOURCES

- Kuijf et al. 2019, IEEE TMI (WMH Challenge)
  WMH Segmentation Challenge: 60 training and 110 test scans, 5 scanners; best method Dice 0.80, second experts 0.77-0.79.

33. Dice is hard on small lesions



KEY POINTS AND WHAT IS SAID

- **10 × 10 voxels, shifted by 1**
  Shift a big lesion by one voxel.
- **Dice 0.90** – 90 of 100 voxels still overlap
  Ninety of a hundred voxels still overlap: Dice is zero point nine.
- **2 × 2 voxels, same shift**
  The same shift on a tiny lesion.
- **Dice 0.50** – only 2 of 4 voxels still overlap
  Only two of four voxels still overlap: Dice is zero point five.
- **small lesions are hard on Dice** – the same small error costs a small lesion much more
  Small lesions are hard on Dice.

34. Try it: shift the mask



> Interactive page: try it in the online deck – mendeltem.github.io/mri-segmentation-lectures/lecture_1118/deck/#try_dice

WHAT IS SAID

Shift the program's mask with the buttons. One voxel costs the big lesion little and the small lesion a lot. Two voxels and the small lesion has no overlap at all.

35. Your turn: segment the lesions



> **Exercise: Lesion exercise** (6 min) – 1 patient, 3 slices. Short tutorial, paint, get your Dice.
> mendeltem.github.io/mri-segmentation-lectures/lecture_1118/exercises/wmh-dice/

KEY POINTS

- **Lesion exercise: 1 patient, 3 slices**
- **tutorial → paint → your Dice**

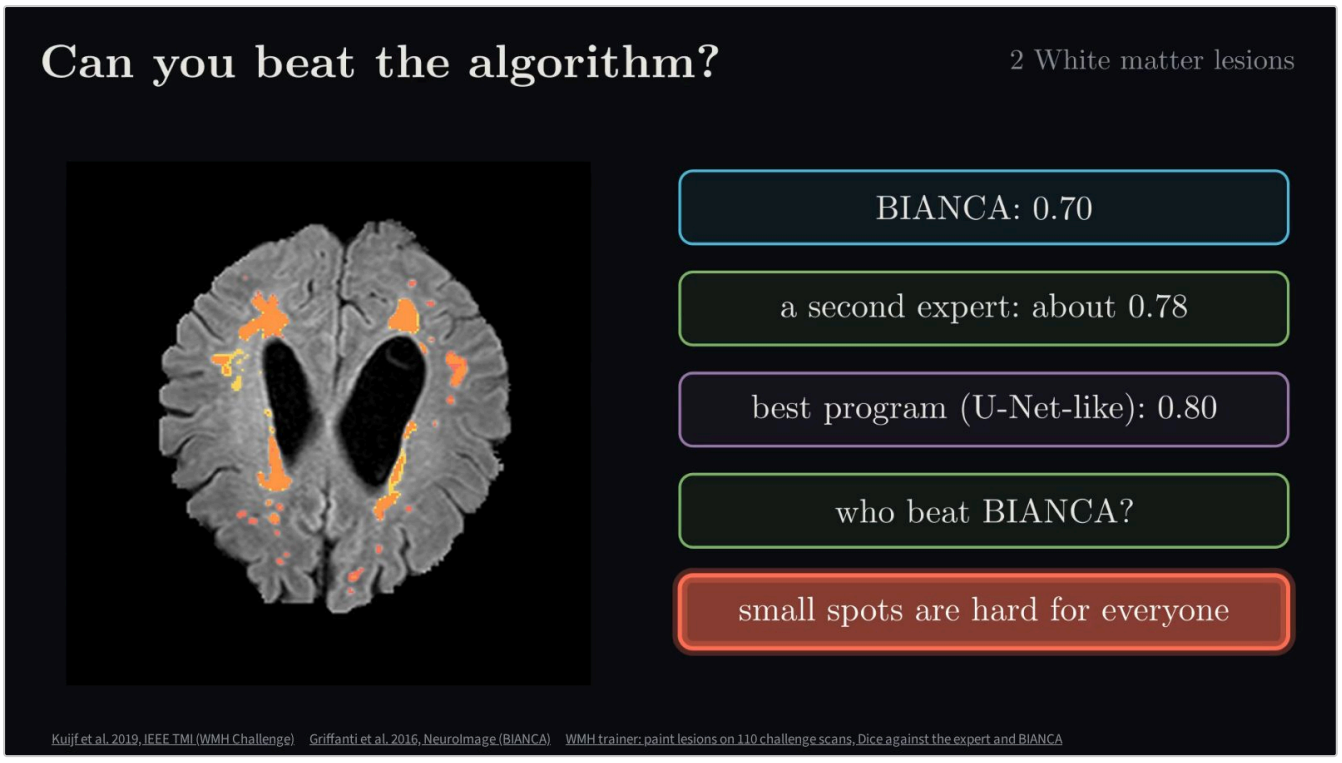
WHAT IS SAID

Your turn. Open the lesion exercise. A short tutorial first, then paint the lesions on three slices of one patient, and get your Dice. You have six minutes.

SOURCES

- Data: WMH Segmentation Challenge, CC BY-NC 4.0
  The public WMH Challenge data (DataverseNL), licence CC BY-NC 4.0; used in the lesion exercise.

36. Can you beat the algorithm?



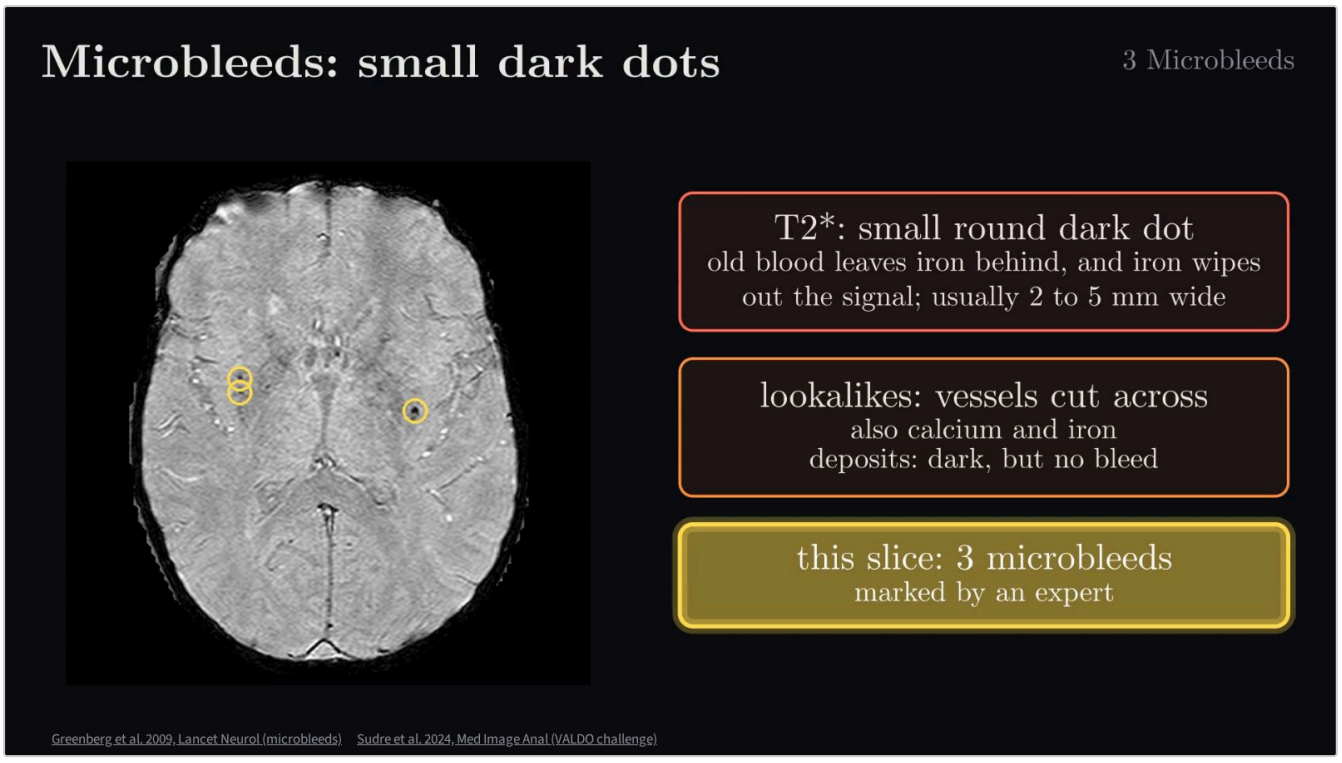
KEY POINTS AND WHAT IS SAID

- **BIANCA: 0.70**
  How did you do? BIANCA got zero point seven,
- **a second expert: about 0.78**
  a second expert about zero point seven eight,
- **best program (U-Net-like): 0.80**
  the best program, a U-Net, zero point eight.
- **who beat BIANCA?**
  Hands up, who beat BIANCA?
- **small spots are hard for everyone**
  The small spots are hard for everyone.

SOURCES

- Kuijf et al. 2019, IEEE TMI (WMH Challenge)
  WMH Segmentation Challenge: 60 training and 110 test scans, 5 scanners; best method Dice 0.80, second experts 0.77-0.79.
- Griffanti et al. 2016, NeuroImage (BIANCA)
  BIANCA: k-nearest-neighbour WMH segmentation in FSL; Dice 0.76 in one cohort, 0.52 in another.
- WMH trainer: paint lesions on 110 challenge scans, Dice against the expert and BIANCA
  Practice at home: paint the WMH on 6 FLAIR slices of each of the 110 public WMH Challenge test scans, then see your Dice against the expert and against BIANCA. Works on phones.

37. Microbleeds: small dark dots



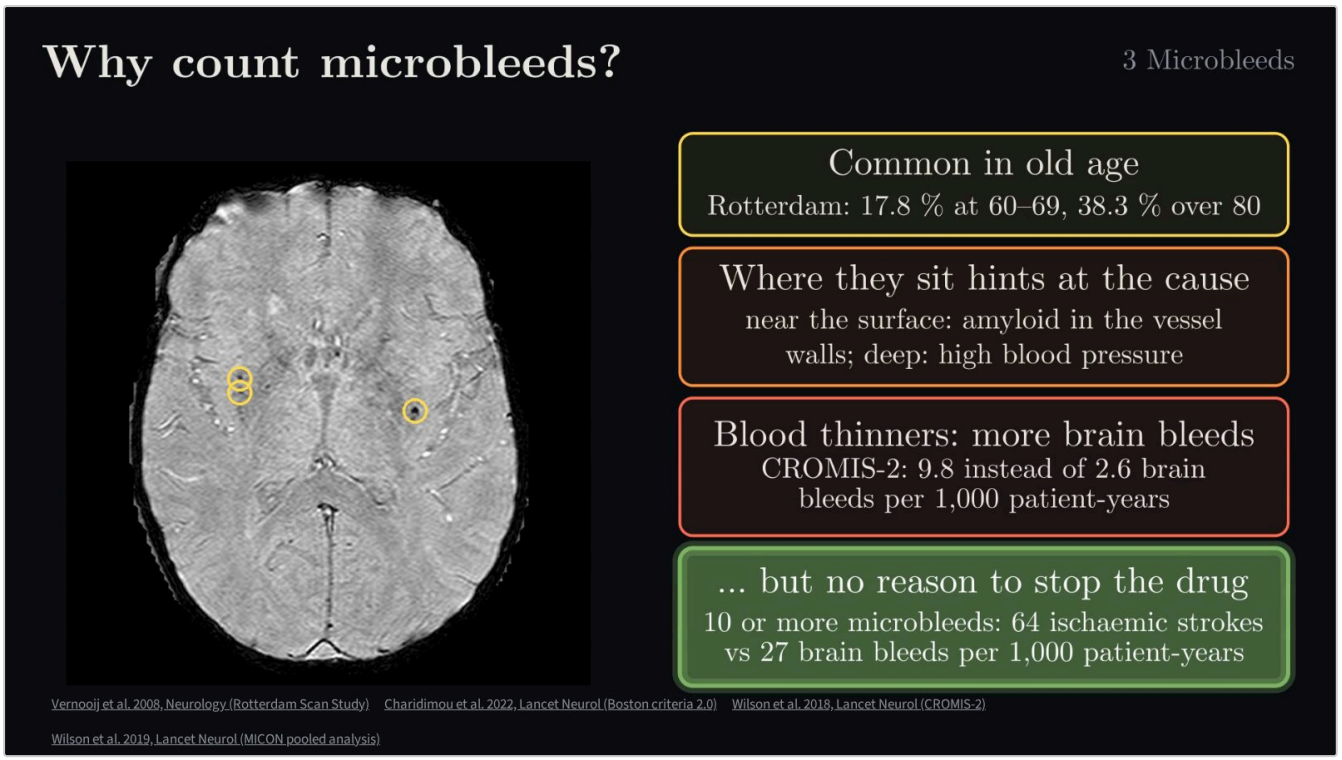
KEY POINTS AND WHAT IS SAID

- **T2*: small round dark dot** – old blood leaves iron behind, and iron wipes out the signal; usually 2 to 5 mm wide
  Part three. A microbleed is a tiny old bleed. It leaves iron behind, and iron disturbs the signal. So on T two star it is a small, round, dark dot, usually two to five millimetres wide.
- **lookalikes: vessels cut across** – also calcium and iron deposits: dark, but no bleed
  Watch for lookalikes, like vessels cut across.
- **this slice: 3 microbleeds** – marked by an expert
  This slice has three.

SOURCES

- Greenberg et al. 2009, Lancet Neurol (microbleeds)
  Review of cerebral microbleeds: definition, detection on T2*/SWI, mimics and clinical meaning.
- Sudre et al. 2024, Med Image Anal (VALDO challenge)
  VALDO challenge (MICCAI 2021): automatic detection of small vessel lesions, including microbleeds.

38. Why count microbleeds?



KEY POINTS AND WHAT IS SAID

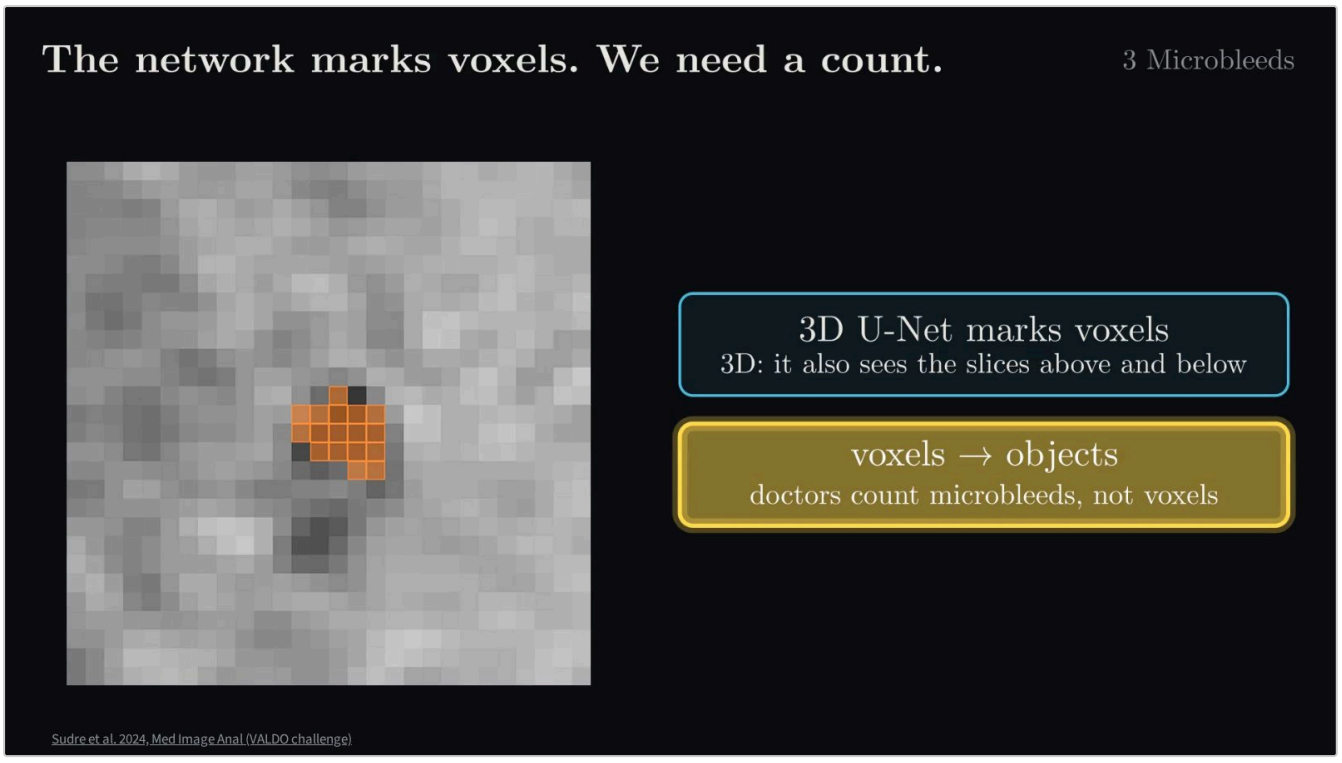
- **Common in old age** – Rotterdam: 17.8 % at 60–69, 38.3 % over 80
  Why count them? First: microbleeds are common. In one large study, every sixth person in their sixties had some, and more than every third over eighty.
- **Where they sit hints at the cause** – near the surface: amyloid in the vessel walls; deep: high blood pressure
  Where they sit hints at the cause. Near the surface of the brain: amyloid in the vessel walls. Deep in the brain: high blood pressure.
- **Blood thinners: more brain bleeds** – CROMIS-2: 9.8 instead of 2.6 brain bleeds per 1,000 patient-years
  And they matter for treatment. On blood thinners, patients with microbleeds had almost four times as many brain bleeds.
- **… but no reason to stop the drug** – 10 or more microbleeds: 64 ischaemic strokes vs 27 brain bleeds per 1,000 patient-years
  Still, the studies say: do not stop the drug for that reason alone. Even with ten or more microbleeds, a new stroke from a blocked vessel stays more likely than a bleed.

SOURCES

- Vernooij et al. 2008, Neurology (Rotterdam Scan Study)
  Rotterdam Scan Study (1,062 people, 1.5 T): microbleeds in 17.8 % at age 60-69 and 38.3 % over 80; lobar location points to amyloid angiopathy, deep location to hypertensive vessel disease.
- Charidimou et al. 2022, Lancet Neurol (Boston criteria 2.0)
  Boston criteria 2.0 for cerebral amyloid angiopathy: strictly lobar haemorrhagic lesions (bleeds, microbleeds, superficial siderosis) make the diagnosis.

- Wilson et al. 2018, Lancet Neurol (CROMIS-2)
  CROMIS-2: anticoagulated patients with atrial fibrillation after stroke or TIA; brain bleeds 9.8 (with microbleeds) vs 2.6 (without) per 1,000 patient-years.
- Wilson et al. 2019, Lancet Neurol (MICON pooled analysis)
  Pooled data of 20,322 patients after ischaemic stroke or TIA: with 10 or more microbleeds still 64 ischaemic strokes vs 27 brain bleeds per 1,000 patient-years.

39. The network marks voxels. We need a count.



KEY POINTS AND WHAT IS SAID

- **3D U-Net marks voxels** – 3D: it also sees the slices above and below
  A three D U-Net looks at the scan as a block, with the slices above and below, and marks every voxel that looks like a bleed.
- **voxels → objects** – doctors count microbleeds, not voxels
  But doctors count bleeds, not voxels. So we group the voxels into objects.

SOURCES

- Sudre et al. 2024, Med Image Anal (VALDO challenge)
  VALDO challenge (MICCAI 2021): automatic detection of small vessel lesions, including microbleeds.

40. Clusters: 26 neighbours



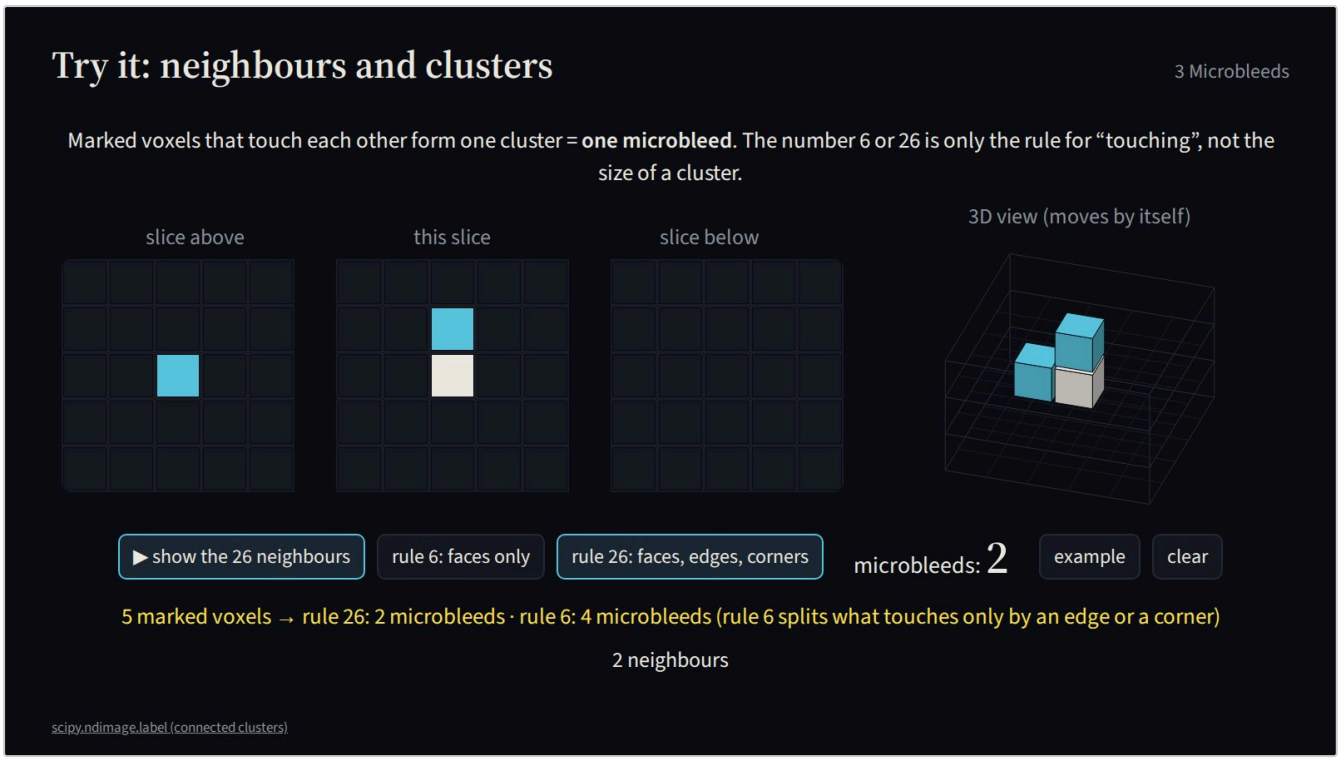
KEY POINTS AND WHAT IS SAID

- **6 share a face**
  Voxels that touch belong together: by a face,
- **12 share an edge**
  an edge,
- **8 share a corner**
  or a corner. Twenty six neighbours.
- **connected voxels = 1 cluster = 1 microbleed**
  One cluster is one microbleed.

SOURCES

- scipy.ndimage.label (connected clusters)
  SciPy function that groups touching voxels into clusters; the neighbourhood (6 or 26) is a choice.

41. Try it: neighbours and clusters



> Interactive page: try it in the online deck – mendeltem.github.io/mri-segmentation-lectures/lecture_1118/deck/#try_26

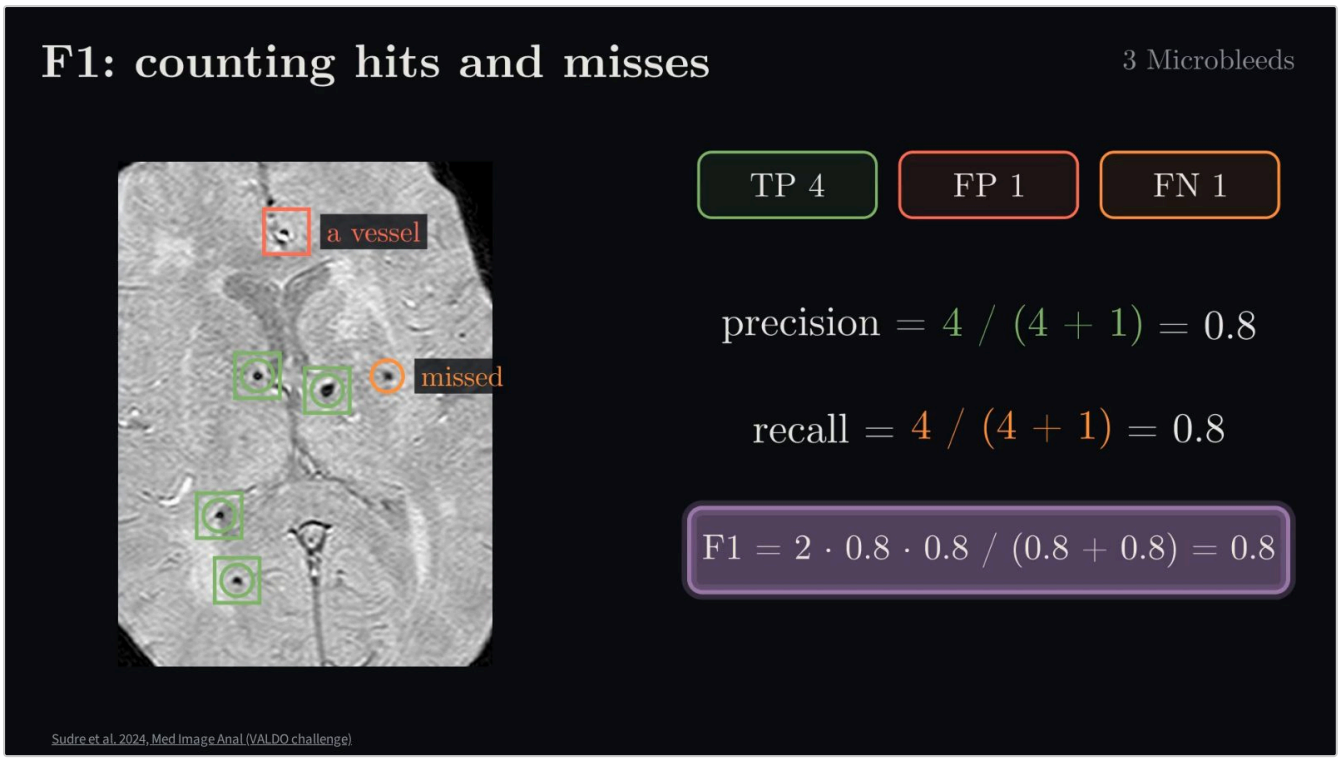
WHAT IS SAID

Click voxels on the three slices. With 26 neighbours, voxels that touch only at an edge or a corner join one cluster; with 6 they stay separate. The count of microbleeds changes with the rule.

SOURCES

- scipy.ndimage.label (connected clusters)
  SciPy function that groups touching voxels into clusters; the neighbourhood (6 or 26) is a choice.

42. F1: counting hits and misses



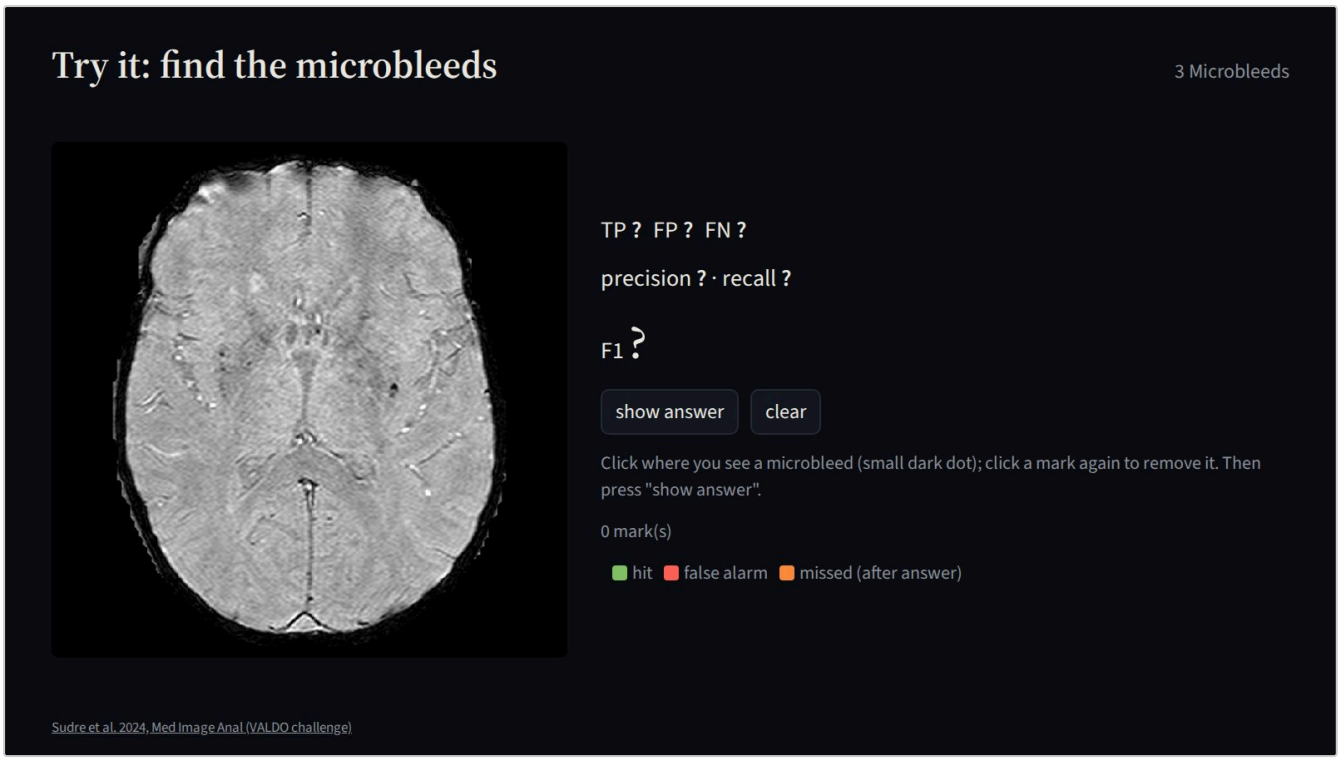
KEY POINTS AND WHAT IS SAID

- **TP: found and real**
  Now we count objects, not voxels. The expert marked five microbleeds: the circles. The program marked five places: the squares. Four are found.
- **FP: false alarm**
  One is a false alarm: a vessel.
- **FN: missed**
  And one bleed is missed.
- **precision = TP / (TP + FP)**
  Precision: how many marks are real.
- **recall = TP / (TP + FN)**
  Recall: how many bleeds we found. F1 combines both.
- **TP 4, FP 1, FN 1 → F1 0.8**
  Four found, one false, one missed: F1 is zero point eight.

SOURCES

- Sudre et al. 2024, Med Image Anal (VALDO challenge)
  VALDO challenge (MICCAI 2021): automatic detection of small vessel lesions, including microbleeds.

43. Try it: find the microbleeds



> Interactive page: try it in the online deck – mendeltem.github.io/mri-segmentation-lectures/lecture_1118/deck/#try_f1

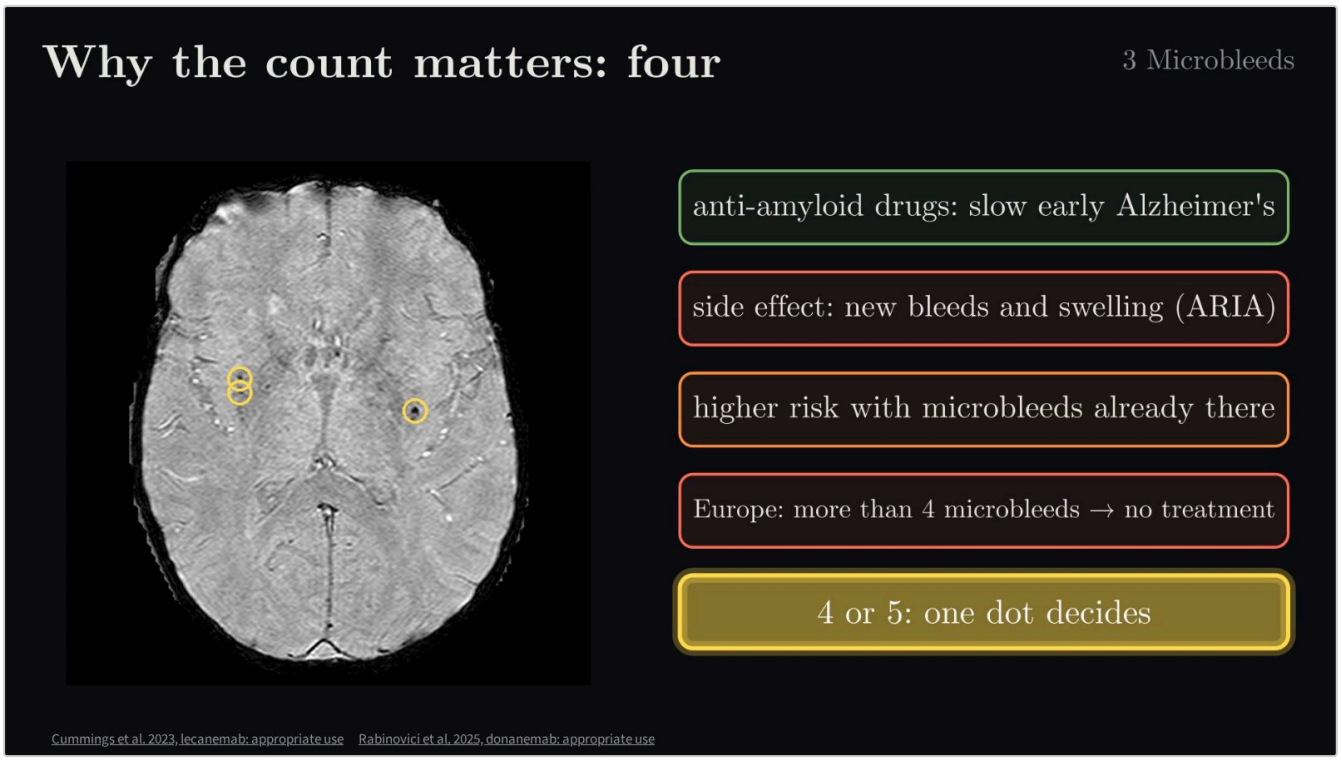
WHAT IS SAID

Click where you see a microbleed. Green: hit. Red: false alarm. 'Show answer' marks the expert's microbleeds and the misses. Precision, recall and F1 update with every click.

SOURCES

- Sudre et al. 2024, Med Image Anal (VALDO challenge)
  VALDO challenge (MICCAI 2021): automatic detection of small vessel lesions, including microbleeds.

44. Why the count matters: four



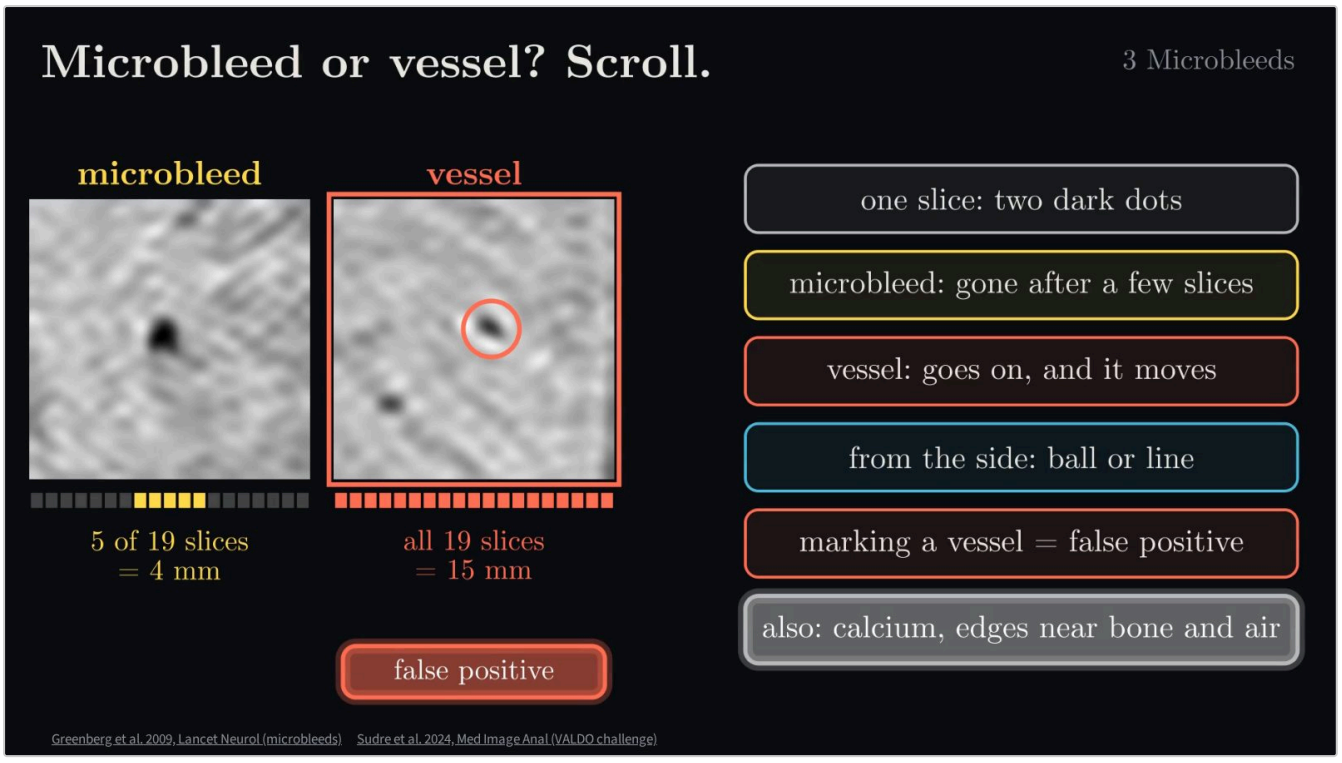
KEY POINTS AND WHAT IS SAID

- **anti-amyloid drugs: slow early Alzheimer's**
  Why count so carefully?
- **side effect: new bleeds and swelling (ARIA)**
  New Alzheimer drugs remove amyloid, but they can cause swelling and new bleeds in the brain.
- **higher risk with microbleeds already there**
  The risk is higher in people who already have microbleeds.
- **Europe: more than 4 microbleeds → no treatment**
  In Europe, with more than four, the drug is not given.
- **4 or 5: one dot decides**
  Four or five: one dot decides.

SOURCES

- Cummings et al. 2023, lecanemab: appropriate use
  Appropriate use recommendations for lecanemab: no treatment with more than 4 microhaemorrhages.
- Rabinovici et al. 2025, donanemab: appropriate use
  Appropriate use recommendations for donanemab: same limit of more than 4 microbleeds; superficial siderosis excludes.

45. Microbleed or vessel? Scroll.



KEY POINTS AND WHAT IS SAID

- **one slice: two dark dots**
  Is every dark dot a microbleed? No. On a single slice, a microbleed and a vessel cut across look the same.
- **microbleed: gone after a few slices**
  So scroll through the slices. The microbleed is gone after a few slices.
- **vessel: goes on, and it moves**
  The vessel goes on, slice after slice, and it moves.
- **from the side: ball or line**
  From the side, the microbleed is a ball, and the vessel is a line.
- **marking a vessel = false positive**
  Mark the vessel, and you have a false positive.
- **also: calcium, edges near bone and air**
  Calcium, and the edges near bone and air, can fool us the same way.

SOURCES

- Greenberg et al. 2009, Lancet Neurol (microbleeds)
  Review of cerebral microbleeds: definition, detection on T2*/SWI, mimics and clinical meaning.
- Sudre et al. 2024, Med Image Anal (VALDO challenge)
  VALDO challenge (MICCAI 2021): automatic detection of small vessel lesions, including microbleeds.

46. Your turn: find the microbleeds



> **Exercise: Microbleed exercise** (9 min) – 3 scans. Click every microbleed, then get your F1.
> mendeltem.github.io/mri-segmentation-lectures/lecture_1118/exercises/cmb-f1/

KEY POINTS

- **Microbleed exercise: 3 scans**
- **click every microbleed → your F1**

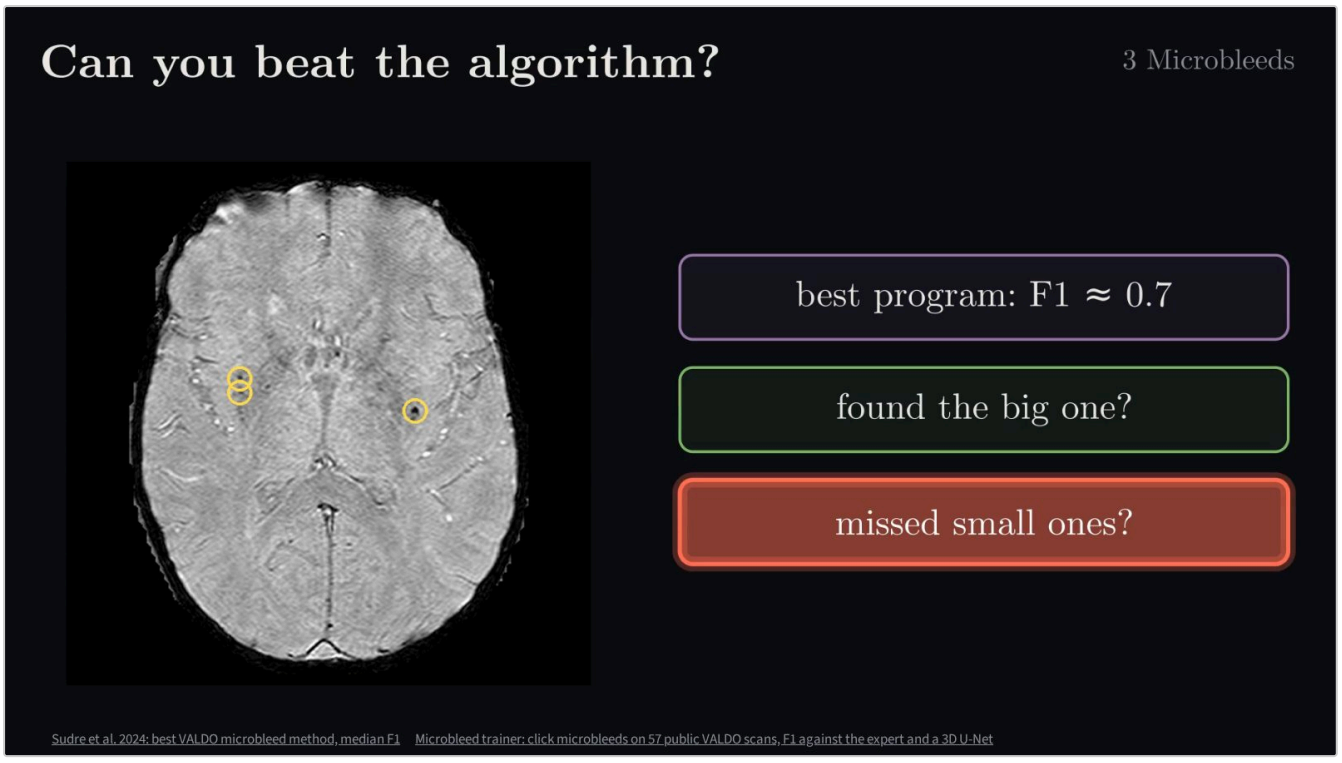
WHAT IS SAID

Last exercise. Three scans. Scroll through the slices and click on every microbleed you find. Then you get your precision, your recall and your F1. You have nine minutes.

SOURCES

- Sudre et al. 2024, Med Image Anal (VALDO challenge)
  VALDO challenge (MICCAI 2021): automatic detection of small vessel lesions, including microbleeds.

47. Can you beat the algorithm?



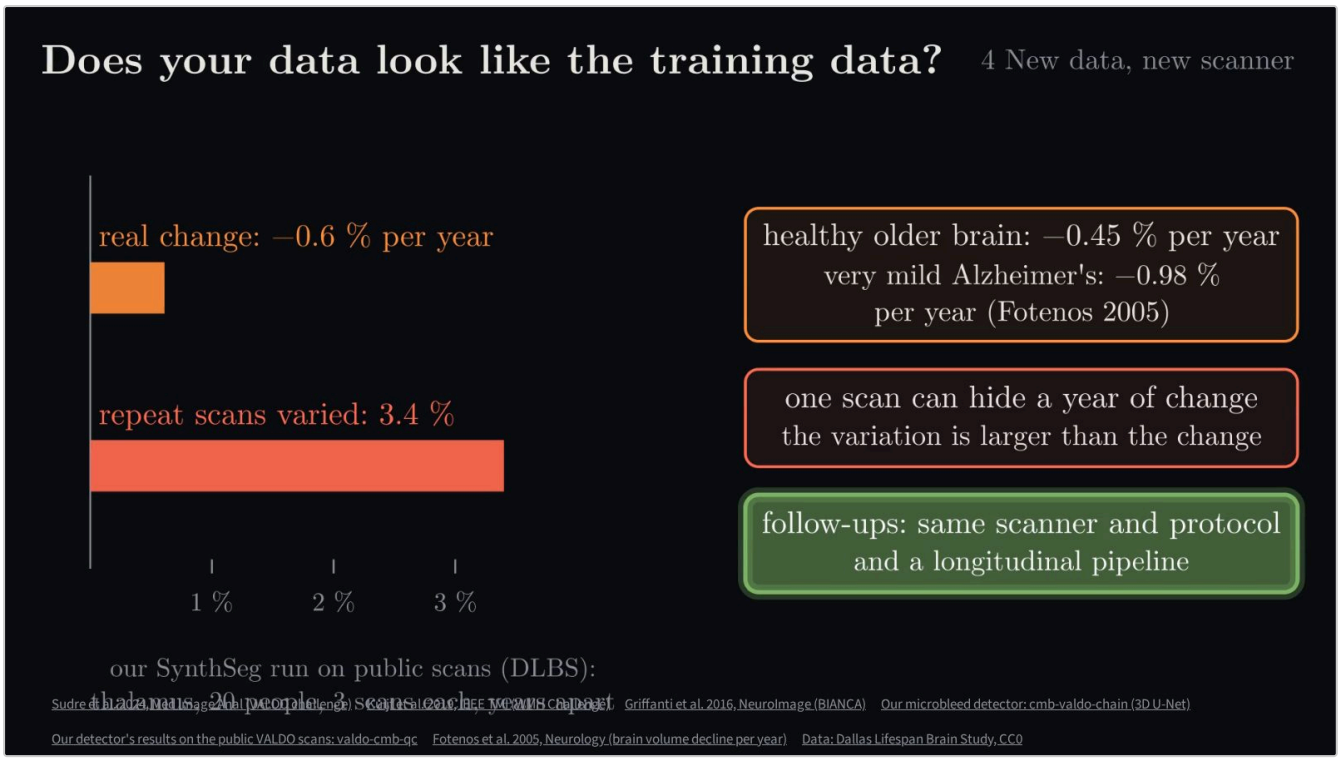
KEY POINTS AND WHAT IS SAID

- **best program: F1 ≈ 0.7**
  The best program reached an F1 of about zero point seven.
- **found the big one?**
  Hands up: who did better? Who found the big one,
- **missed small ones?**
  and who missed small ones?

SOURCES

- Sudre et al. 2024: best VALDO microbleed method, median F1
  VALDO results: the best microbleed method reached a median F1 of about 0.68 per scan.
- Microbleed trainer: click microbleeds on 57 public VALDO scans, F1 against the expert and a 3D U-Net
  Practice at home: click every microbleed on the 57 public VALDO cross-validation scans, then see your F1 against the expert and against a 3D U-Net.

48. Does your data look like the training data?



KEY POINTS AND WHAT IS SAID

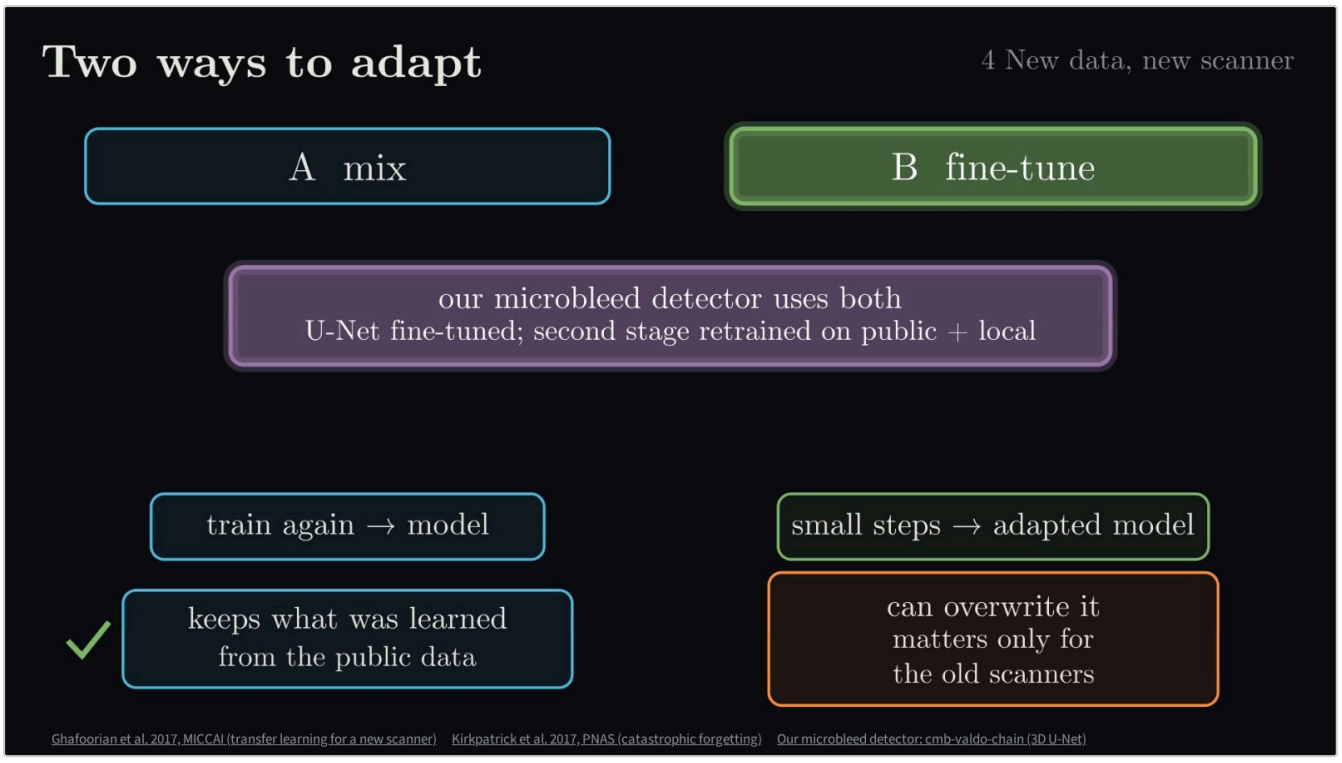
- **what was it trained on?** – people, scanner, sequence, slice thickness, number and size of the lesions
  Before you use a trained tool, ask: what was it trained on? Check the people, the scanner, the sequence, the slice thickness, and the number and size of the lesions.
- **one challenge, three sites** – 1.5 or 3 T; slices from 0.8 to 4 mm
  Even inside one challenge, the three sites differ: one point five or three tesla, slices from under one millimetre to four.
- **one site: 3 scans hold 86 % of its microbleeds** – 91 of 106
  At one site, three scans hold most of the microbleeds.
- **and your scans?** – scanner? sequence? slice thickness? patients?
  And your scans?
- **BIANCA trained on 3 scanners**
  A second example: white matter lesions. BIANCA was trained on scans from three scanners.
- **the same scanners: Dice 0.70** – 90 test scans
  On test scans from these scanners, the Dice was zero point seven.
- **2 new scanners: Dice 0.66** – 20 test scans: 1.5 T GE 0.62, 3 T Philips 0.71
  On two new scanners, zero point six six.
- **small gap here** – but you only know it after you test
  One new scanner was worse, the other just as good. The gap is small here. But you only know it after you test.

- **over 5 mm: 6 of 8 found**
  And size matters. Our microbleed detector, on public scans it had not seen in training. Over five millimetres, six of eight were found.
- **2–5 mm: 79 of 121 found (65 %)**
  Two to five millimetres: sixty five percent.
- **under 2 mm: only 4 of 10 found** – size = diameter of a ball with the same volume
  Under two millimetres: only four of ten. The small ones fail first.
- **healthy older brain: −0.45 % per year** – very mild Alzheimer's: −0.98 % per year (Fotenos 2005)
  And for follow-ups: a brain loses often less than one percent per year.
- **our run, public scans: thalamus −0.6 %/year** – repeat scans varied by 3.4 %
  In our run on public data, the thalamus shrank about zero point six percent per year, but repeat scans varied by about three point four percent.
- **follow-ups: same scanner and protocol** – and a longitudinal pipeline
  So use the same scanner and protocol, and a longitudinal pipeline.

SOURCES

- Sudre et al. 2024, Med Image Anal (VALDO challenge)
  VALDO challenge (MICCAI 2021): automatic detection of small vessel lesions, including microbleeds.
- Kuijf et al. 2019, IEEE TMI (WMH Challenge)
  WMH Segmentation Challenge: 60 training and 110 test scans, 5 scanners; best method Dice 0.80, second experts 0.77-0.79.
- Griffanti et al. 2016, NeuroImage (BIANCA)
  BIANCA: k-nearest-neighbour WMH segmentation in FSL; Dice 0.76 in one cohort, 0.52 in another.
- Our microbleed detector: cmb-valdo-chain (3D U-Net)
  Our 3D U-Net microbleed detector, trained on the public VALDO data; code and decision rules on GitHub. On the slide: its out-of-fold results on the 57 public cross-validation scans.
- Our detector's results on the public VALDO scans: valdo-cmb-qc
  Quality-control pages of our detector on the public VALDO scans: every detection and every miss, scan by scan.
- Fotenos et al. 2005, Neurology (brain volume decline per year)
  370 adults aged 18-97: whole-brain volume (corrected for head size) fell by 0.45 % per year in older people without dementia, and by 0.98 % per year with very mild Alzheimer's disease.
- Data: Dallas Lifespan Brain Study, CC0
  Dallas Lifespan Brain Study: public MRI data (CC0), used for our grey-matter pictures.

49. Two ways to adapt



KEY POINTS AND WHAT IS SAID

- **A mix** – your marked scans + the public data, train again; needs the training data
  Your scans differ, and the score drops. There are two ways to adapt a tool. Way one: mix. Put a few of your own marked scans into the training pool, together with the public data, and train again. For this, you need the training data.
- **B fine-tune** – only the model (downloaded weights); go on training in small steps
  Way two: fine-tune. You only have the trained model: the downloaded weights. You go on training it on your marked scans, in small steps.
- **mixing keeps what was learned**
  Mixing keeps what the model learned from the public data.
- **fine-tuning can overwrite it** – matters only if the old scanners still count
  Fine-tuning can overwrite it. That matters only if the model must still work on the old scanners.
- **our microbleed detector uses both** – U-Net fine-tuned; second stage retrained on public + local
  Our microbleed detector uses both ways: the U-Net is fine-tuned, and its second stage is retrained on public and local examples together.

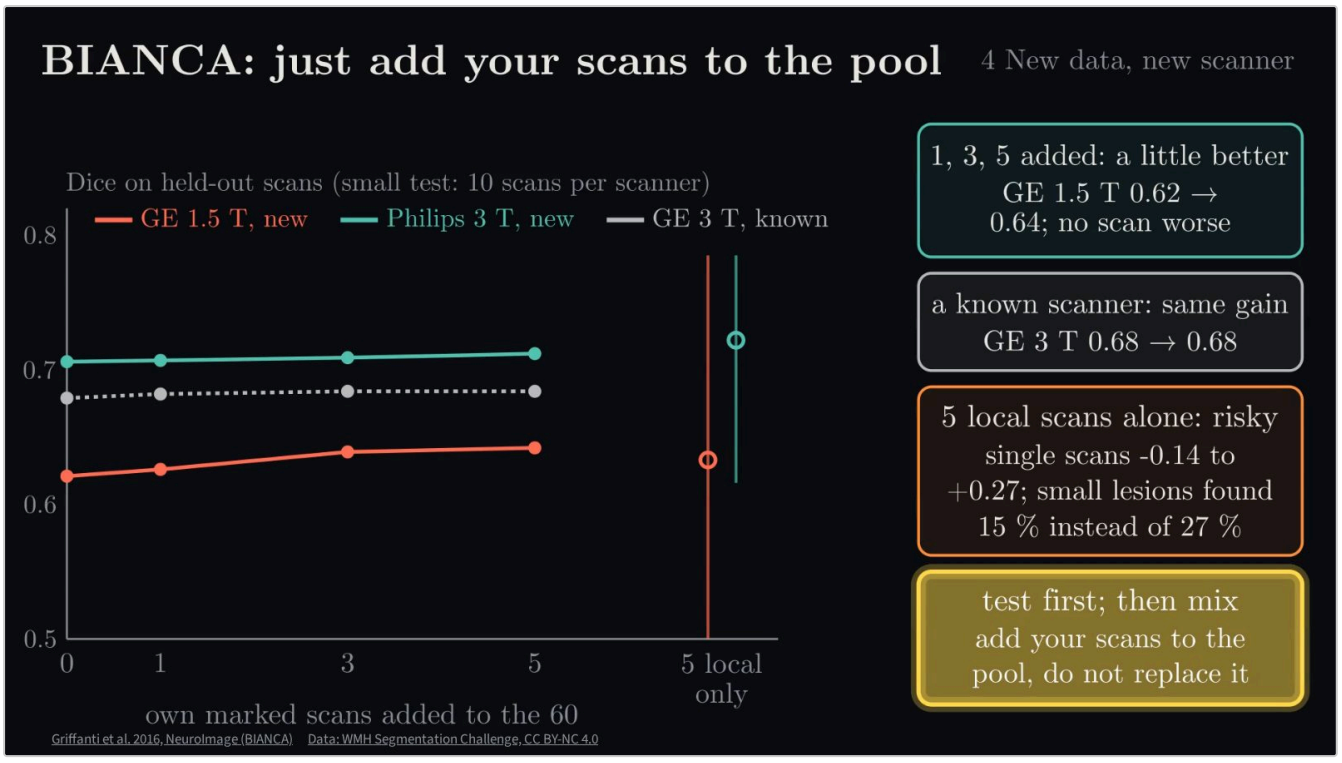
SOURCES

- Ghafoorian et al. 2017, MICCAI (transfer learning for a new scanner)
  White matter lesions: a network trained on old data, fine-tuned with only 2 scans from the new scanner, reached Dice 0.63 and clearly beat the same network trained from scratch on those 2 scans.
- Kirkpatrick et al. 2017, PNAS (catastrophic forgetting)

A network trained on a new task in sequence tends to forget the old one (catastrophic forgetting); the paper proposes a way to protect the weights that mattered for the old task.

- Our microbleed detector: cmb-valdo-chain (3D U-Net)
  Our 3D U-Net microbleed detector, trained on the public VALDO data; code and decision rules on GitHub. On the slide: its out-of-fold results on the 57 public cross-validation scans.

50. BIANCA: just add your scans to the pool



KEY POINTS AND WHAT IS SAID

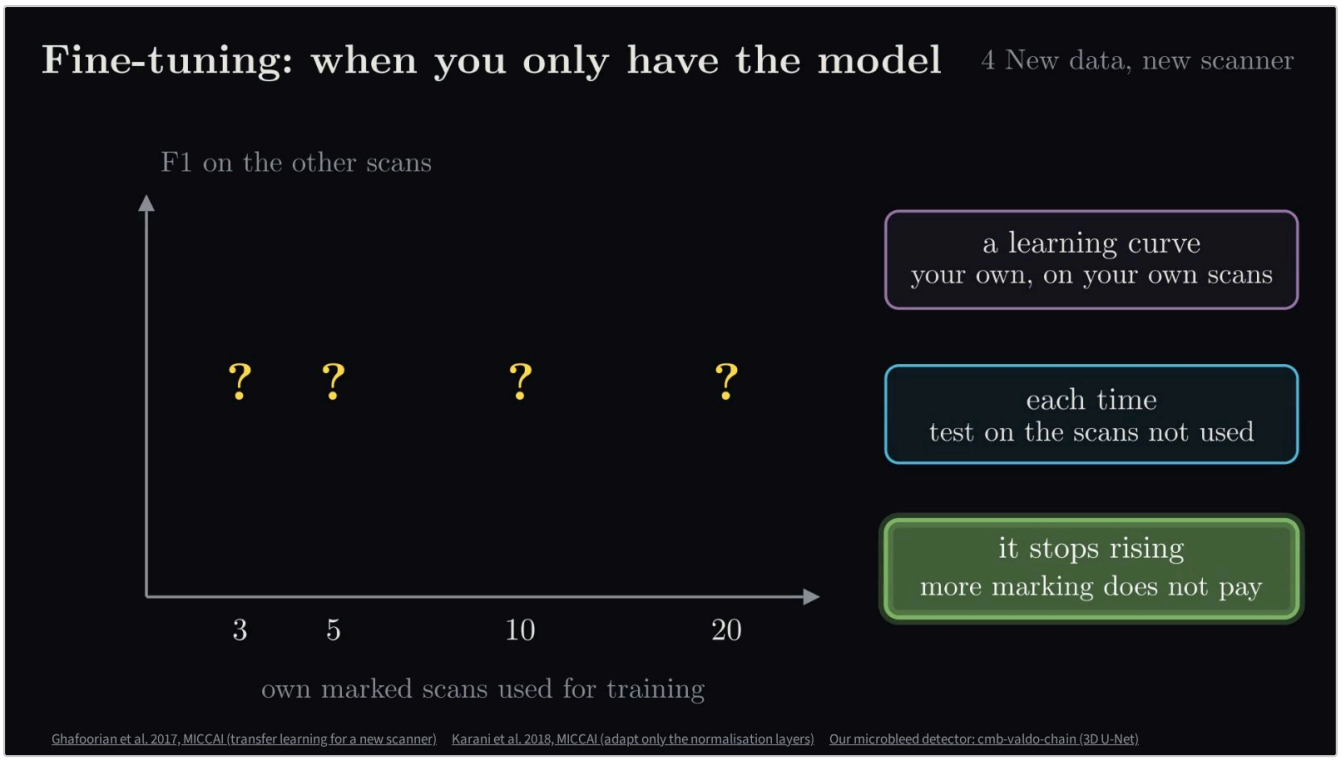
- **the model = the stored points** – BIANCA keeps the marked example voxels
  For BIANCA, mixing is the natural way. Its model is nothing but the stored example points.
- **new scanner: brightness shifts** – a lesion voxel lands among healthy points (drawn, not measured)
  On a new scanner, the brightness shifts, so a lesion voxel can land among healthy points, and its neighbours vote wrong.
- **add marked points of the new scanner**
  Add marked points from the new scanner,
- **the 40 nearest vote right again**
  and the vote is right again. Retraining takes less than a minute.
- **2 new scanners, 0 added: 0.62 and 0.71** – small test: 10 scans per scanner, held out in turn
  A small test on the two new scanners of the challenge, ten scans each. Without local scans, the Dice was zero point six two and zero point seven one.
- **1, 3, 5 added: a little better** – GE 1.5 T 0.62 → 0.64; no scan worse
  Adding one, three or five marked scans helped only a little, and no scan got worse.
- **a known scanner: the same small gain**
  A scanner BIANCA already knew gained about the same.
- **5 local scans alone: risky** – single scans −0.14 to +0.27; fewer small lesions found
  Trained on the five local scans alone, single scans jumped up and down, and fewer small lesions were found.

- **test first; then mix, do not replace**
  So: test first. If you adapt, mix your scans into the pool. Do not replace it.

SOURCES

- Griffanti et al. 2016, NeuroImage (BIANCA)
  BIANCA: k-nearest-neighbour WMH segmentation in FSL; Dice 0.76 in one cohort, 0.52 in another.
- Data: WMH Segmentation Challenge, CC BY-NC 4.0
  The public WMH Challenge data (DataverseNL), licence CC BY-NC 4.0; used in the lesion exercise.

51. Fine-tuning: when you only have the model



KEY POINTS AND WHAT IS SAID

- **a trained model knows the target** – here: what a microbleed looks like
  Fine-tuning, step by step. A trained model already knows what a microbleed looks like.
- **fine-tuning: start from its weights** – train a little more, in small steps, on a few own marked scans
  Fine-tuning keeps this knowledge: it starts from the model's weights, and trains a little more, in small steps, on a few of your own marked scans. So it learns your scanner's contrast and noise.
- **from scratch: random weights** – needs many more marked scans
  Training from scratch starts from random weights, and needs many more marked scans.
- **1 a new threshold** – needs only own scans without microbleeds
  The options, cheapest first. One: choose a new threshold on your own scans. For this you only need scans without microbleeds.
- **2 adapt the normalisation layers**
  Two: adapt only the normalisation layers.
- **3 fine-tune the whole model**
  Three: fine-tune the whole model.
- **4 train from scratch**
  Four: train from scratch.
- **always: test on own scans** – scans not used for training
  And always test on your own scans that were not used for training.
- **how many marked scans? measure it** – a learning curve

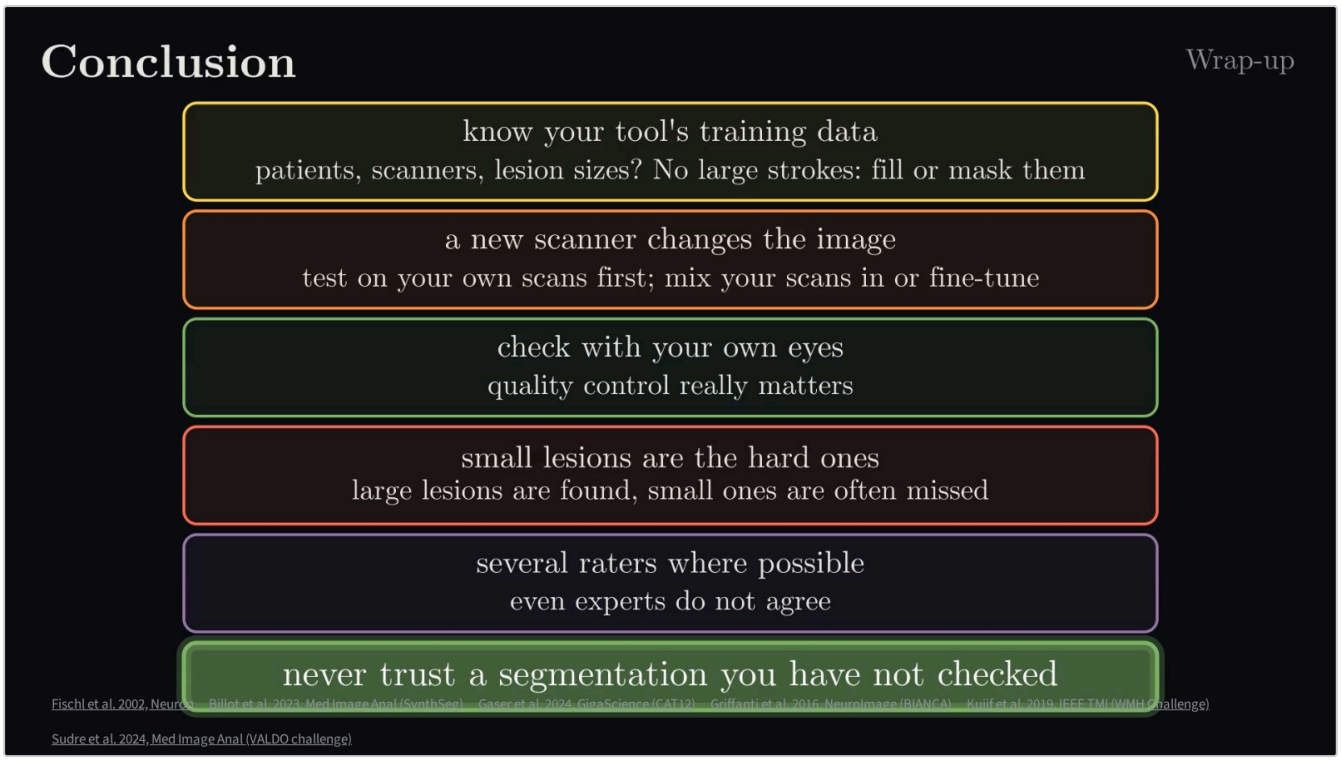
How many marked scans do you need? Measure it, with a learning curve.

- **train with 3, 5, 10, 20 scans** – test each time on the rest
  Train with three, five, ten, twenty scans, and each time test on the rest.
- **the curve stops rising: enough**
  Where the curve stops rising, more marking does not pay.

SOURCES

- Ghafoorian et al. 2017, MICCAI (transfer learning for a new scanner)
  White matter lesions: a network trained on old data, fine-tuned with only 2 scans from the new scanner, reached Dice 0.63 and clearly beat the same network trained from scratch on those 2 scans.
- Karani et al. 2018, MICCAI (adapt only the normalisation layers)
  One network for several scanners: the filters are shared, only the normalisation layers are learned per scanner, from about 4 labelled scans each.
- Our microbleed detector: cmb-valdo-chain (3D U-Net)
  Our 3D U-Net microbleed detector, trained on the public VALDO data; code and decision rules on GitHub. On the slide: its out-of-fold results on the 57 public cross-validation scans.

52. Conclusion



KEY POINTS AND WHAT IS SAID

- **know your tool's training data** – patients, scanners, lesion sizes? No large strokes: fill or mask them

  To conclude. First, know your tool's training data: which patients, which scanners, which lesion sizes? None of these tools was trained on large stroke lesions, so fill or mask the lesion.

- **a new scanner changes the image** – test on your own scans first; mix your scans in or fine-tune

  Second, a new scanner changes the image. Test the tool on your own scans first. If it drops, mix your scans into the training, or fine-tune.

- **check with your own eyes** – quality control really matters

  Third, check with your own eyes: quality control really matters.

- **small lesions are the hard ones** – large lesions are found, small ones are often missed

  Fourth, small lesions are the problem. Large lesions are found, the small ones are often missed.

- **several raters where possible** – even experts do not agree

  And fifth, if you can, use several raters, because even experts do not agree.

- **never trust a segmentation you have not checked**

  And one rule for all: never trust a segmentation you have not checked.

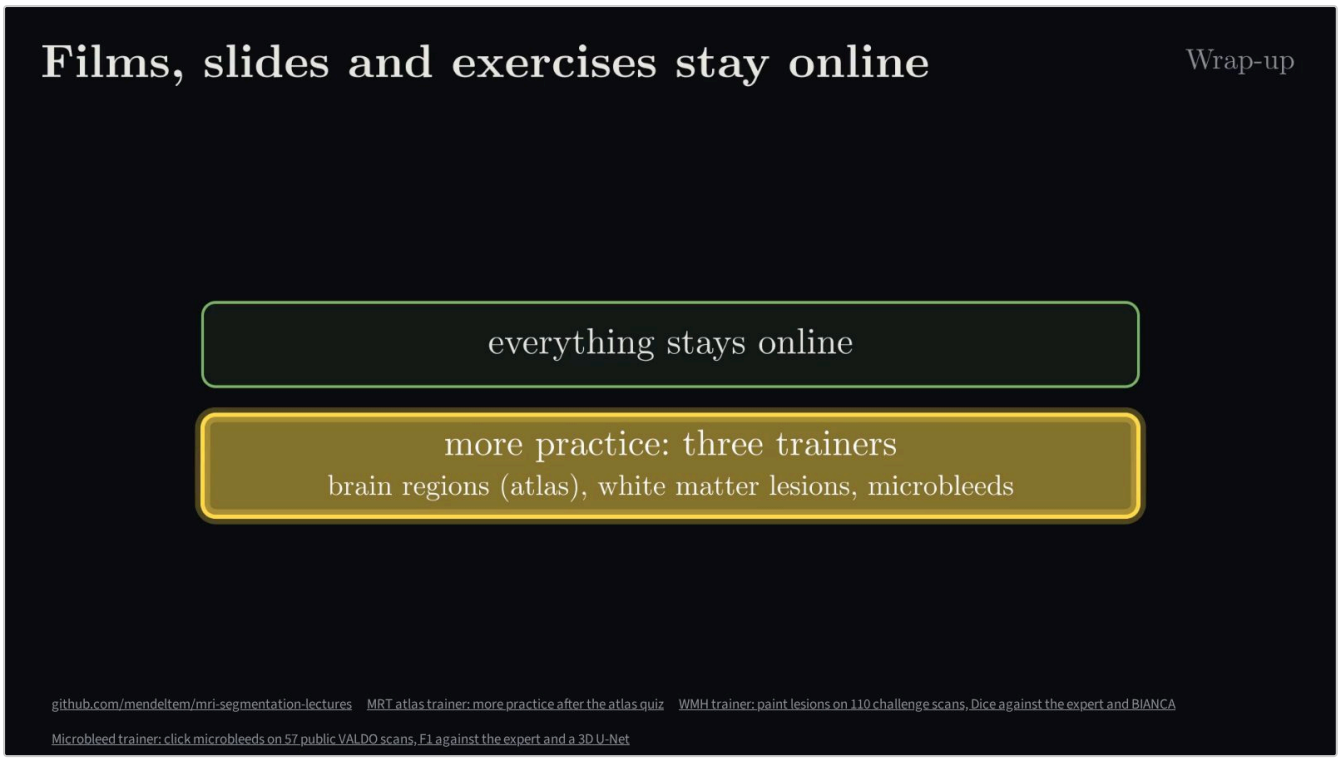
SOURCES

- Fischl et al. 2002, Neuron

  FreeSurfer's whole-brain segmentation of deep structures with a probabilistic atlas (37 labels).

- Billot et al. 2023, Med Image Anal (SynthSeg)
  SynthSeg: a network trained only on synthetic images; segments scans of any contrast and resolution.
- Gaser et al. 2024, GigaScience (CAT12)
  CAT12, the computational anatomy toolbox for SPM: its average template was made from the public IXI data set of healthy adults (the paper's figure registers 555 brains).
- Griffanti et al. 2016, NeuroImage (BIANCA)
  BIANCA: k-nearest-neighbour WMH segmentation in FSL; Dice 0.76 in one cohort, 0.52 in another.
- Kuijf et al. 2019, IEEE TMI (WMH Challenge)
  WMH Segmentation Challenge: 60 training and 110 test scans, 5 scanners; best method Dice 0.80, second experts 0.77-0.79.
- Sudre et al. 2024, Med Image Anal (VALDO challenge)
  VALDO challenge (MICCAI 2021): automatic detection of small vessel lesions, including microbleeds.

53. Films, slides and exercises stay online



KEY POINTS AND WHAT IS SAID

- **everything stays online**
  Everything stays online.
- **more practice: three trainers** – brain regions (atlas), white matter lesions, microbleeds
  For more practice, there are three trainers: brain regions, lesions, and microbleeds. Thank you.

SOURCES

- github.com/mendeltem/mri-segmentation-lectures
  The lecture's GitHub repository with slides, films and exercises.
- MRT atlas trainer: more practice after the atlas quiz
  The full atlas trainer for home: explore, quiz and study modes with many more brain regions (interface in German).
- WMH trainer: paint lesions on 110 challenge scans, Dice against the expert and BIANCA
  Practice at home: paint the WMH on 6 FLAIR slices of each of the 110 public WMH Challenge test scans, then see your Dice against the expert and against BIANCA. Works on phones.
- Microbleed trainer: click microbleeds on 57 public VALDO scans, F1 against the expert and a 3D U-Net
  Practice at home: click every microbleed on the 57 public VALDO cross-validation scans, then see your F1 against the expert and against a 3D U-Net.

All sources

1. Bethlehem et al. 2022, Nature (brain charts) – slide 3
   Brain charts for the human lifespan: 123,984 scans of 101,457 people, from 115 days after conception to 100 years; gives centiles like a growth chart.
   https://doi.org/10.1038/s41586-022-04554-y

2. Billot et al. 2023, Med Image Anal (SynthSeg) – slide 9, 10, 11, 12, 18, 52
   SynthSeg: a network trained only on synthetic images; segments scans of any contrast and resolution.
   https://doi.org/10.1016/j.media.2023.102789

3. Billot et al. 2023, PNAS (robust SynthSeg, QC) – slide 13, 14
   SynthSeg+ (robust version) with automatic QC scores per region; below 0.65 a region counts as failed.
   https://doi.org/10.1073/pnas.2216399120

4. Charidimou et al. 2022, Lancet Neurol (Boston criteria 2.0) – slide 38
   Boston criteria 2.0 for cerebral amyloid angiopathy: strictly lobar haemorrhagic lesions (bleeds, microbleeds, superficial siderosis) make the diagnosis.
   https://doi.org/10.1016/S1474-4422(22)00208-3

5. Cummings et al. 2023, lecanemab: appropriate use – slide 21, 44
   Appropriate use recommendations for lecanemab: no treatment with more than 4 microhaemorrhages.
   https://doi.org/10.14283/jpad.2023.30

6. Dale, Fischl, Sereno 1999, NeuroImage – slide 4, 7
   FreeSurfer part I: segmentation and the white and pial surface of the cortex.
   https://doi.org/10.1006/nimg.1998.0395

7. Data: Aphasia Recovery Cohort, CC0 – slide 14, 17, 18
   Aphasia Recovery Cohort: public stroke MRI with expert lesion masks (CC0).
   https://openneuro.org/datasets/ds004884

8. Data: Dallas Lifespan Brain Study, CC0 – slide 2, 13, 17, 18, 48
   Dallas Lifespan Brain Study: public MRI data (CC0), used for our grey-matter pictures.
   https://doi.org/10.18112/openneuro.ds004856.v1.3.0

9. Data: WMH Segmentation Challenge, CC BY-NC 4.0 – slide 20, 22, 24, 25, 26, 27, 28, 31, 35, 50
   The public WMH Challenge data (DataverseNL), licence CC BY-NC 4.0; used in the lesion exercise.
   https://doi.org/10.34894/AECRSD

10. Debette & Markus 2010, BMJ – slide 20
    Meta-analysis: white matter hyperintensities raise the risk of stroke (about 3x) and dementia (about 2x).
    https://doi.org/10.1136/bmj.c3666

11. Desikan et al. 2006, NeuroImage – slide 6
    Automatic labelling of 34 cortical regions per hemisphere, learned from 40 hand-labelled brains.
    https://doi.org/10.1016/j.neuroimage.2006.01.021

12. Dice 1945, Ecology – slide 31
    Dice's 1945 overlap coefficient, first used to measure how often two species occur together.
    https://doi.org/10.2307/1932409

13. Fischl et al. 2002, Neuron – slide 52
    FreeSurfer's whole-brain segmentation of deep structures with a probabilistic atlas (37 labels).
    https://doi.org/10.1016/S0896-6273(02)00569-X

14. Fischl, Dale 2000, PNAS (cortical thickness) – slide 5

How FreeSurfer measures cortical thickness from the two surfaces; differences between people below 0.5 mm.
https://doi.org/10.1073/pnas.200033797

15. Fischl, Sereno, Dale 1999, NeuroImage – slide 6
    FreeSurfer part II: the cortex inflated and mapped to a sphere; folding patterns aligned between brains.
    https://doi.org/10.1006/nimg.1998.0396

16. Fotenos et al. 2005, Neurology (brain volume decline per year) – slide 48
    370 adults aged 18-97: whole-brain volume (corrected for head size) fell by 0.45 % per year in older people without dementia, and by 0.98 % per year with very mild Alzheimer's disease.
    https://doi.org/10.1212/01.WNL.0000154530.72969.11

17. FreeSurfer: recon-all – slide 5, 7, 8, 12, 18
    FreeSurfer documentation of recon-all, the command that runs the whole pipeline.
    https://surfer.nmr.mgh.harvard.edu/fswiki/recon-all

18. Gaser et al. 2024, GigaScience (CAT12) – slide 52
    CAT12, the computational anatomy toolbox for SPM: its average template was made from the public IXI data set of healthy adults (the paper's figure registers 555 brains).
    https://doi.org/10.1093/gigascience/giae049

19. Ghafoorian et al. 2017, MICCAI (transfer learning for a new scanner) – slide 49, 51
    White matter lesions: a network trained on old data, fine-tuned with only 2 scans from the new scanner, reached Dice 0.63 and clearly beat the same network trained from scratch on those 2 scans.
    https://doi.org/10.1007/978-3-319-66179-7_59

20. github.com/mendeltem/mri-segmentation-lectures – slide 53
    The lecture's GitHub repository with slides, films and exercises.
    https://github.com/mendeltem/mri-segmentation-lectures

21. Gouw et al. 2008, Cerebrovasc Dis (visual scales vs volume) – slide 21
    Measuring the change of lesion volume was reliable (ICC 0.88); judging progression with visual scales was less reliable (20 scan pairs).
    https://doi.org/10.1159/000113863

22. Grabner et al. 2006, MICCAI (MNI152 standard brain) – slide 26
    The MNI152 standard brain (average of 152 healthy brains): scans are aligned to it, so every voxel gets a position that means the same place in every person.
    https://doi.org/10.1007/11866763_8

23. Greenberg et al. 2009, Lancet Neurol (microbleeds) – slide 37, 45
    Review of cerebral microbleeds: definition, detection on T2*/SWI, mimics and clinical meaning.
    https://doi.org/10.1016/S1474-4422(09)70013-4

24. Griffanti et al. 2016, NeuroImage (BIANCA) – slide 23, 24, 25, 26, 27, 30, 36, 48, 50, 52
    BIANCA: k-nearest-neighbour WMH segmentation in FSL; Dice 0.76 in one cohort, 0.52 in another.
    https://doi.org/10.1016/j.neuroimage.2016.07.018

25. Inzitari et al. 2009, BMJ (LADIS study) – slide 21
    LADIS: older people living independently, 3 years follow-up; disability or death per year 10.5 % (mild), 15.1 % (moderate), 29.5 % (severe white matter changes).
    https://doi.org/10.1136/bmj.b2477

26. Jack et al. 2018, Alzheimers Dement (AT(N) research framework) – slide 3
    Research definition of Alzheimer's disease by biomarkers: amyloid (A), tau (T), neurodegeneration (N). Atrophy on MRI is one (N) marker; not specific, and not meant for routine practice.
    https://doi.org/10.1016/j.jalz.2018.02.018

27. Karani et al. 2018, MICCAI (adapt only the normalisation layers) – slide 51
One network for several scanners: the filters are shared, only the normalisation layers are learned per scanner, from about 4 labelled scans each.
https://doi.org/10.1007/978-3-030-00928-1_54

28. Kirkpatrick et al. 2017, PNAS (catastrophic forgetting) – slide 49
A network trained on a new task in sequence tends to forget the old one (catastrophic forgetting); the paper proposes a way to protect the weights that mattered for the old task.
https://doi.org/10.1073/pnas.1611835114

29. Kuijf et al. 2019, IEEE TMI (WMH Challenge) – slide 22, 29, 30, 32, 36, 48, 52
WMH Segmentation Challenge: 60 training and 110 test scans, 5 scanners; best method Dice 0.80, second experts 0.77-0.79.
https://doi.org/10.1109/TMI.2019.2905770

30. Microbleed trainer: click microbleeds on 57 public VALDO scans, F1 against the expert and a 3D U-Net – slide 47, 53
Practice at home: click every microbleed on the 57 public VALDO cross-validation scans, then see your F1 against the expert and against a 3D U-Net.
https://mendeltem.github.io/cmb_detection_trainer/

31. MRT atlas trainer: more practice after the atlas quiz – slide 19, 53
The full atlas trainer for home: explore, quiz and study modes with many more brain regions (interface in German).
https://mendeltem.github.io/mrt-atlas-trainer/

32. Nachev et al. 2008, NeuroImage (enantiomorphic: mirror the healthy hemisphere) – slide 18
Lesions distort the normalisation of brain scans. The enantiomorphic method replaces the lesion with the mirror image of the healthy homologous region; tested on healthy brains with artificial lesions, it beat cost-function masking.
https://doi.org/10.1016/j.neuroimage.2007.10.002

33. O'Brien et al. 2011, Psychiatry Res (adjusting for head size) – slide 3
Heads differ in size, so do brain structures: volumes must be adjusted for head (intracranial) size. Compares three ways (proportion, residual, covariance) and recommends a general model.
https://doi.org/10.1016/j.pscychresns.2011.01.007

34. Our detector's results on the public VALDO scans: valdo-cmb-qc – slide 48
Quality-control pages of our detector on the public VALDO scans: every detection and every miss, scan by scan.
https://github.com/mendeltem/valdo-cmb-qc

35. Our microbleed detector: cmb-valdo-chain (3D U-Net) – slide 48, 49, 51
Our 3D U-Net microbleed detector, trained on the public VALDO data; code and decision rules on GitHub. On the slide: its out-of-fold results on the 57 public cross-validation scans.
https://github.com/mendeltem/cmb-valdo-chain

36. Rabinovici et al. 2025, donanemab: appropriate use – slide 44
Appropriate use recommendations for donanemab: same limit of more than 4 microbleeds; superficial siderosis excludes.
https://doi.org/10.1016/j.tjpad.2025.100150

37. Ronneberger et al. 2015, MICCAI (U-Net) – slide 29
U-Net: the encoder-decoder network with skip connections for biomedical image segmentation.
https://doi.org/10.1007/978-3-319-24574-4_28

38. Rosen et al. 2018, NeuroImage (Euler number) – slide 15
The Euler number of FreeSurfer's surface flags unusable scans very well (AUC 0.98-0.99).
https://doi.org/10.1016/j.neuroimage.2017.12.059

39. Schmaal et al. 2016, Mol Psychiatry (ENIGMA depression) – slide 3
ENIGMA consortium: 1,728 patients with depression and 7,199 controls from 15 samples; hippocampus 1.24 % smaller (Cohen's d = -0.14).
https://doi.org/10.1038/mp.2015.69

40. scipy.ndimage.label (connected clusters) – slide 40, 41
SciPy function that groups touching voxels into clusters; the neighbourhood (6 or 26) is a choice.
https://docs.scipy.org/doc/scipy/reference/generated/scipy.ndimage.label.html

41. Sormani et al. 2014, Ann Neurol (atrophy in MS trials) – slide 3
Meta-analysis of 13 multiple sclerosis trials (more than 13,500 patients): the more a drug slowed brain atrophy, the more it slowed disability (R2 = 0.48).
https://doi.org/10.1002/ana.24018

42. SPRINT MIND 2019, JAMA (blood pressure and white matter lesions) – slide 21
SPRINT MIND: intensive vs standard blood pressure control; lesion volume 4.57 to 5.49 cm3 vs 4.40 to 5.85 cm3 after about 4 years (difference -0.54 cm3, small).
https://doi.org/10.1001/jama.2019.10551

43. Sudre et al. 2024, Med Image Anal (VALDO challenge) – slide 37, 39, 42, 43, 45, 46, 47, 48, 52
VALDO challenge (MICCAI 2021): automatic detection of small vessel lesions, including microbleeds.
https://doi.org/10.1016/j.media.2023.103029

44. Vernooij et al. 2008, Neurology (Rotterdam Scan Study) – slide 38
Rotterdam Scan Study (1,062 people, 1.5 T): microbleeds in 17.8 % at age 60-69 and 38.3 % over 80; lobar location points to amyloid angiopathy, deep location to hypertensive vessel disease.
https://doi.org/10.1212/01.wnl.0000307750.41970.d9

45. Wardlaw et al. 2013, Lancet Neurol (STRIVE) – slide 20
STRIVE consensus: names, definitions and imaging standards for small vessel disease (WMH, microbleeds).
https://doi.org/10.1016/S1474-4422(13)70124-8

46. Wilson et al. 2018, Lancet Neurol (CROMIS-2) – slide 38
CROMIS-2: anticoagulated patients with atrial fibrillation after stroke or TIA; brain bleeds 9.8 (with microbleeds) vs 2.6 (without) per 1,000 patient-years.
https://doi.org/10.1016/S1474-4422(18)30145-5

47. Wilson et al. 2019, Lancet Neurol (MICON pooled analysis) – slide 38
Pooled data of 20,322 patients after ischaemic stroke or TIA: with 10 or more microbleeds still 64 ischaemic strokes vs 27 brain bleeds per 1,000 patient-years.
https://doi.org/10.1016/S1474-4422(19)30197-8

48. WMH trainer: paint lesions on 110 challenge scans, Dice against the expert and BIANCA – slide 36, 53
Practice at home: paint the WMH on 6 FLAIR slices of each of the 110 public WMH Challenge test scans, then see your Dice against the expert and against BIANCA. Works on phones.
https://mendeltem.github.io/wmh_trainer/